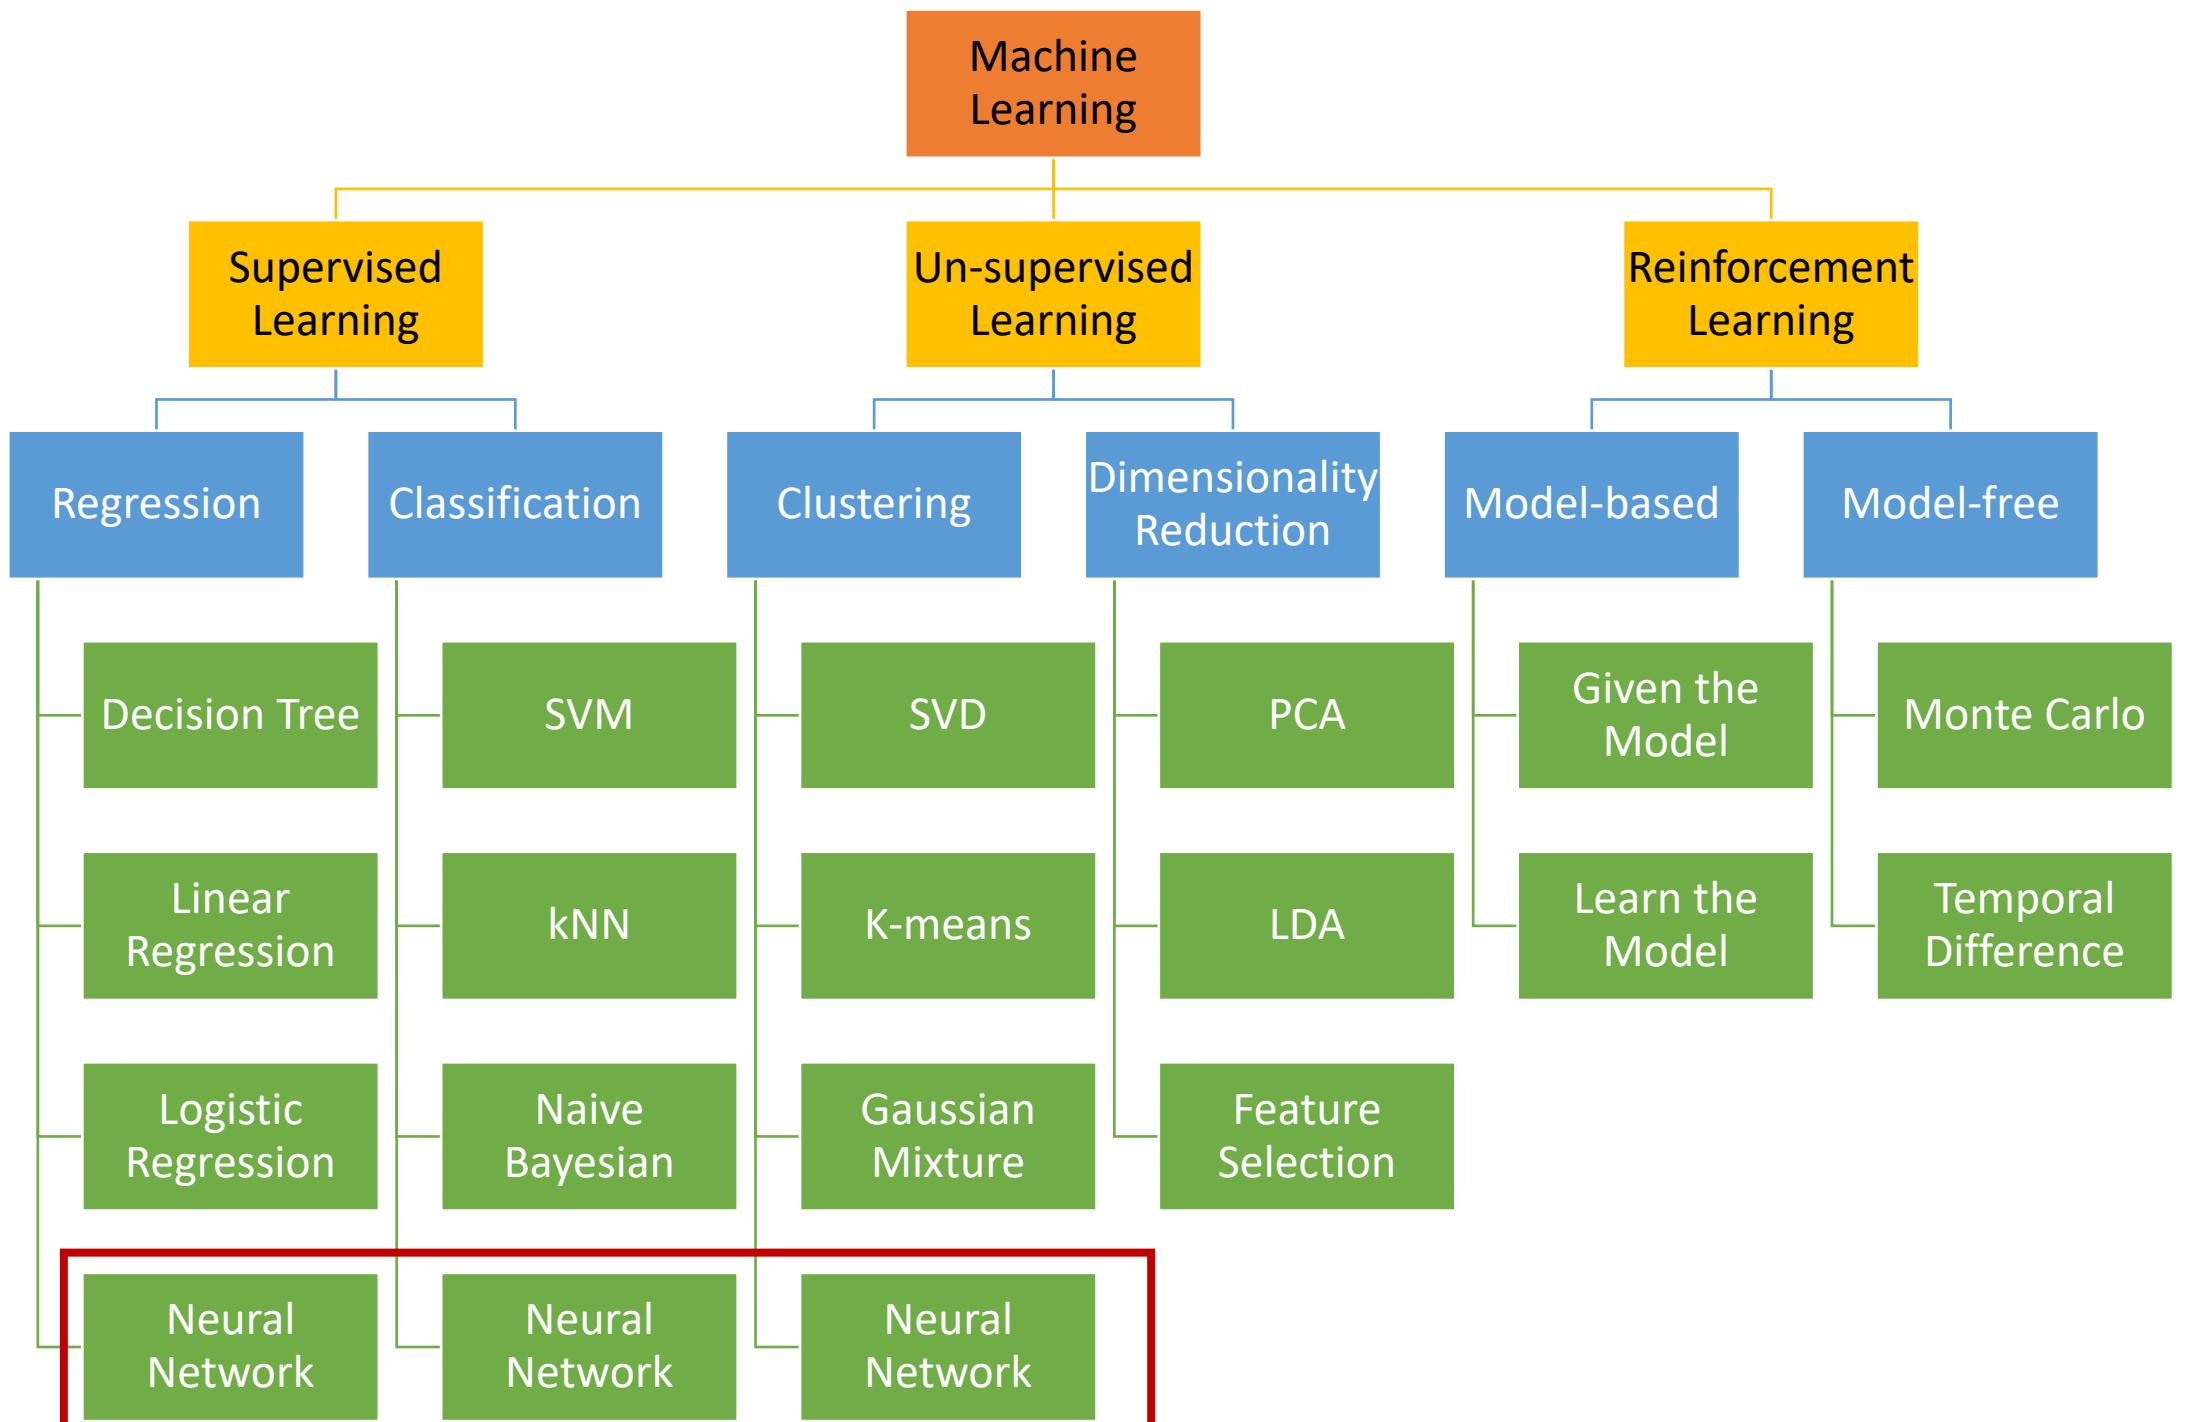
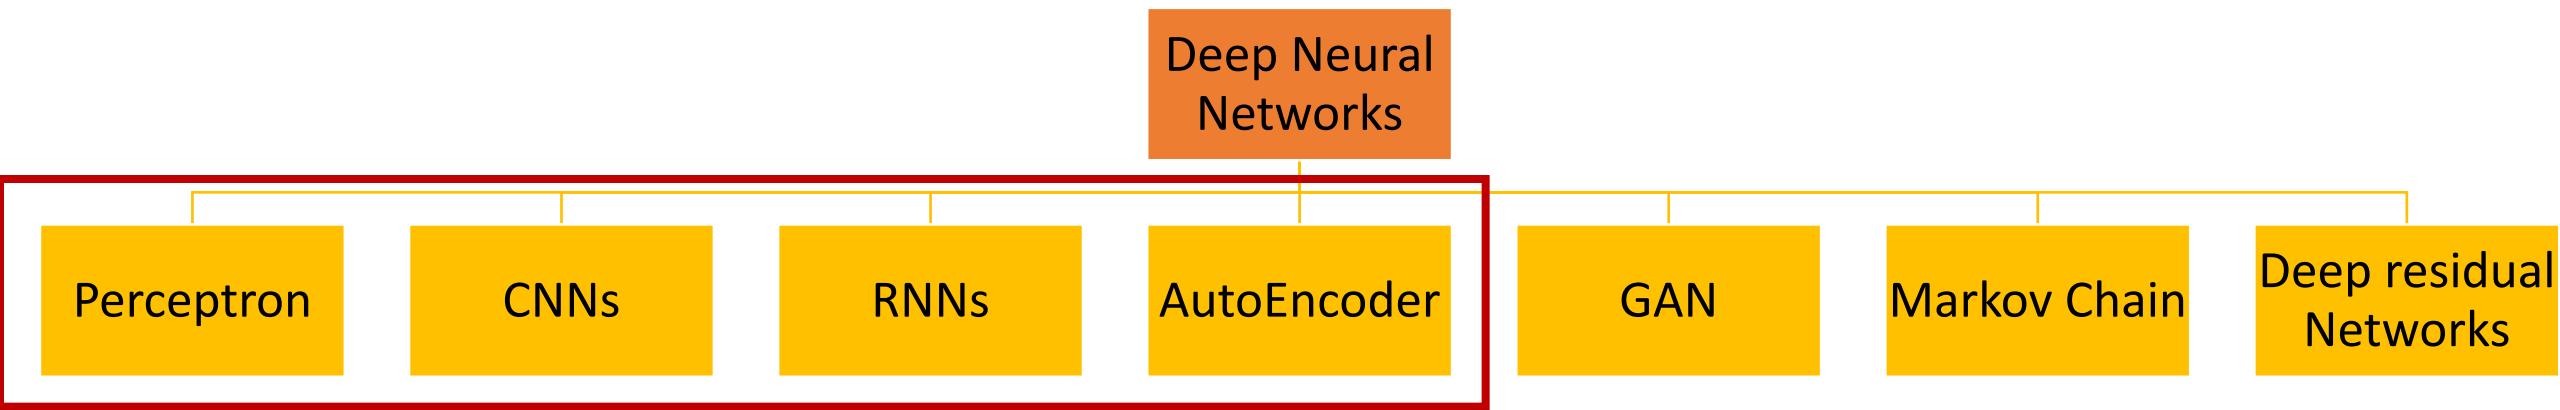


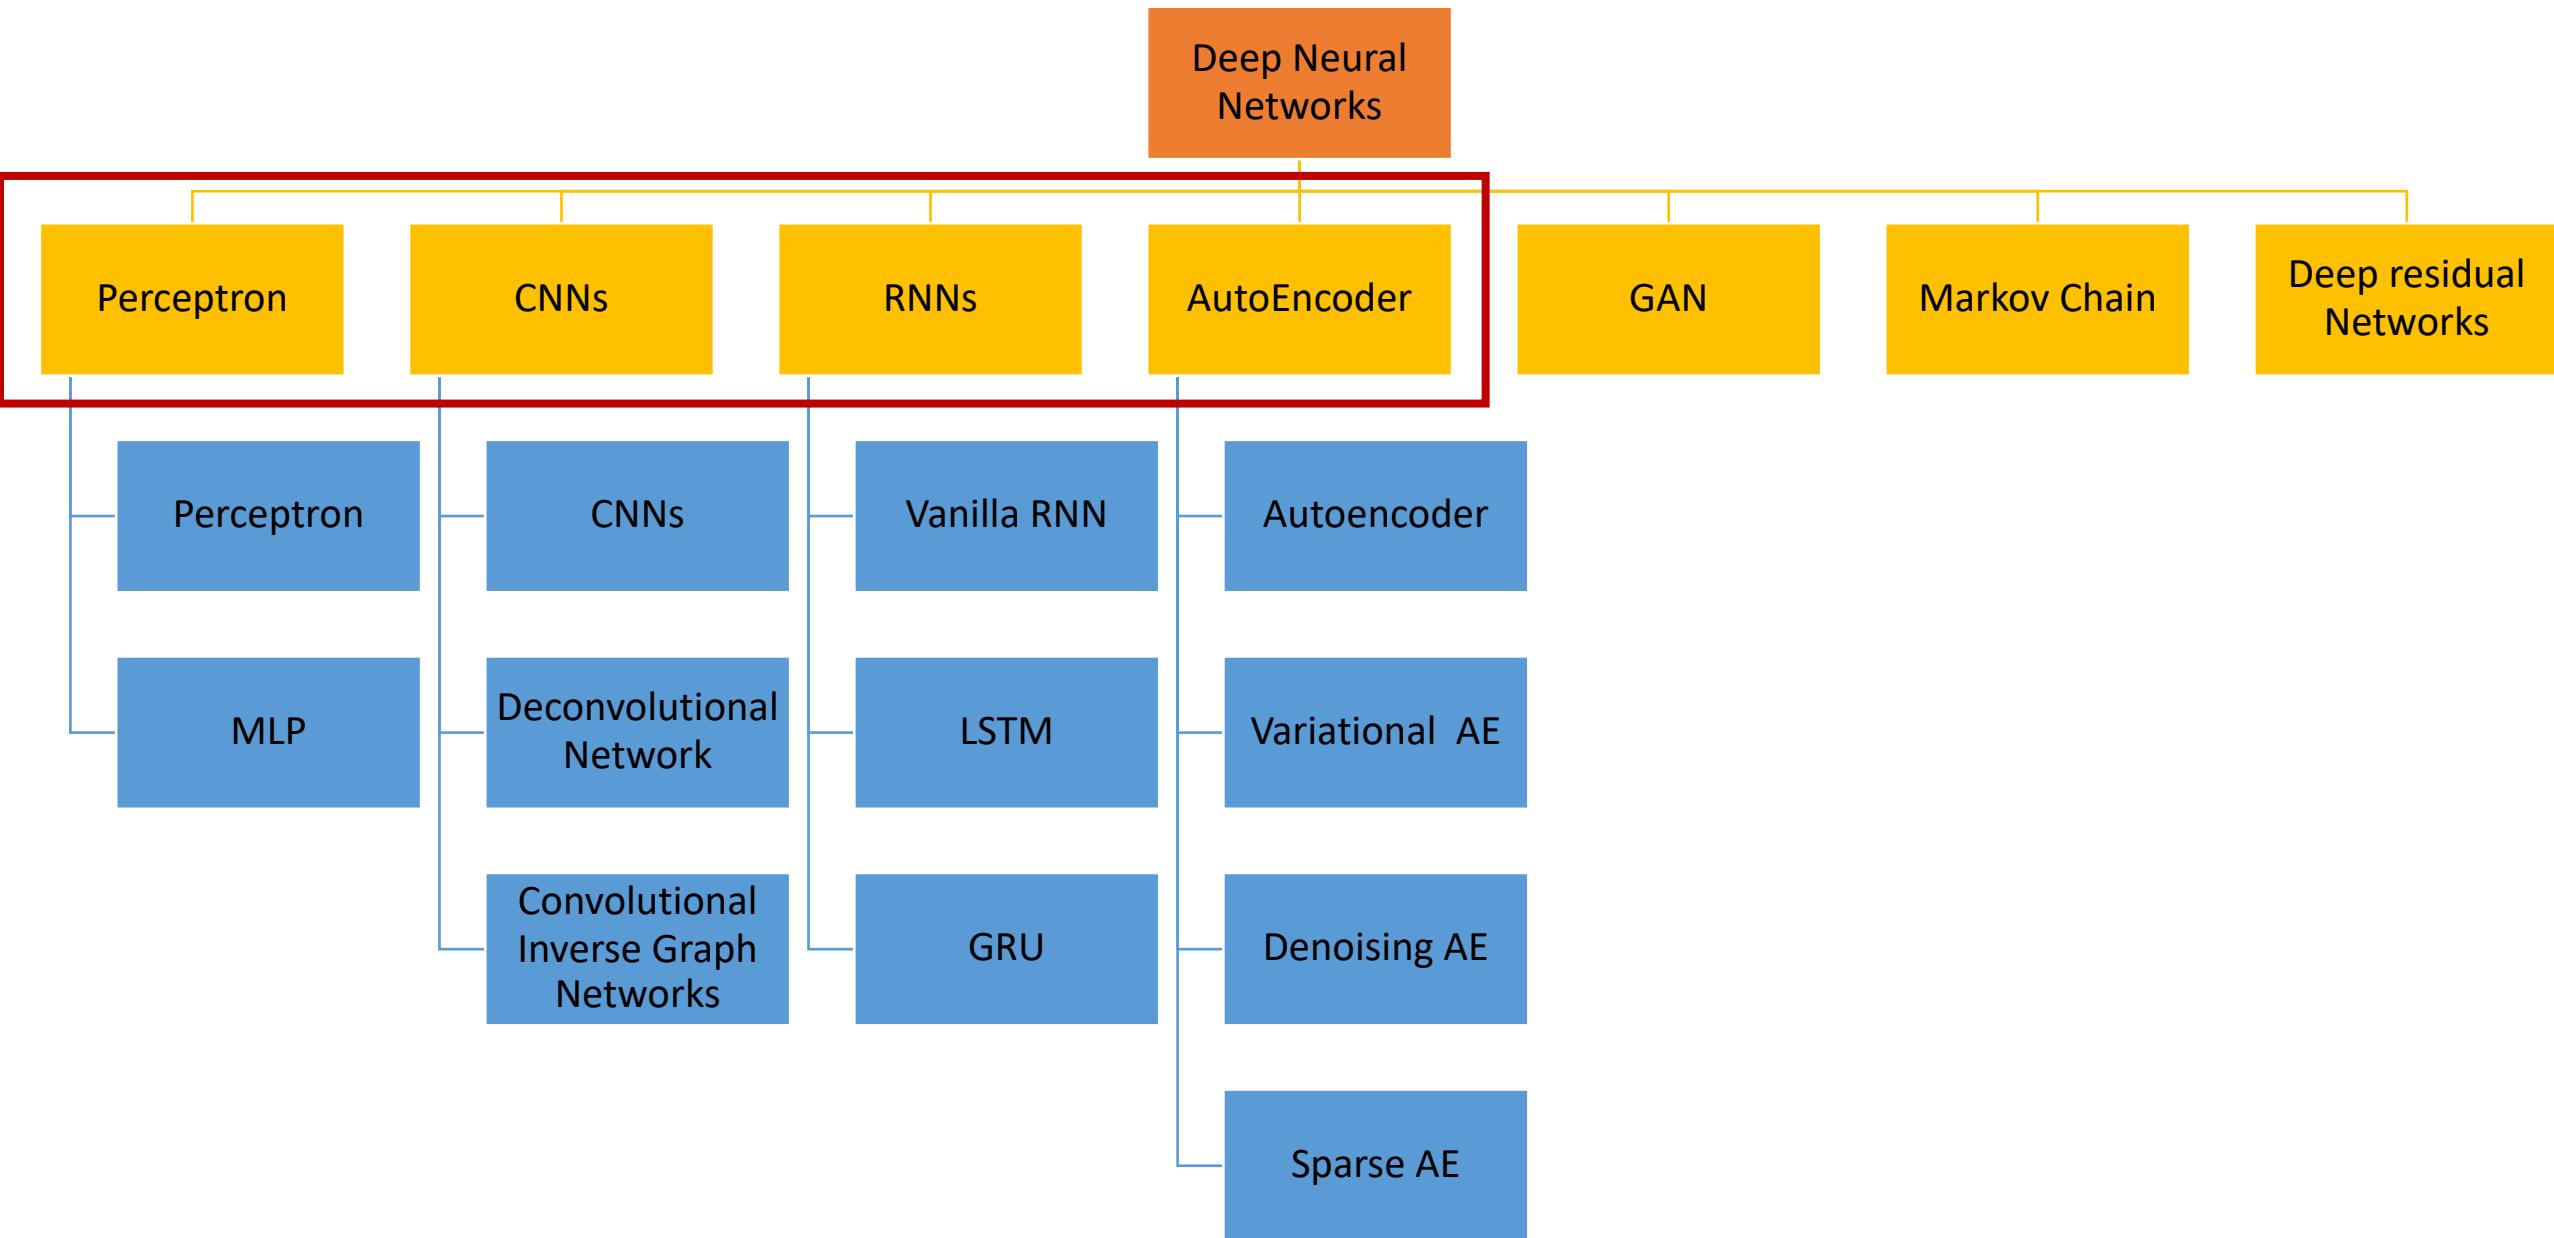
Summer Bootcamp 2021

Applied Machine Learning Deep Neural Networks: **RECAP**

Ngan Le
thile@uark.edu

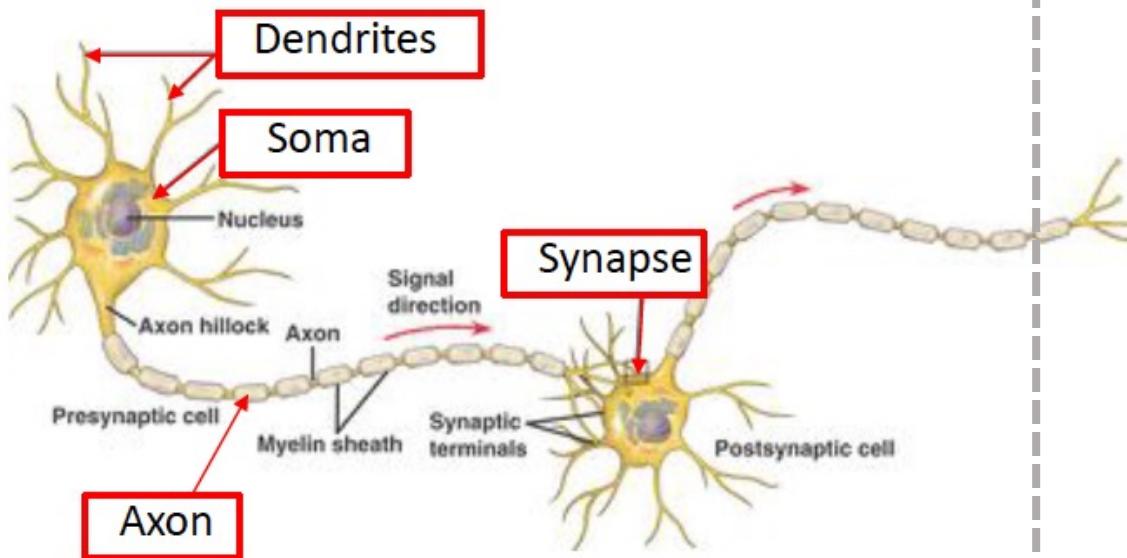




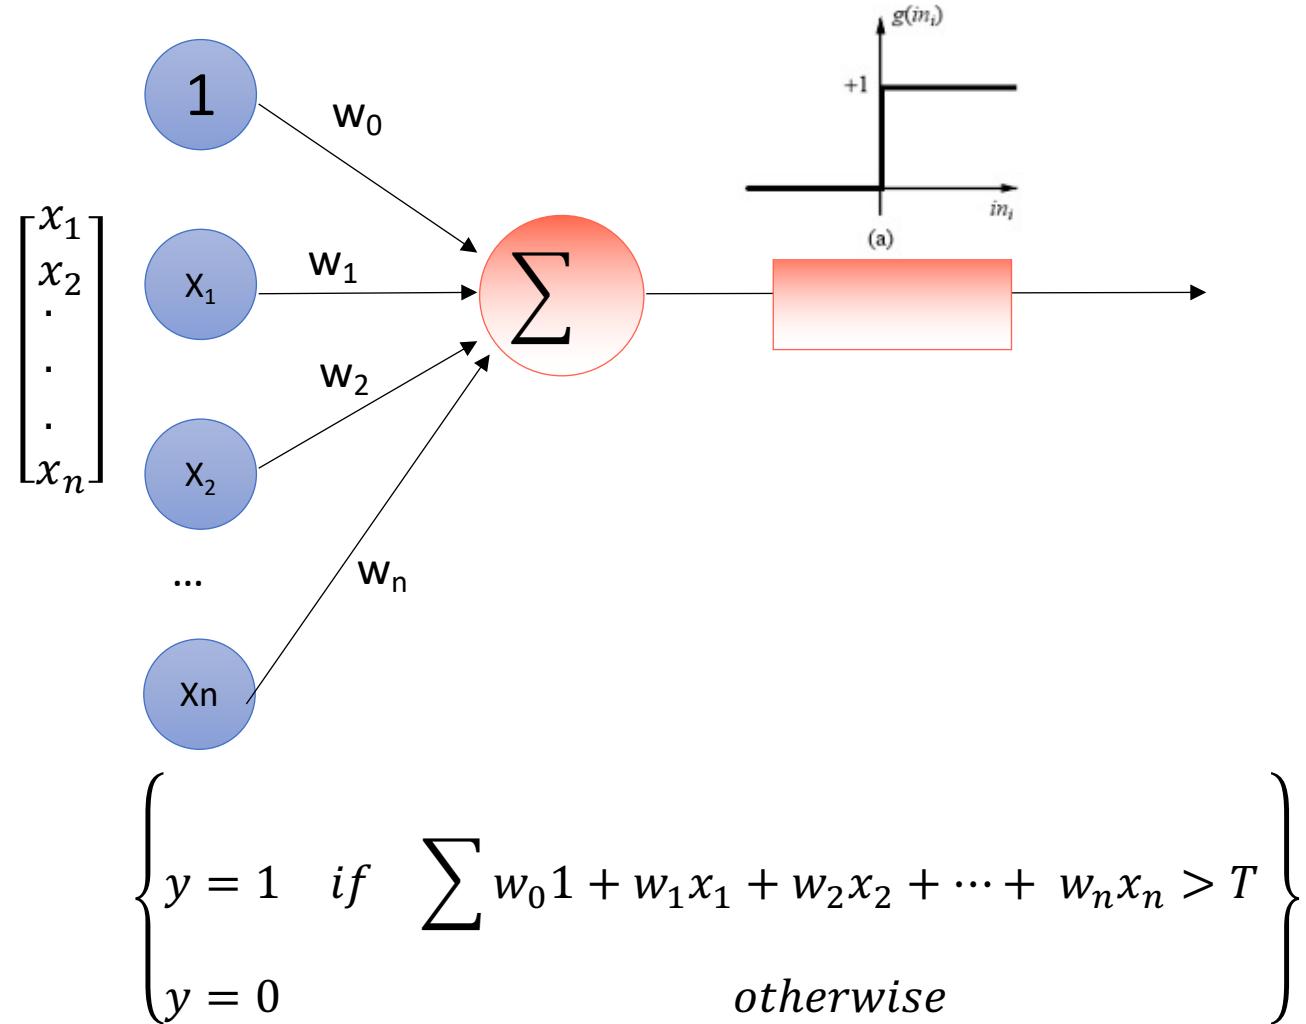


Neural Network

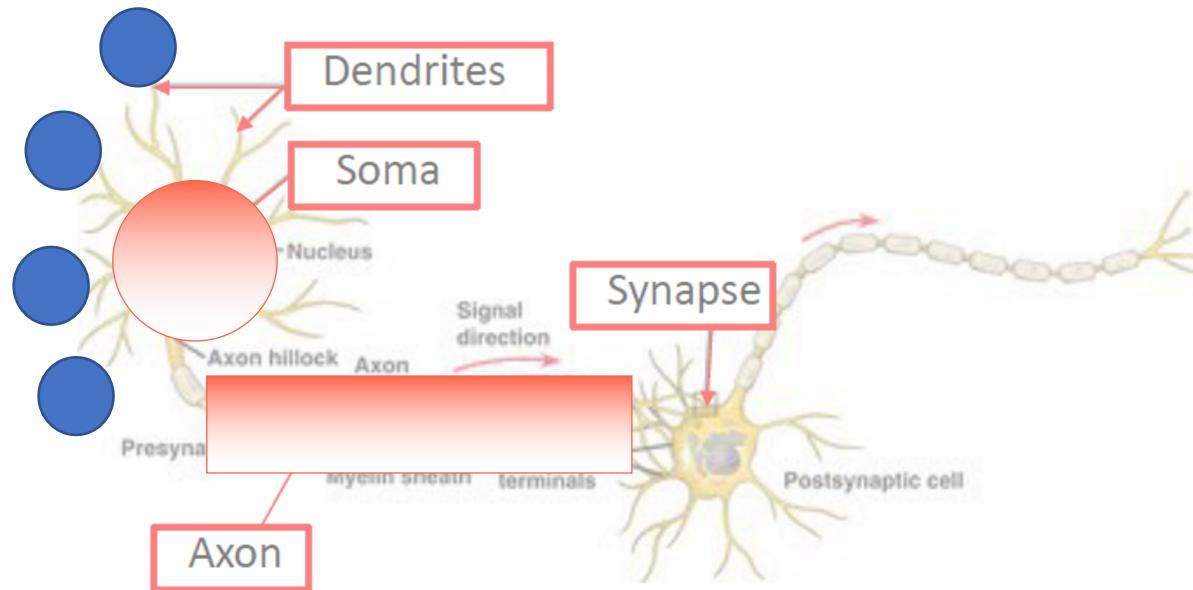
Perceptron



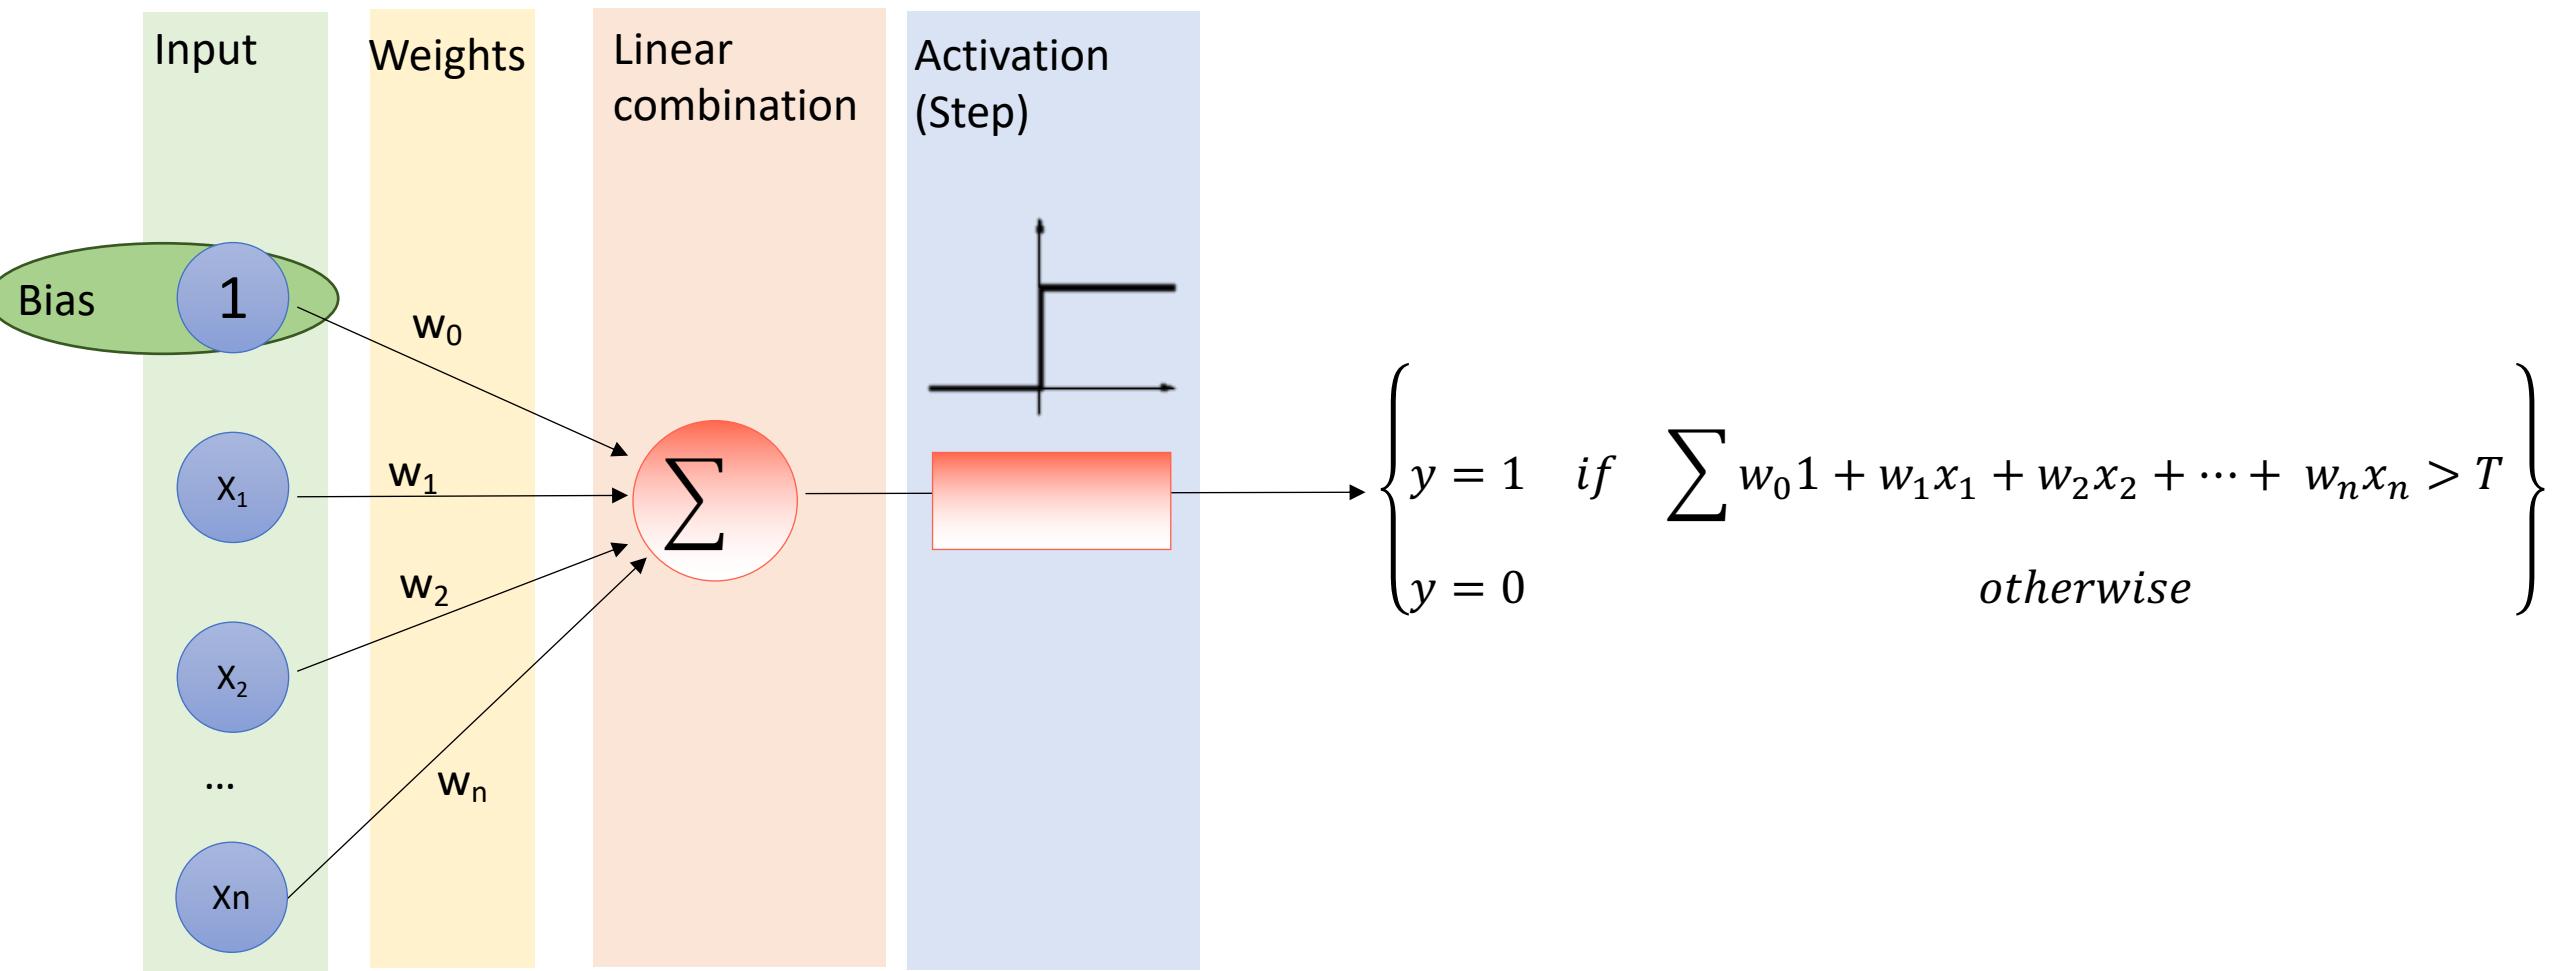
Cells that fire together, wire together



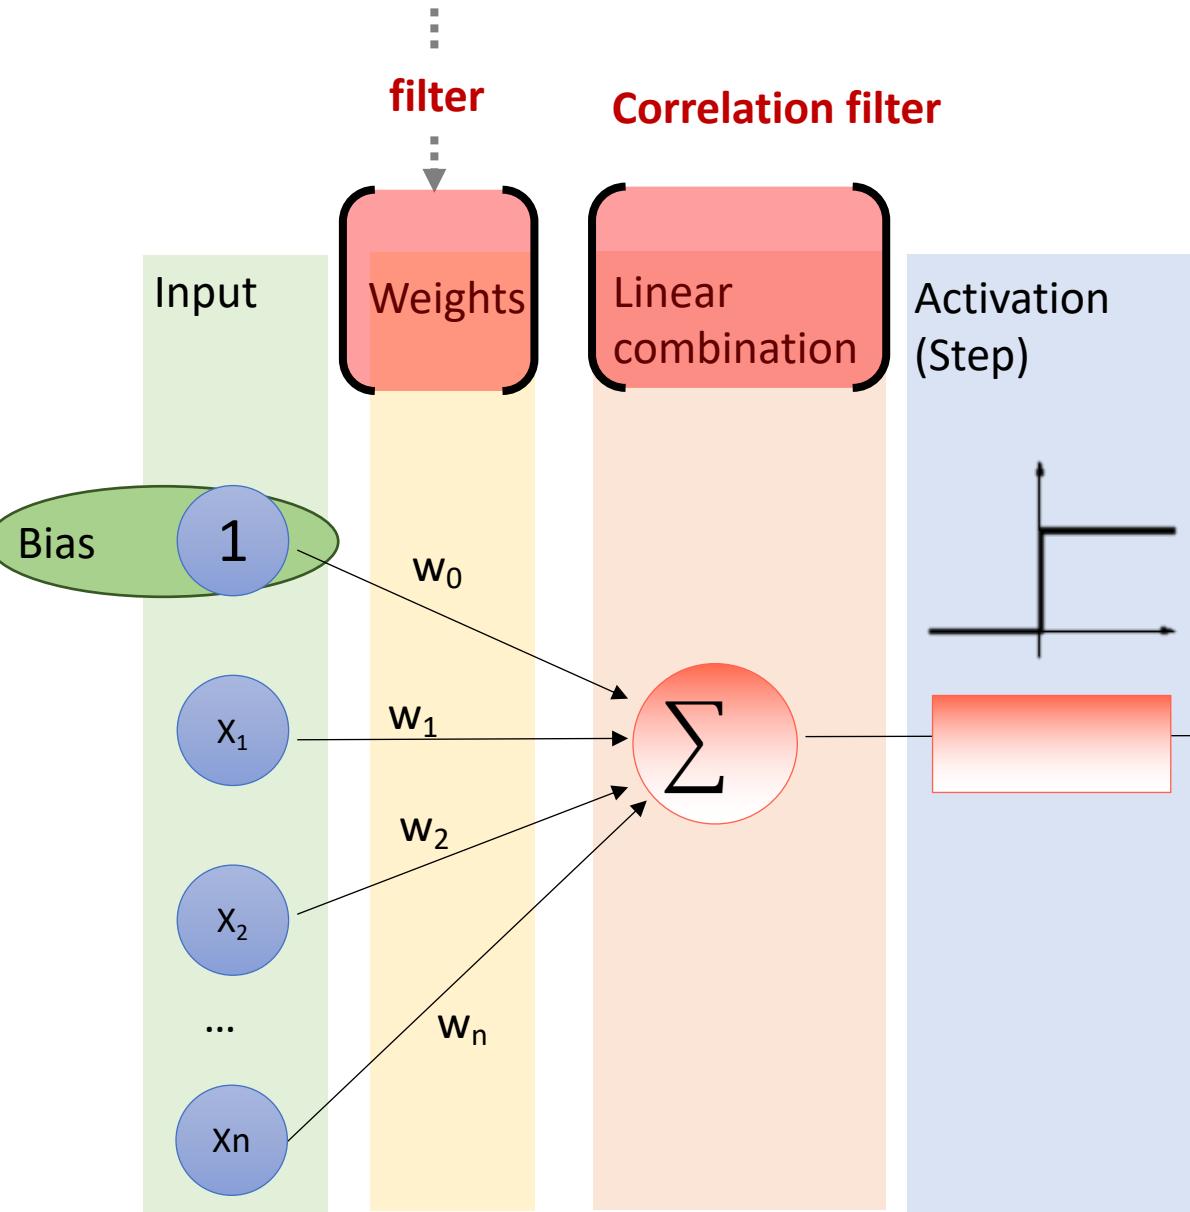
Perceptron



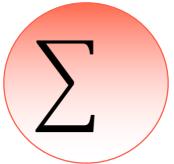
Cells that fire together, wire together



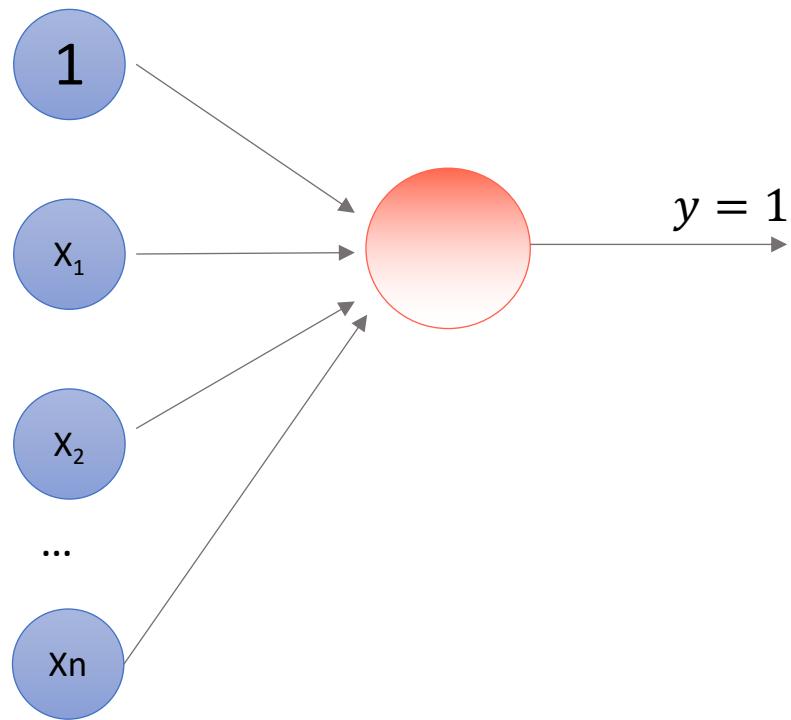
Learnt by training



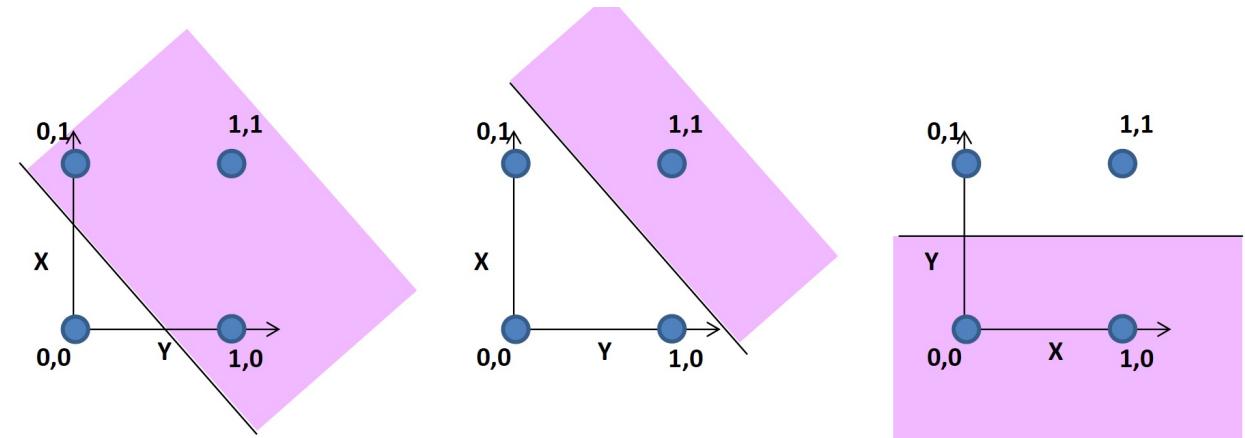
$$\left\{ \begin{array}{ll} y = 1 & \text{if } w_0 1 + w_1 x_1 + w_2 x_2 + \dots + w_n x_n > T \\ y = 0 & \text{otherwise} \end{array} \right.$$

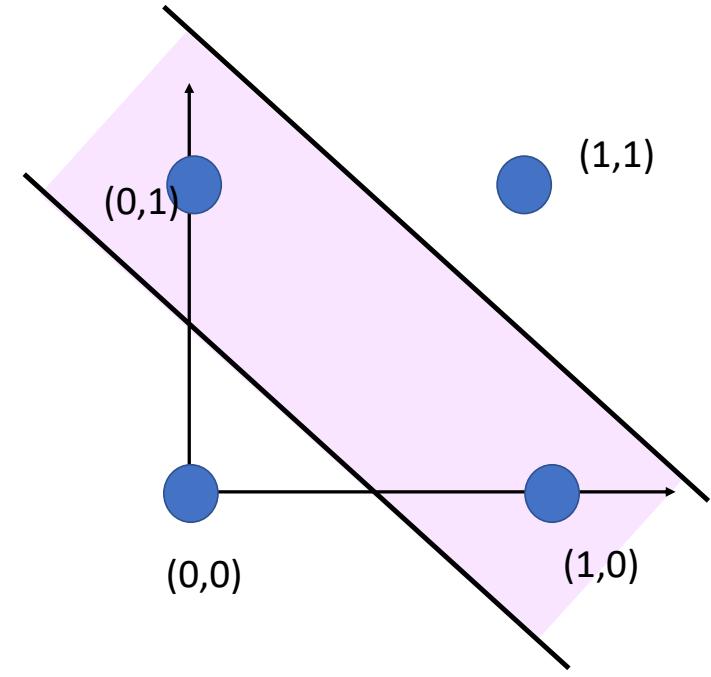
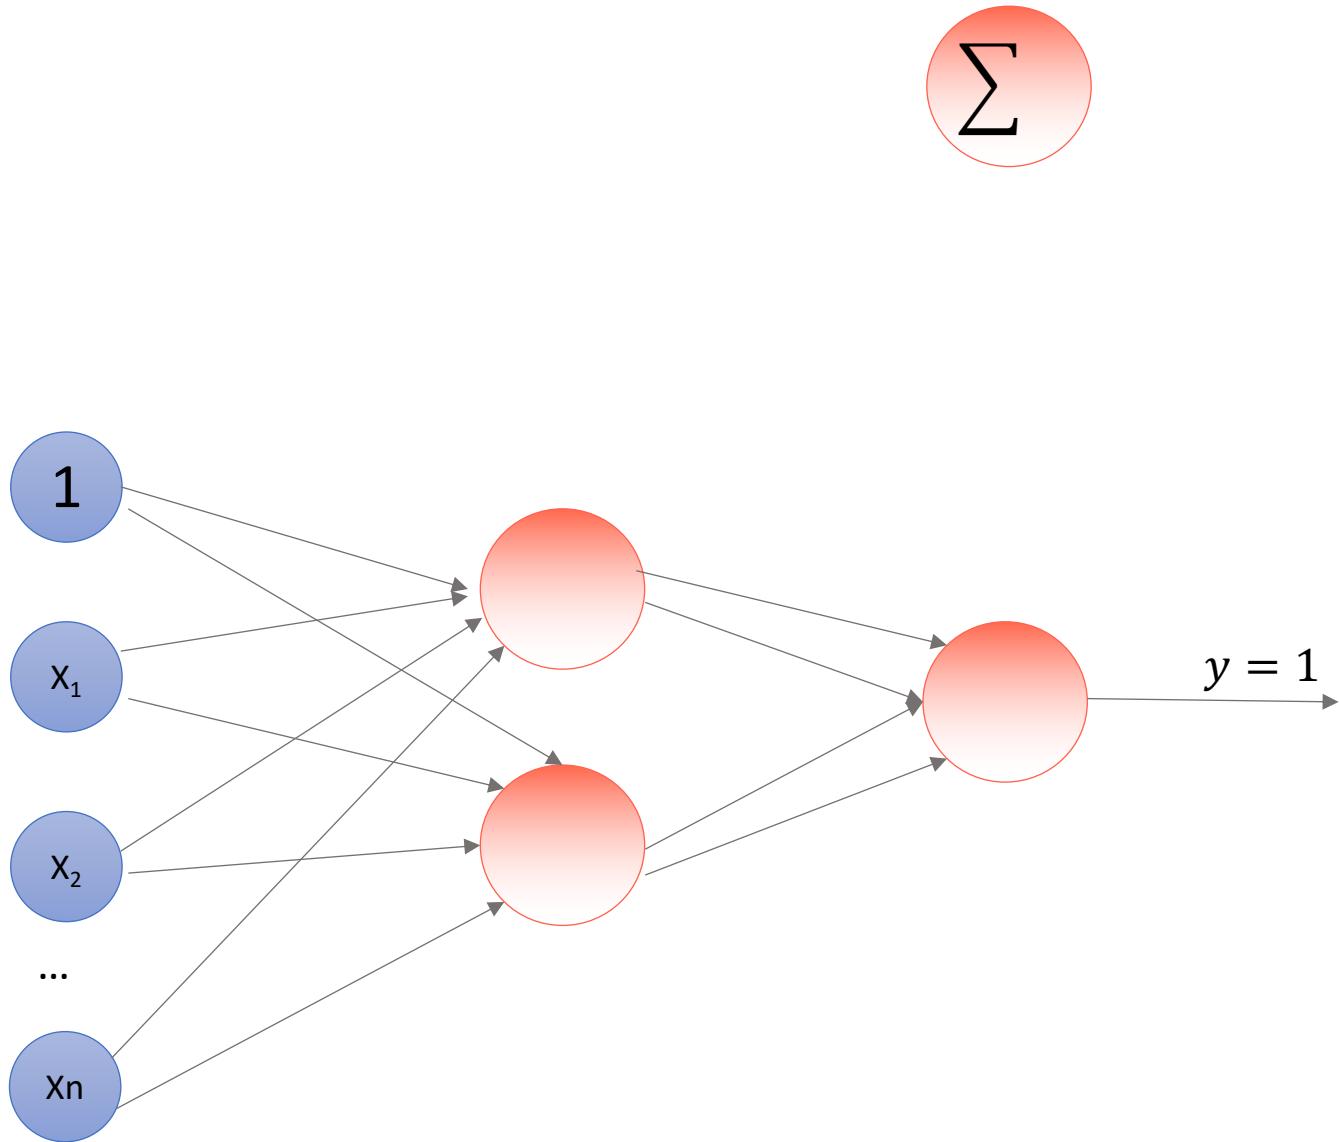


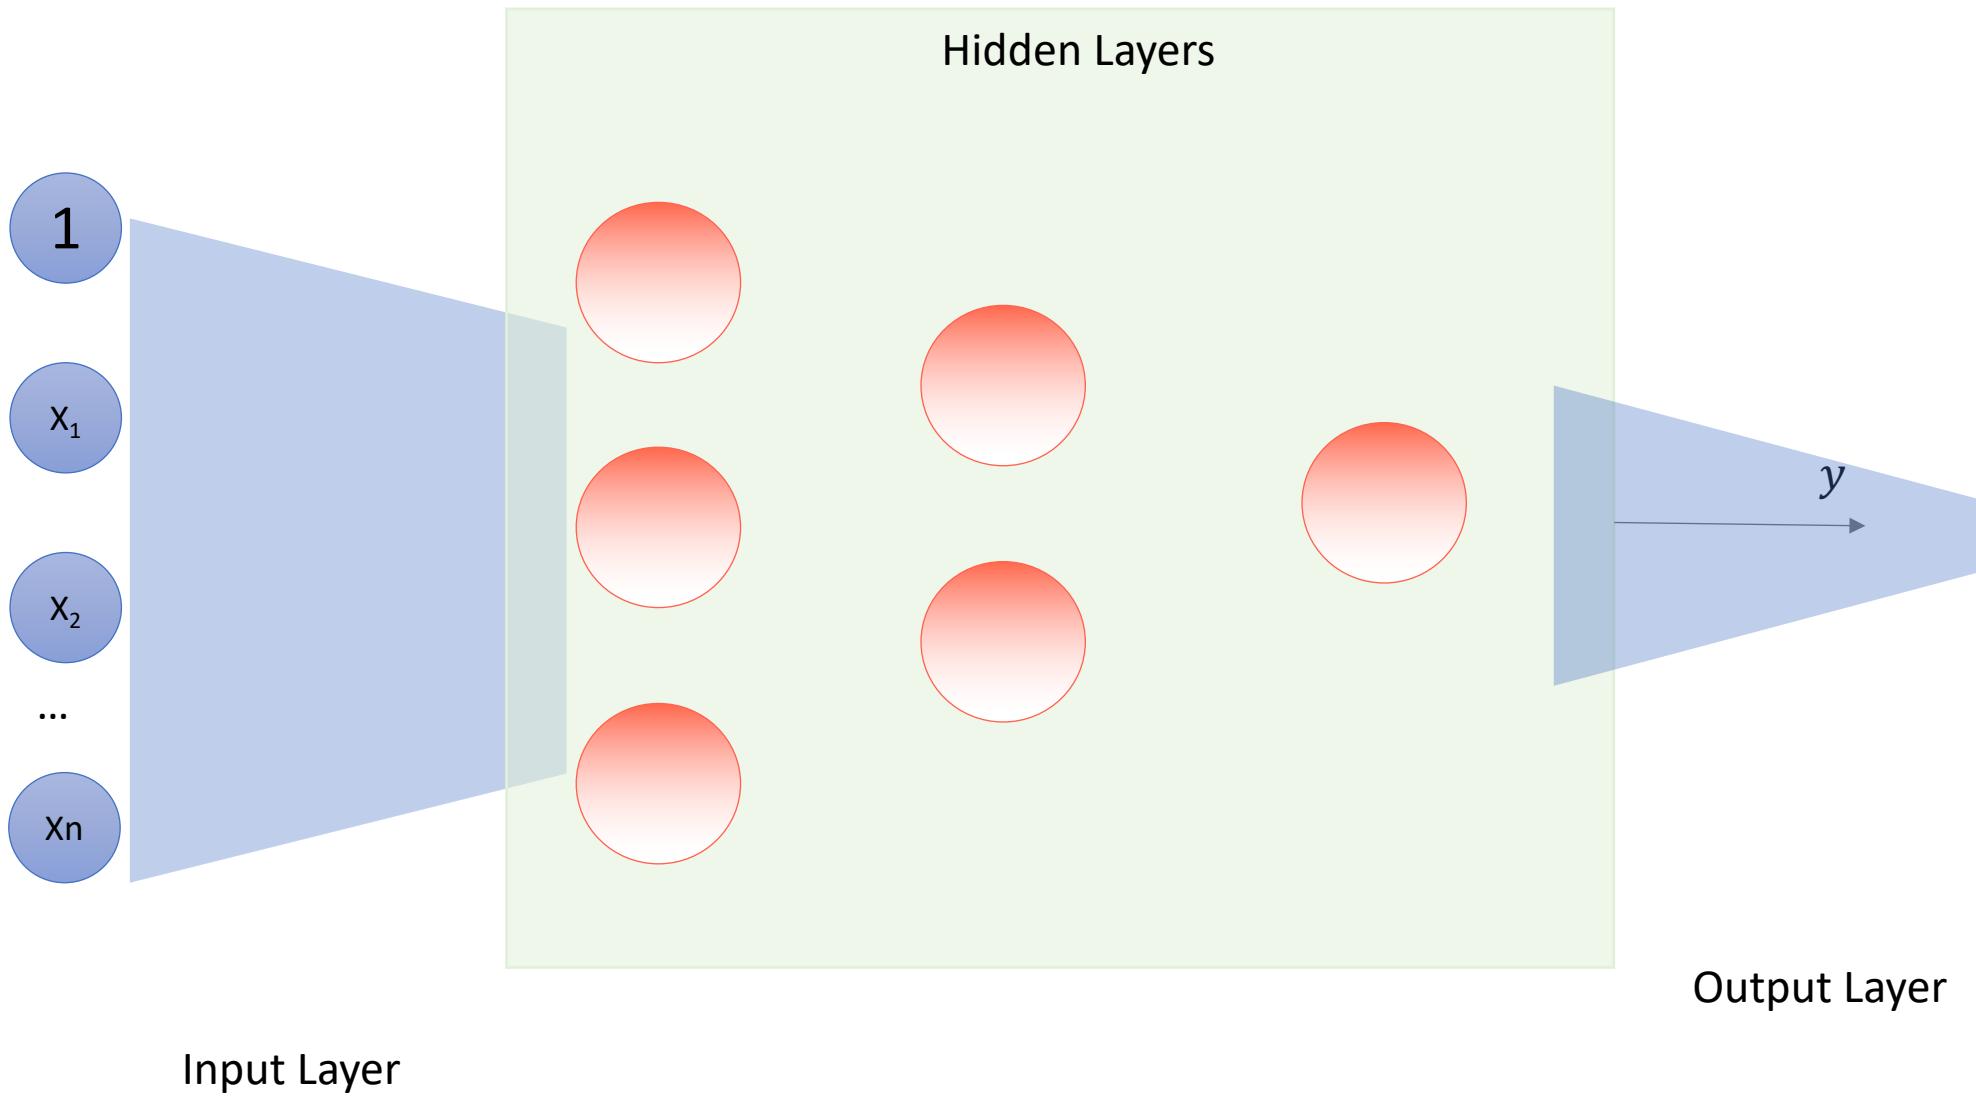
- How many perceptrons?
- How is perceptron's arrange?



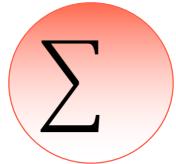
\sum



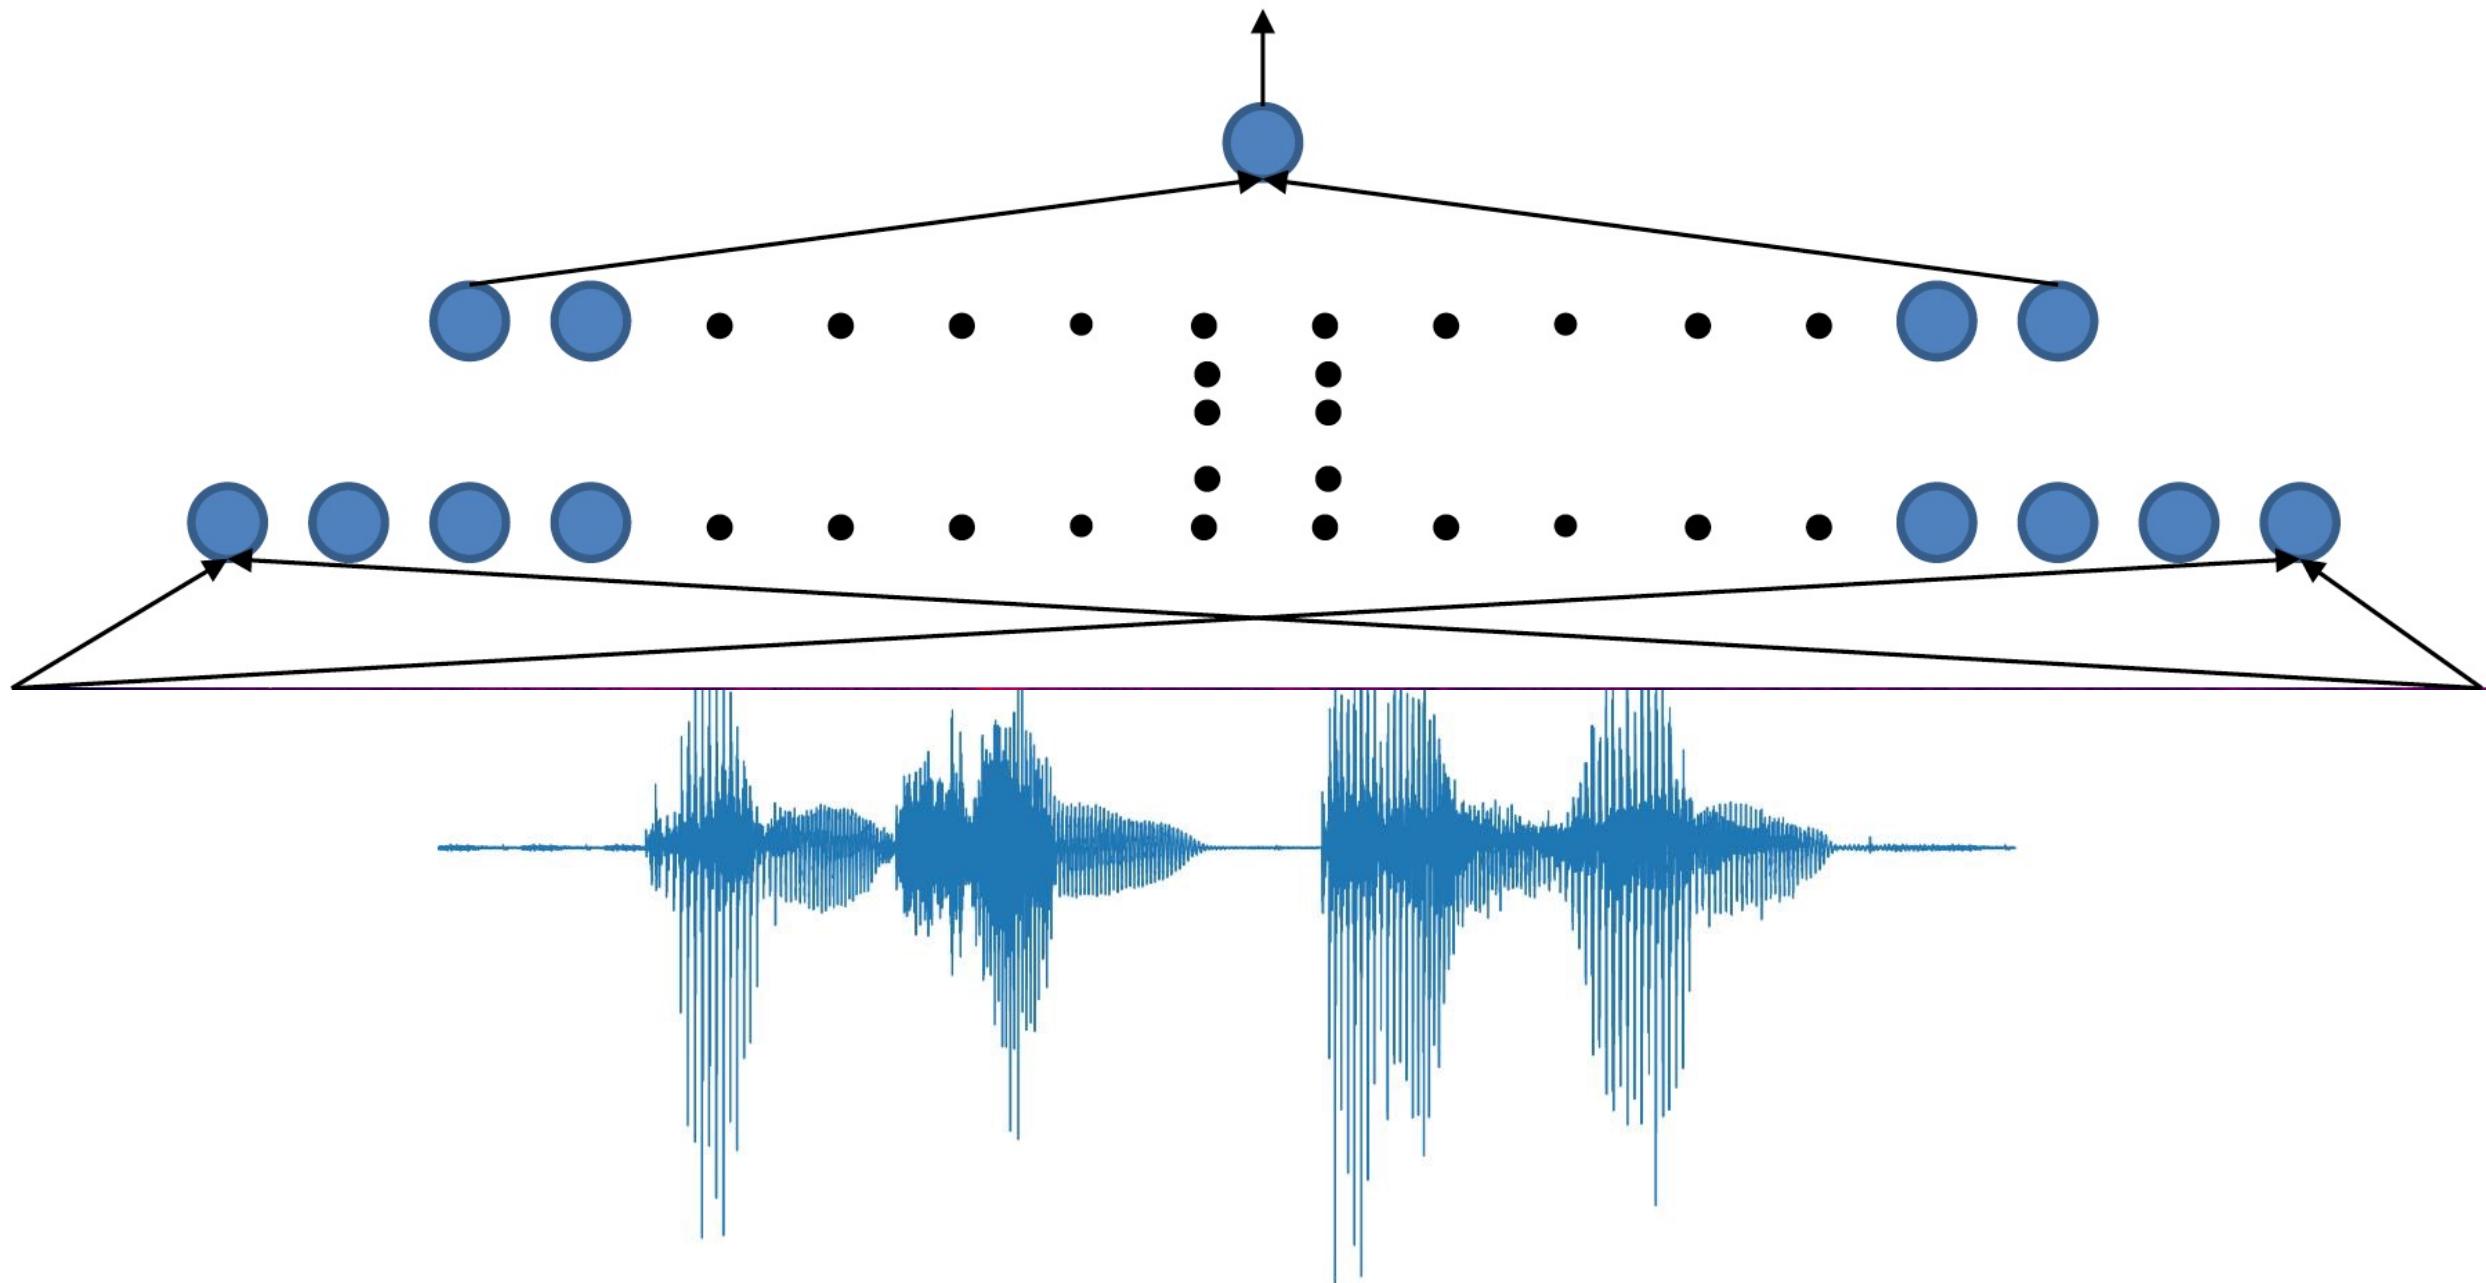




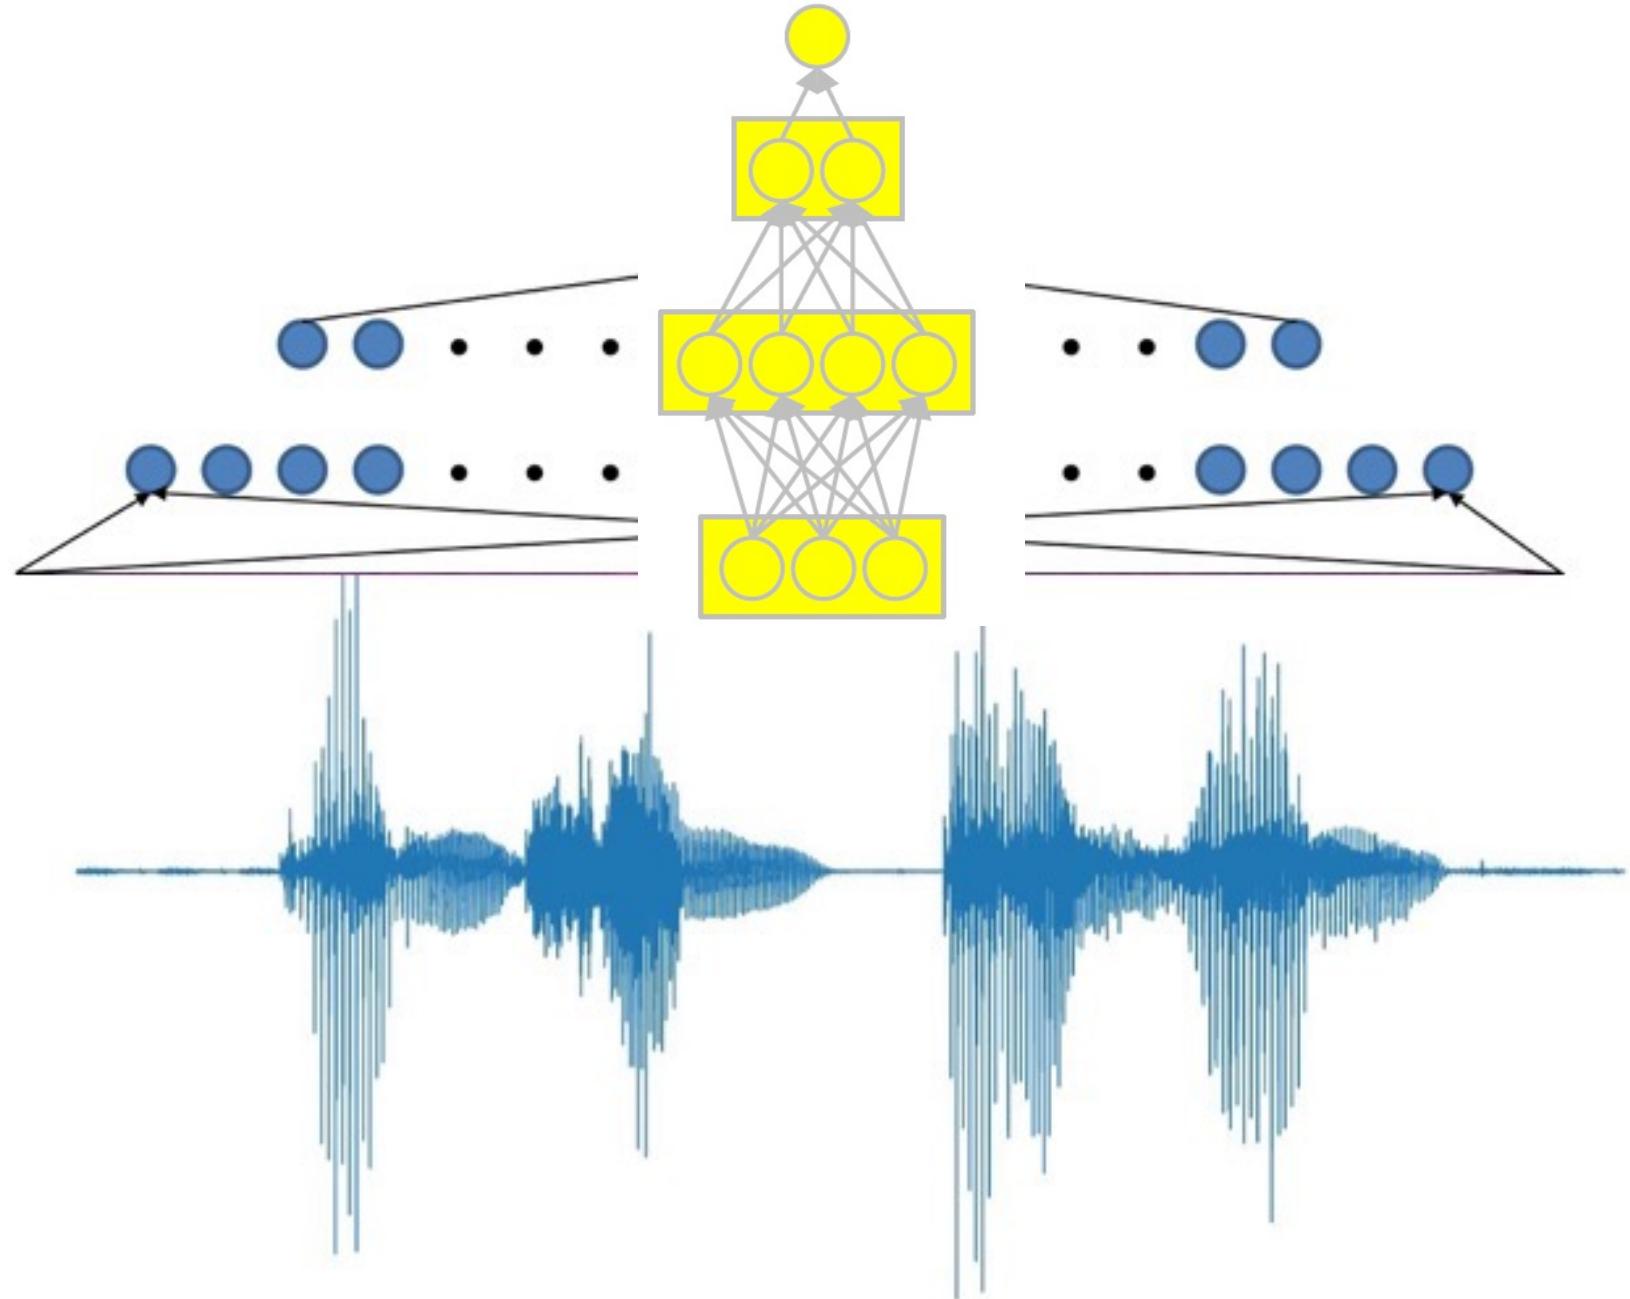
Multi Layer Perceptron (MLP)



- Which information that one perceptron cover?

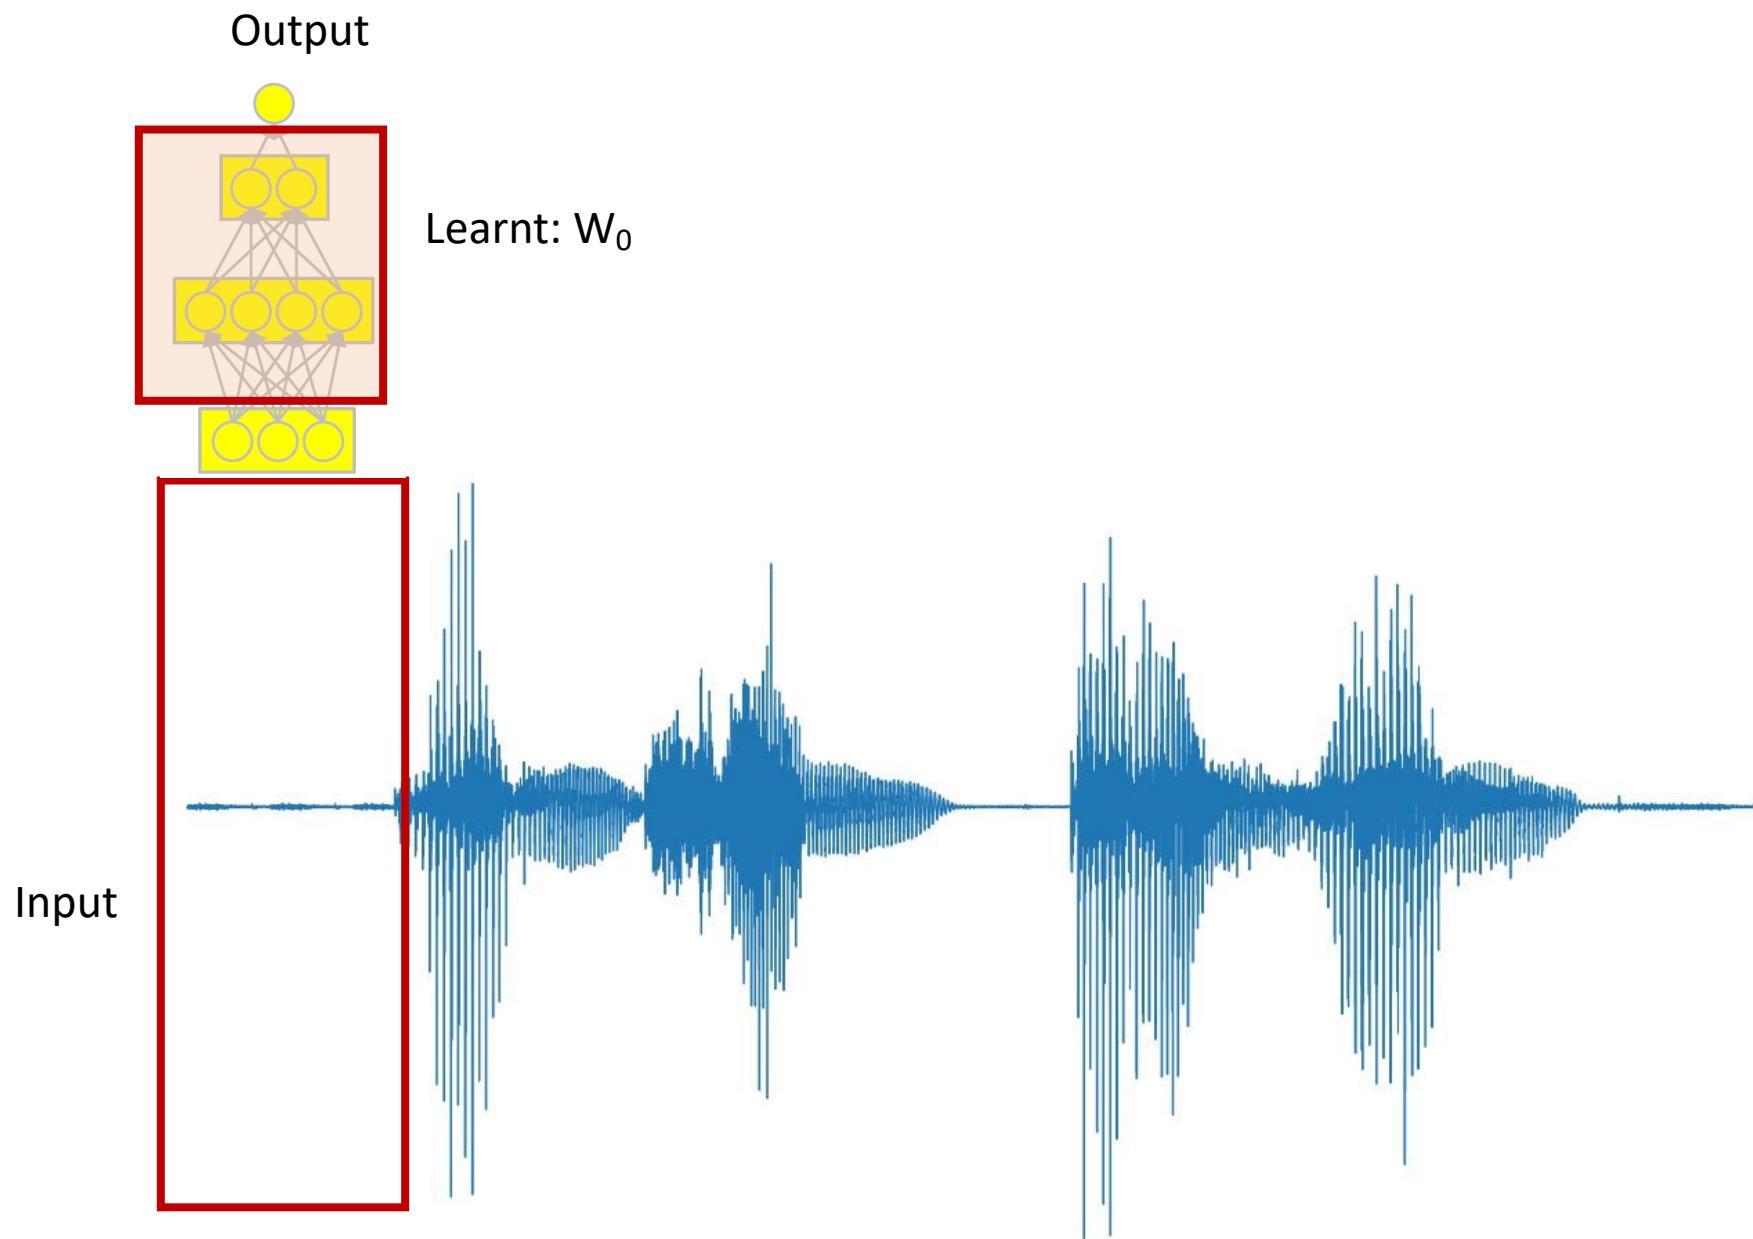


Solution for?

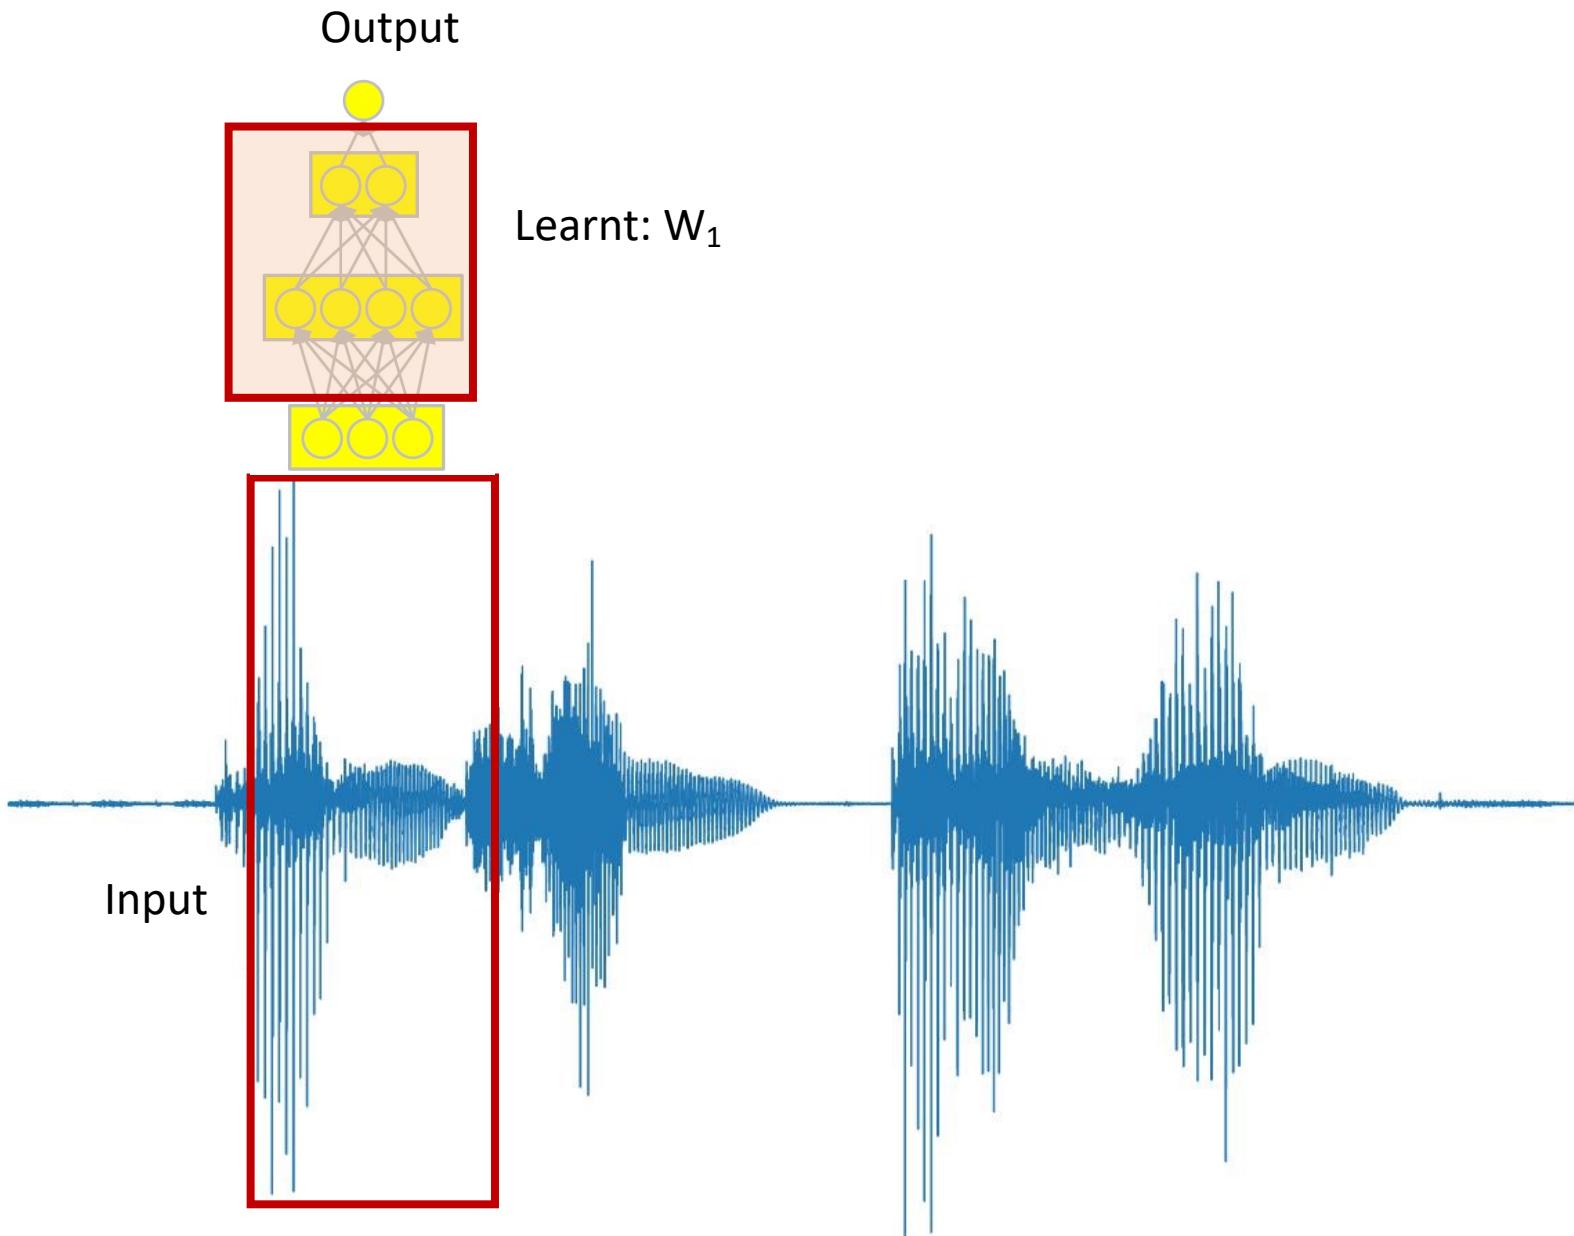


Limitations?

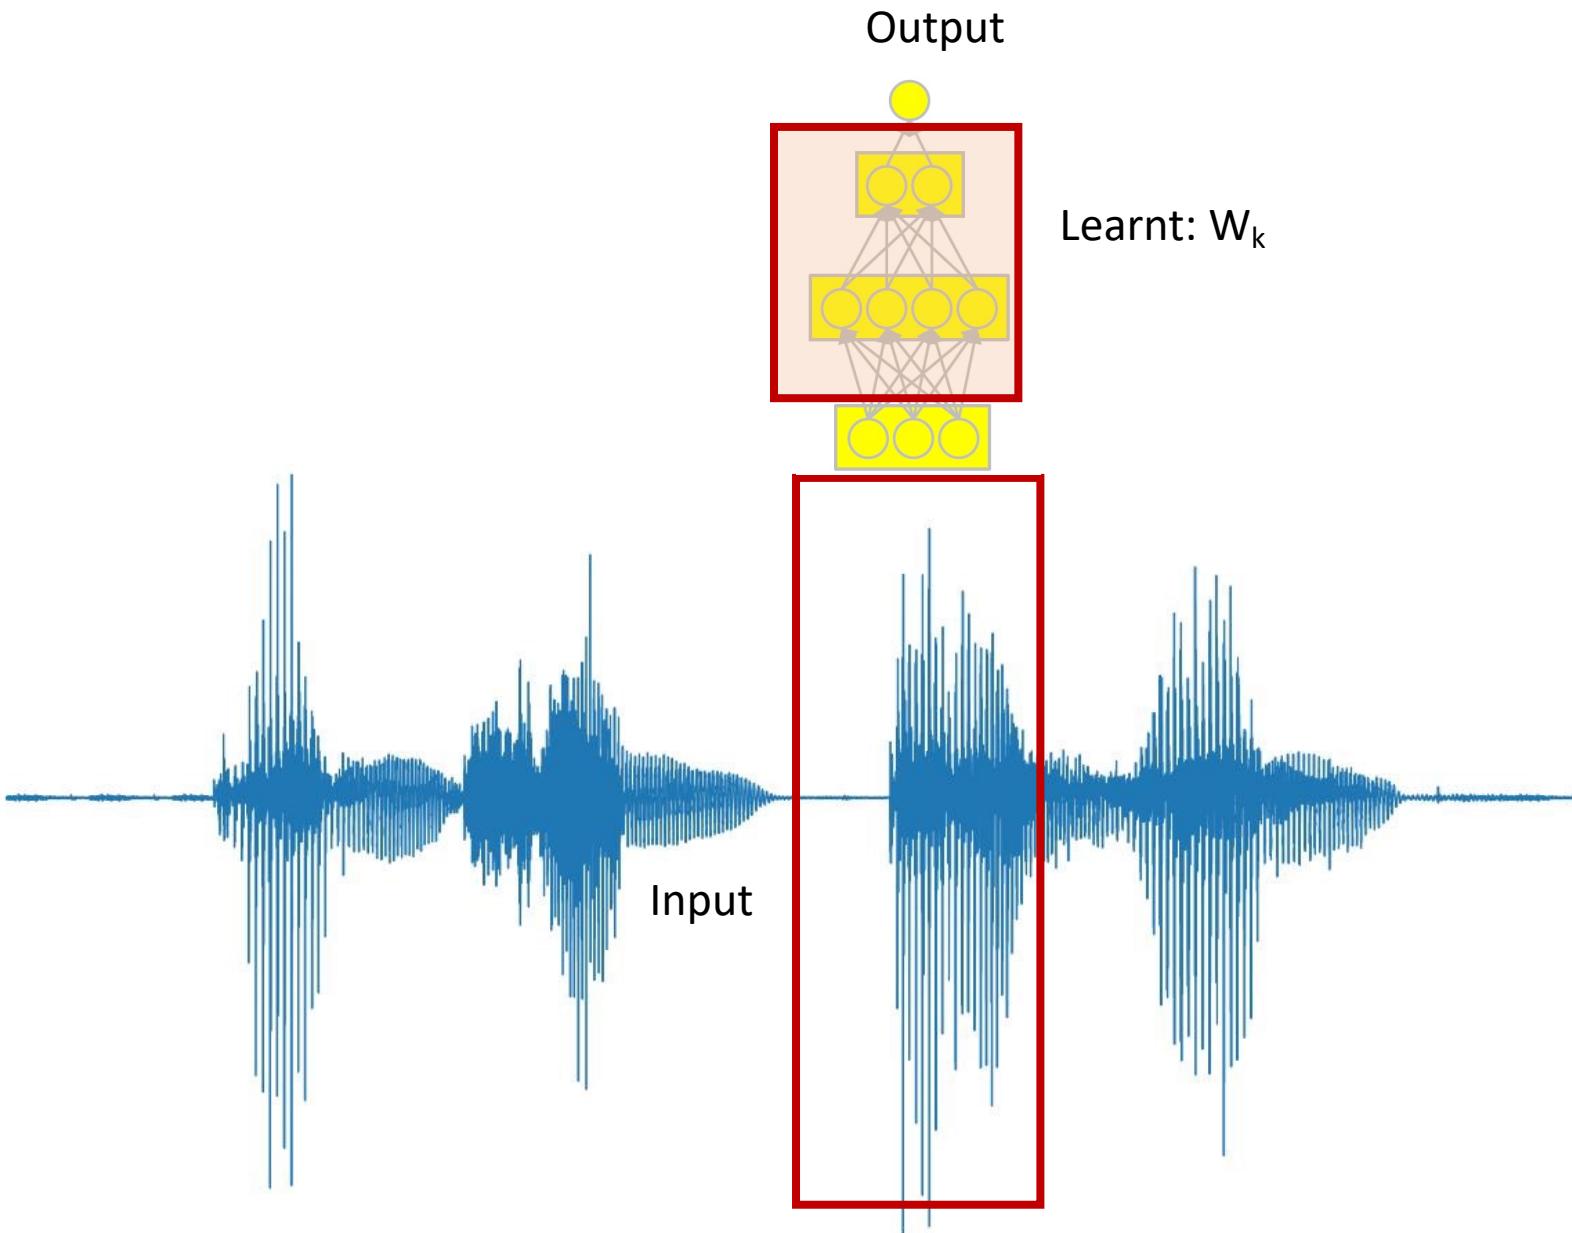
Search Pattern: Scanning



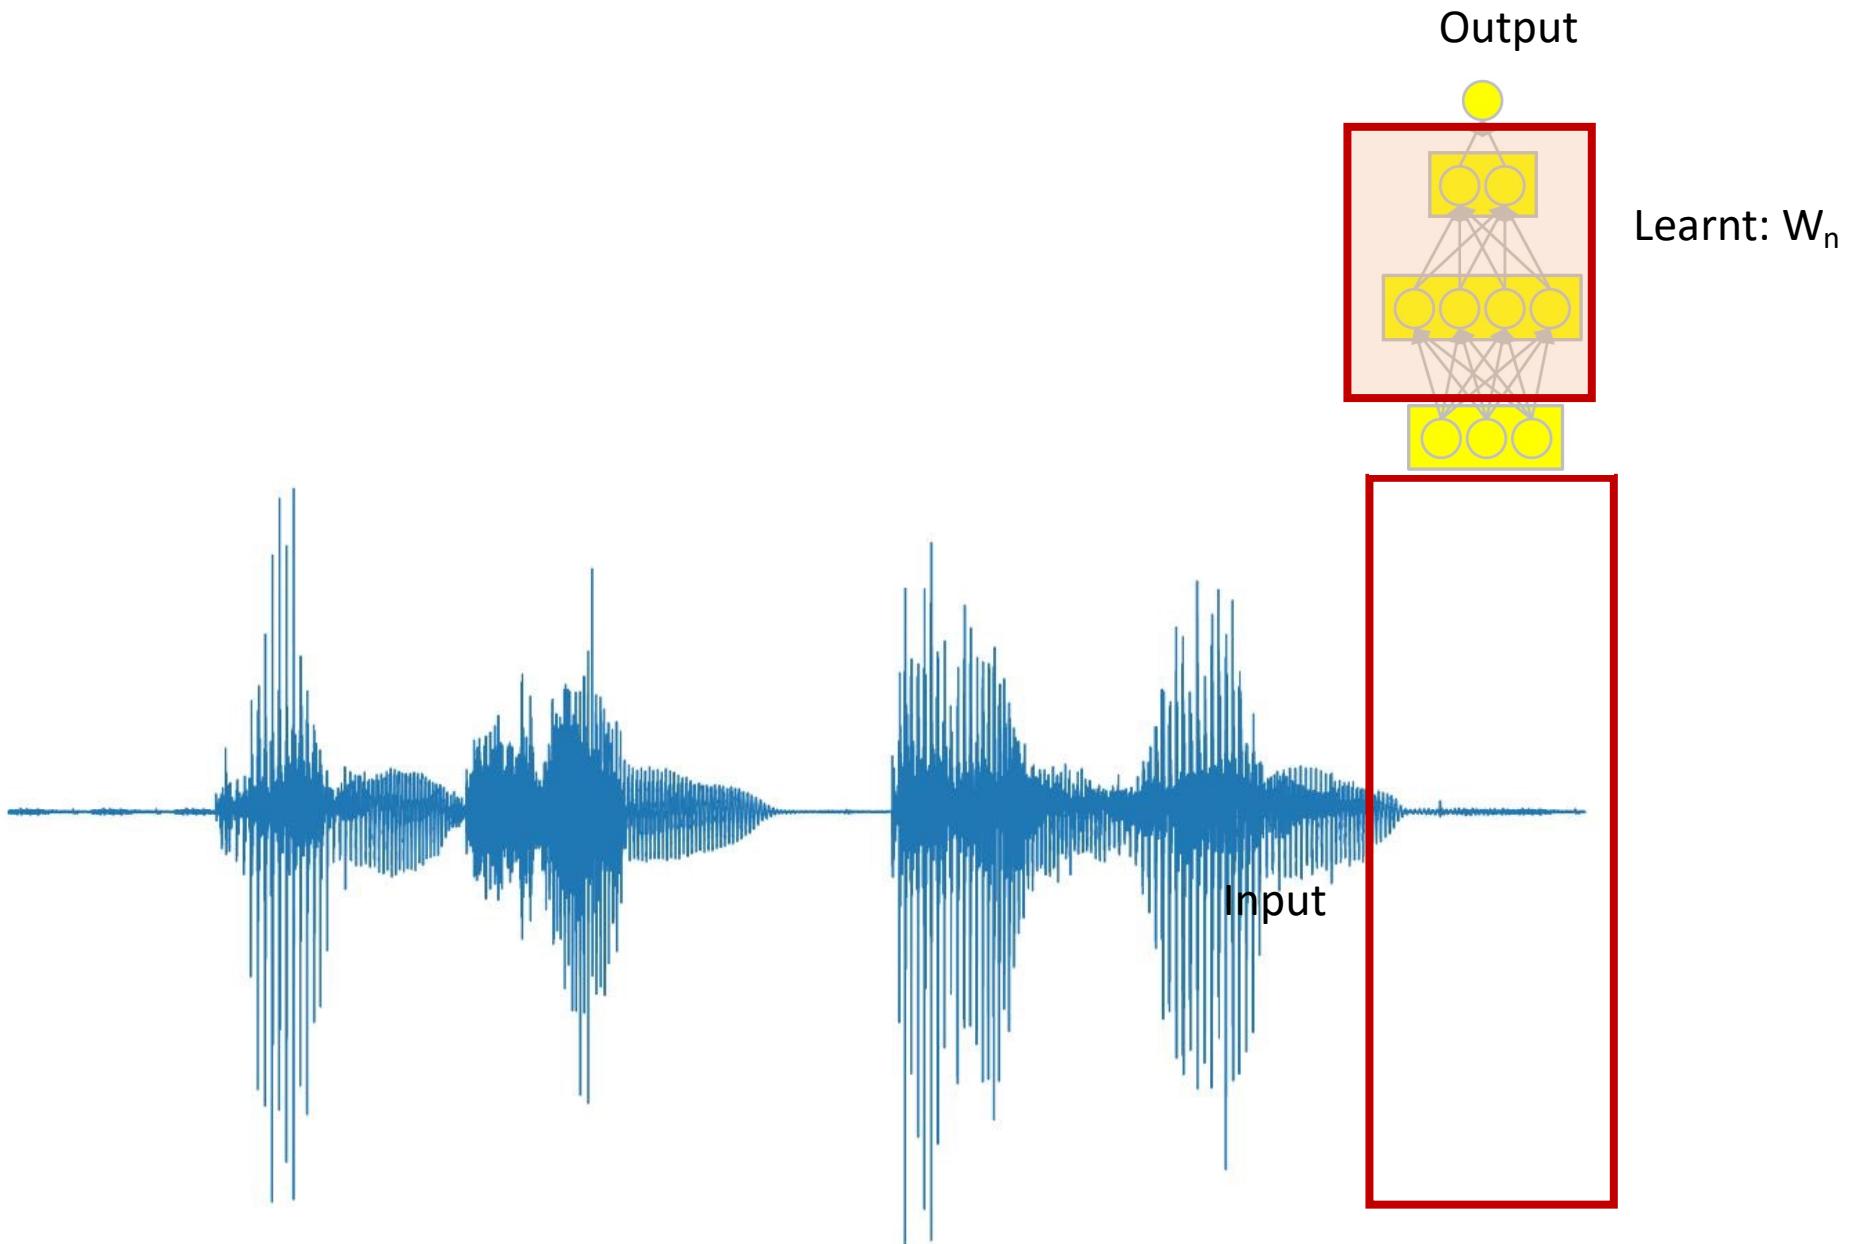
Search Pattern: Scanning



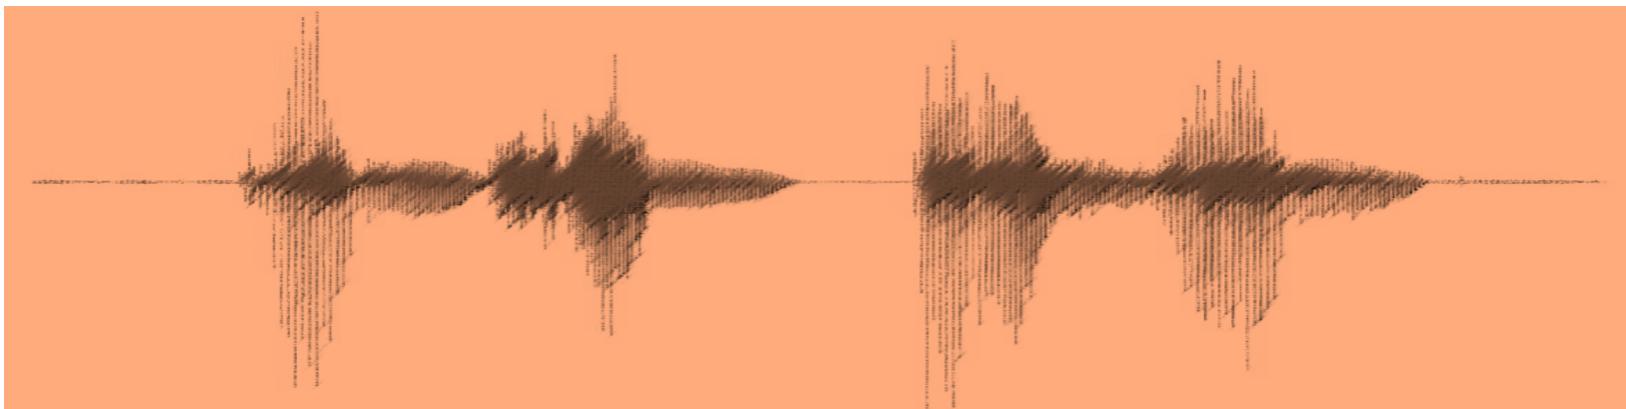
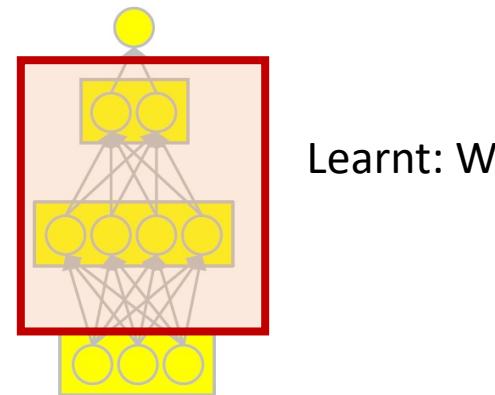
Search Pattern: Scanning

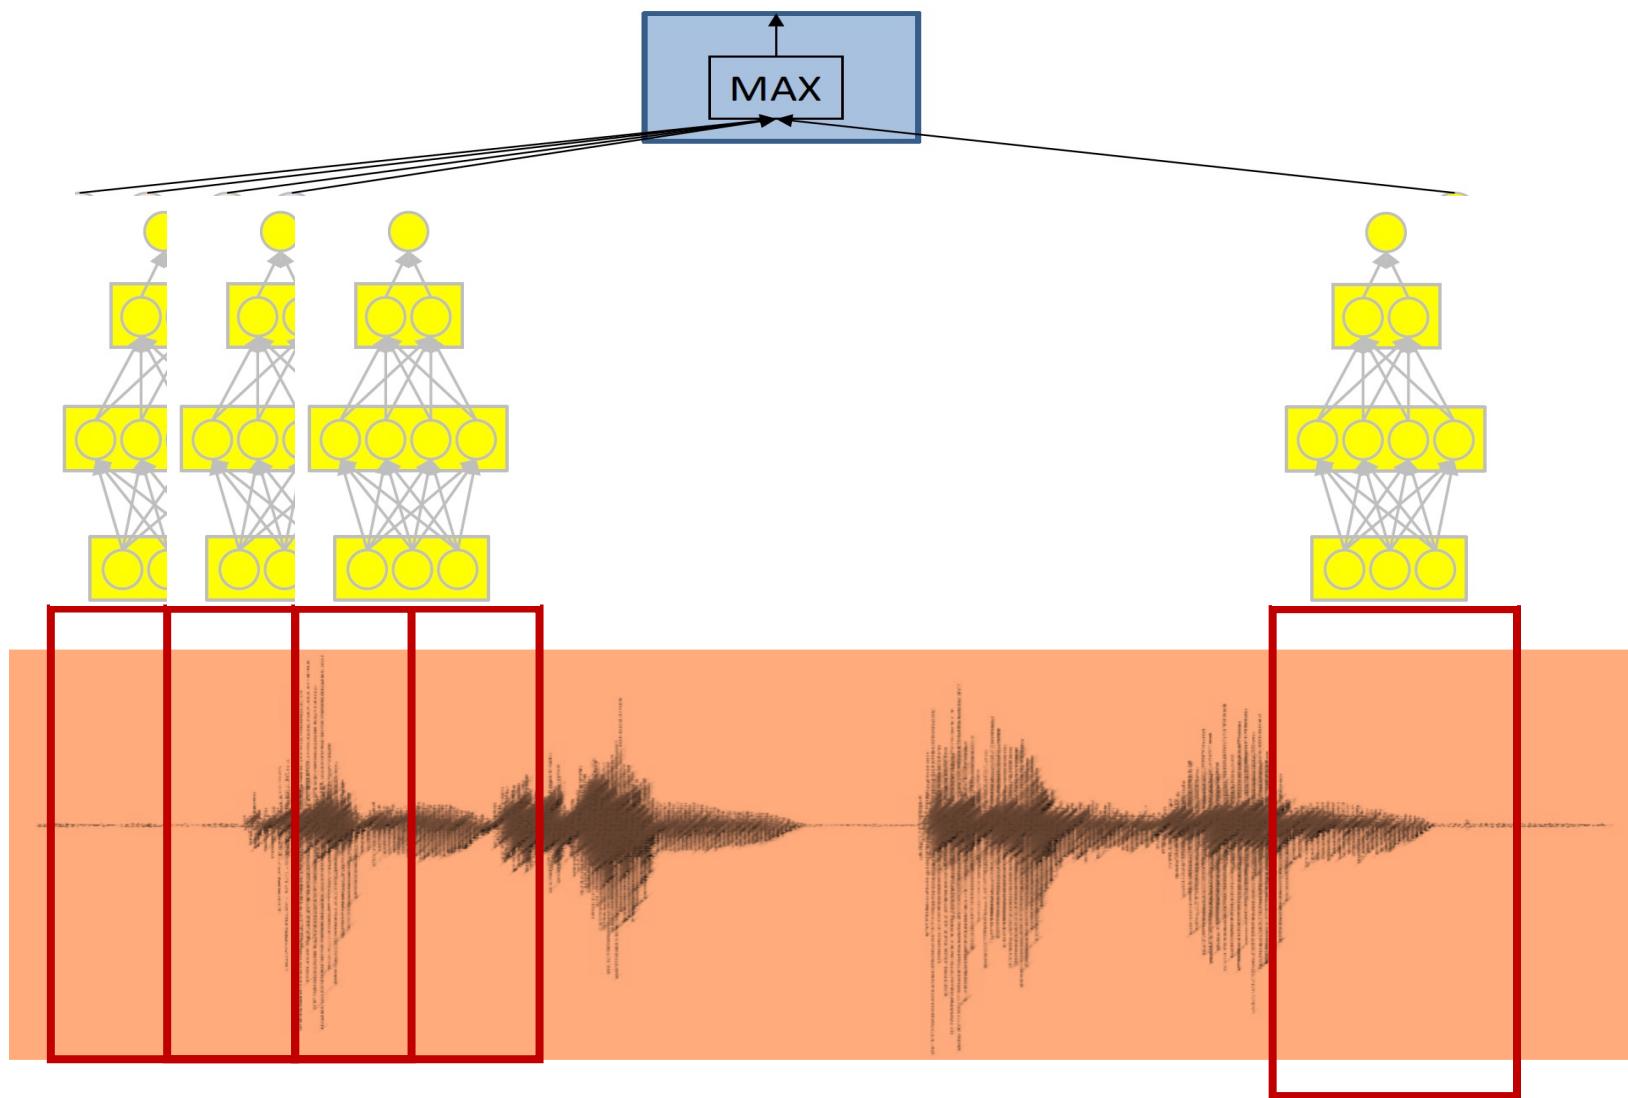


Search Pattern: Scanning

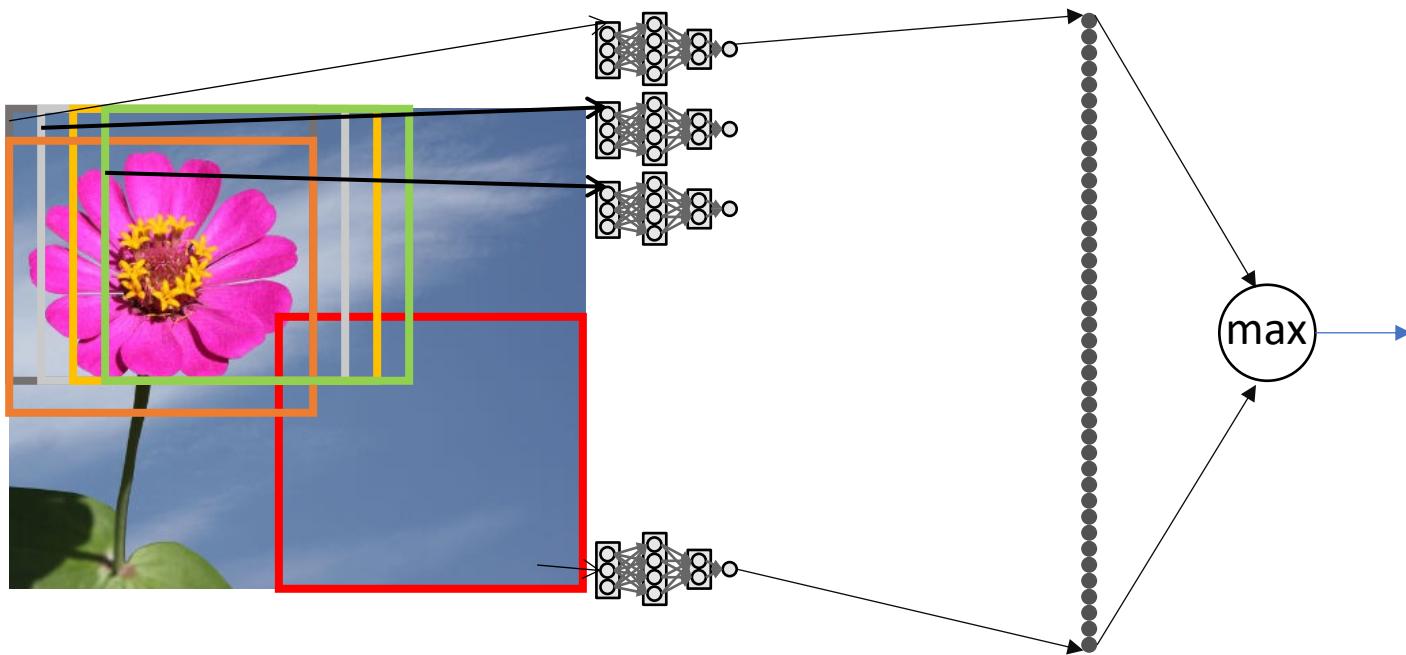


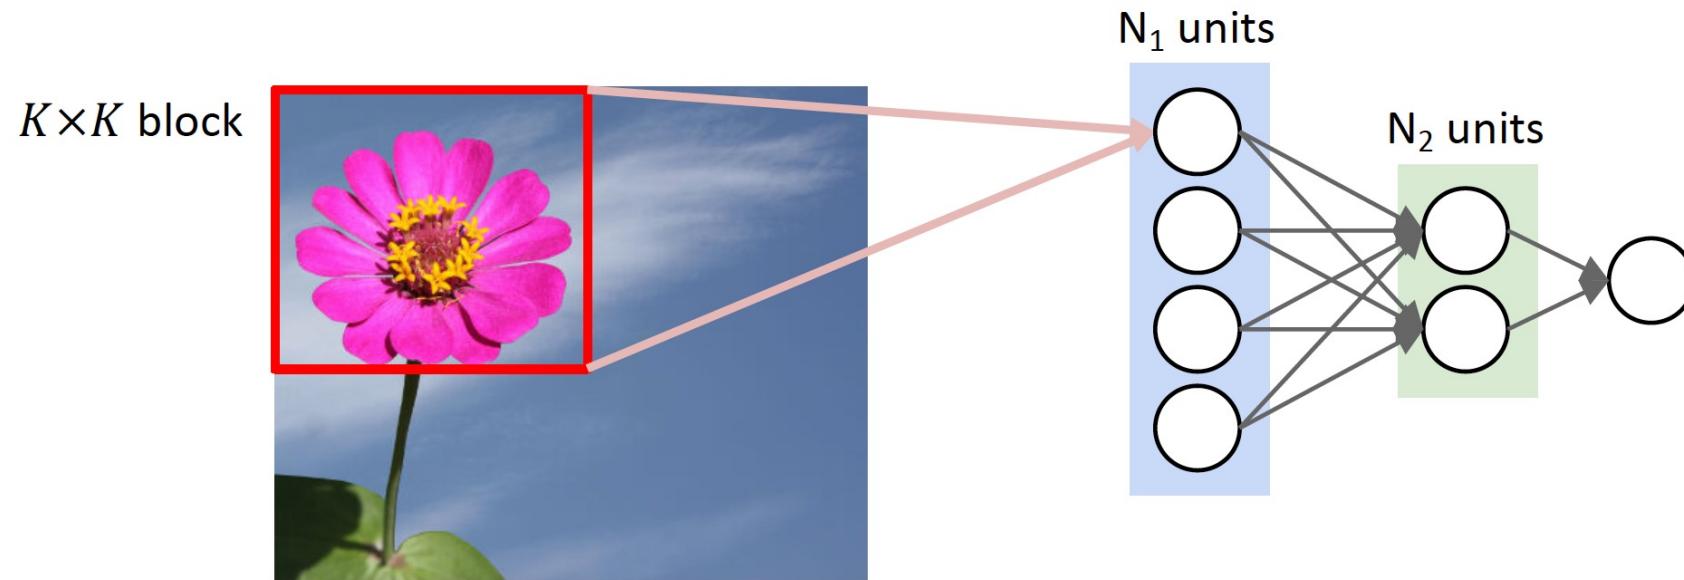
Search Pattern: Scanning

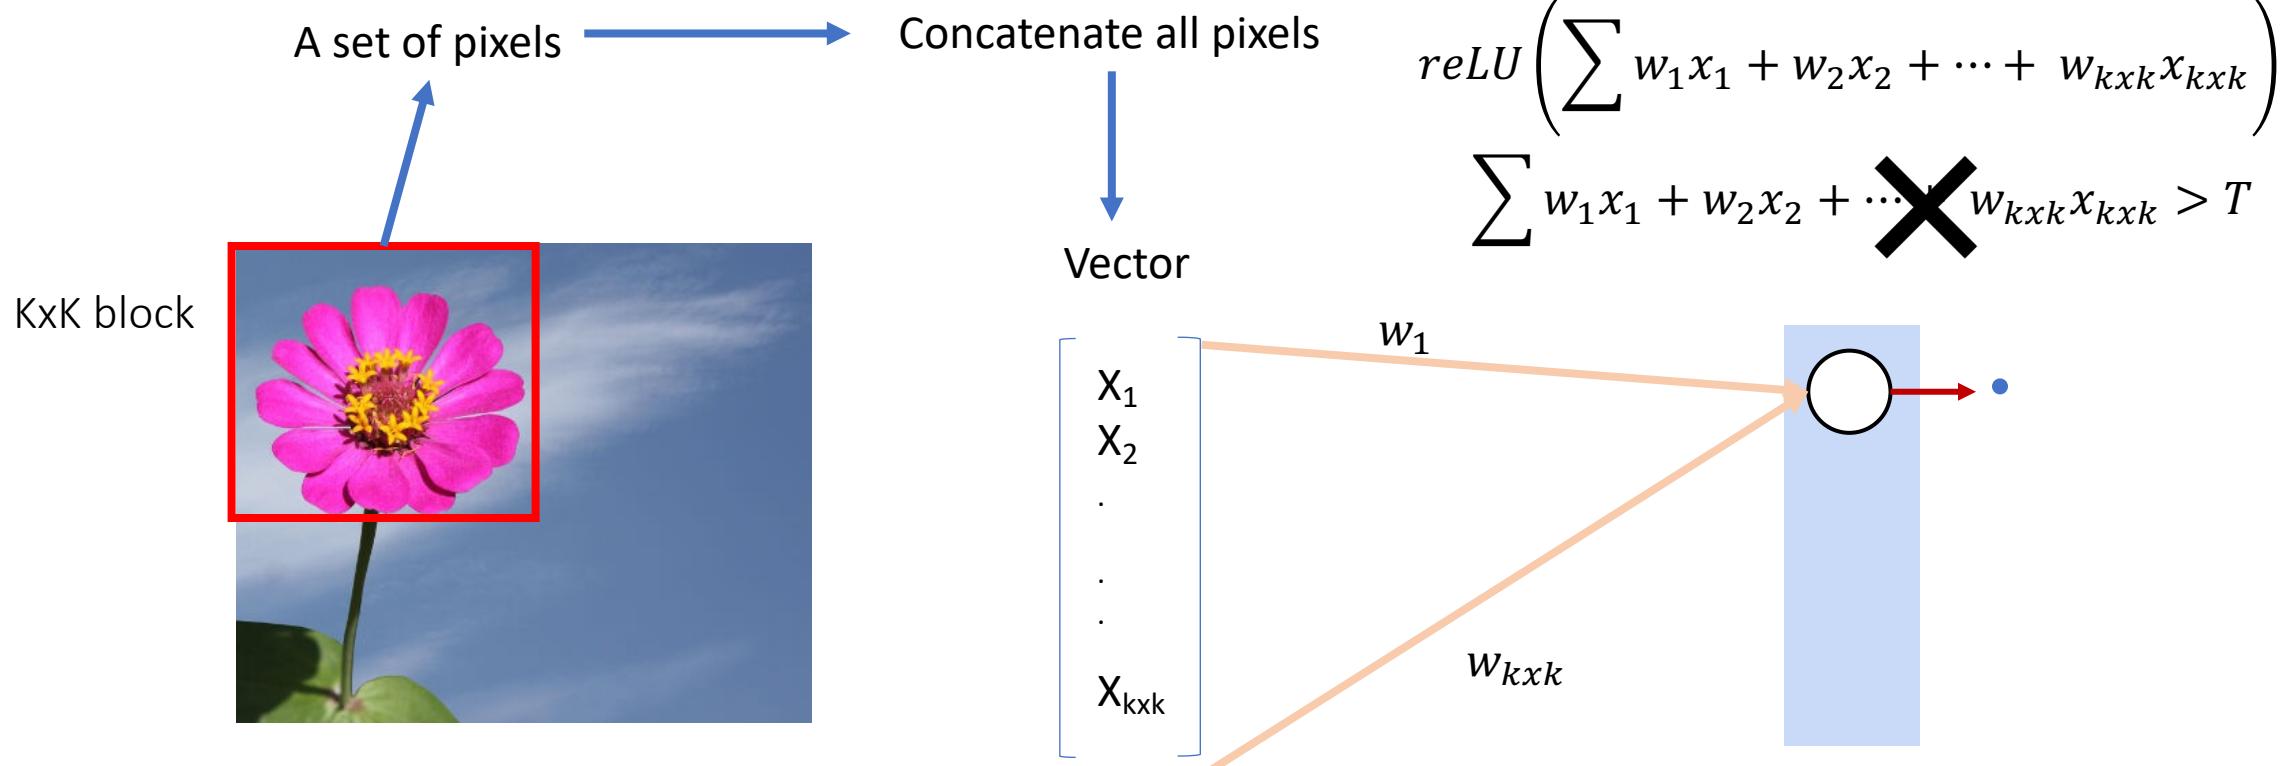


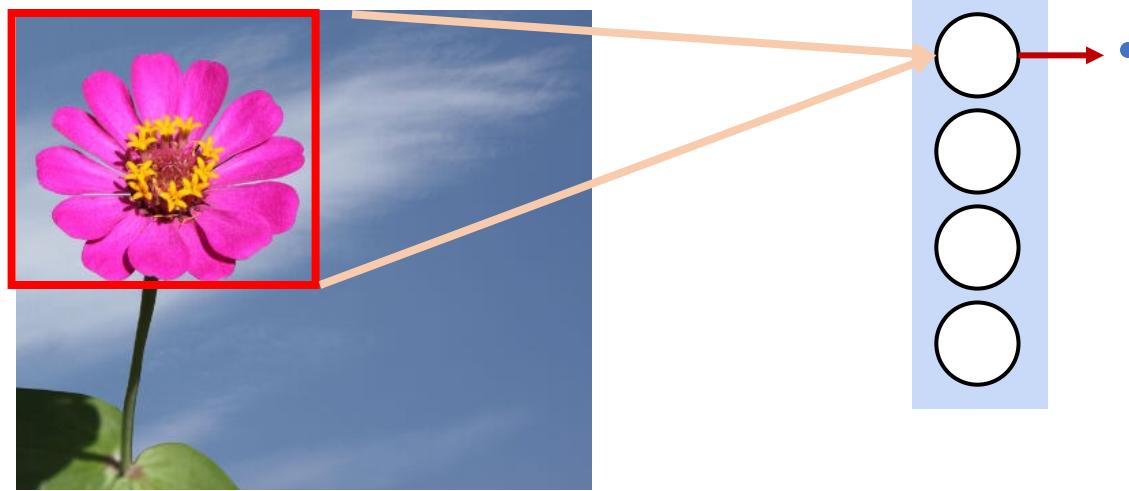


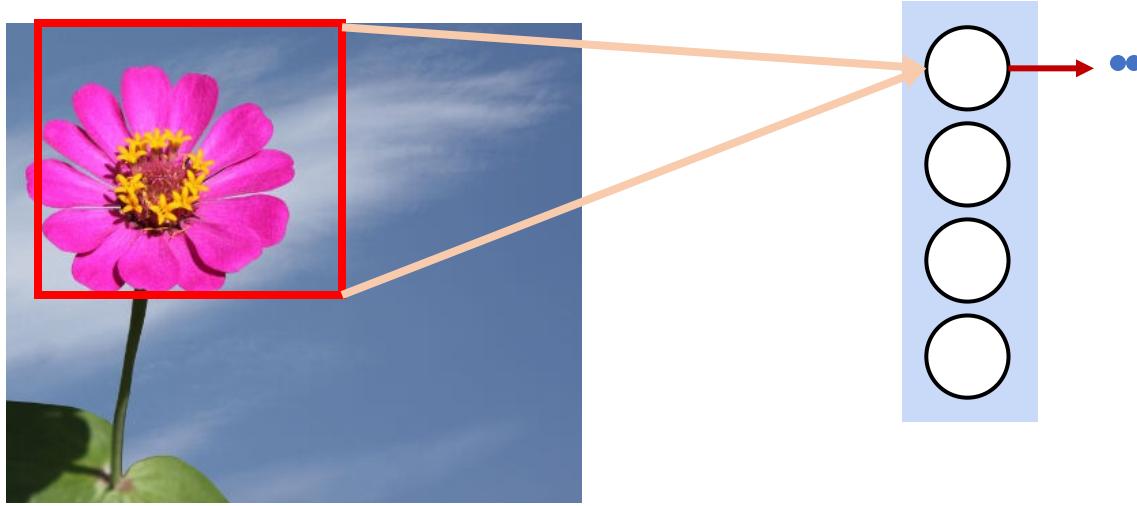
“Look” for the target object at each position

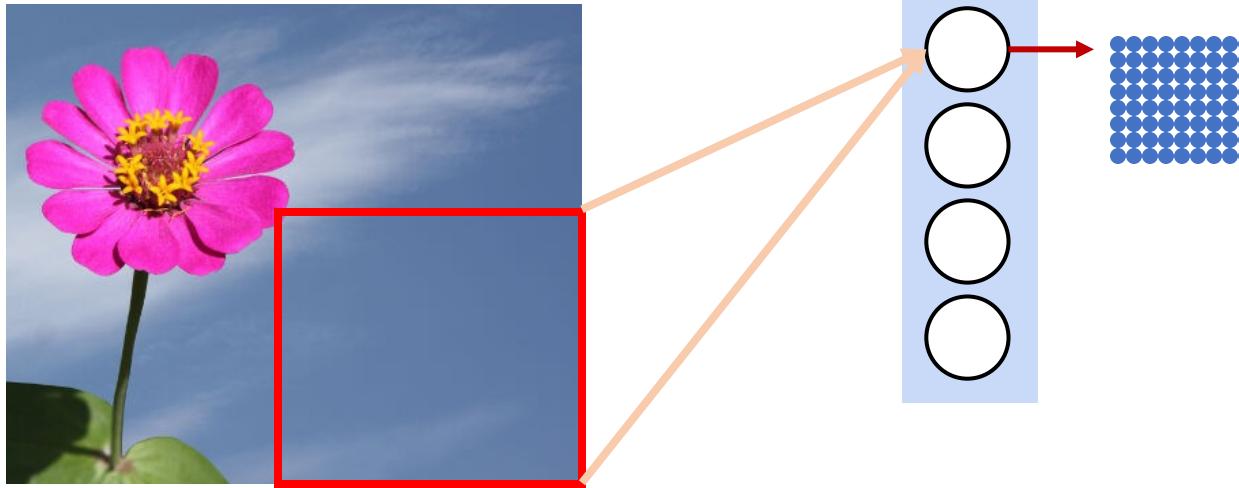


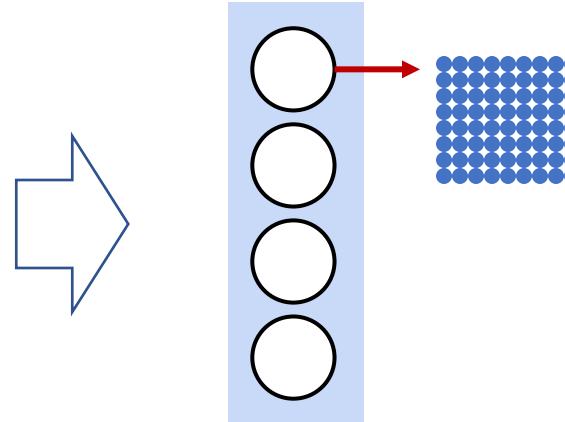


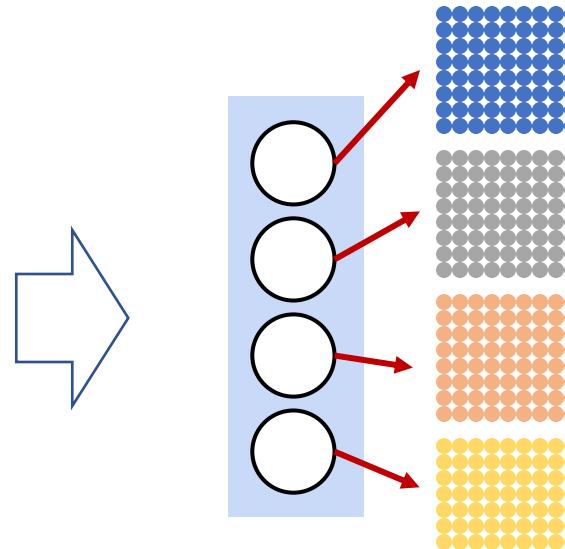


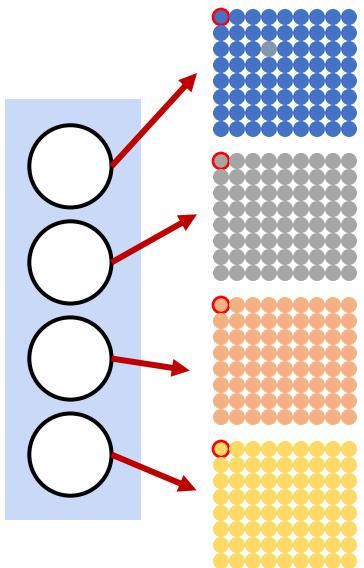


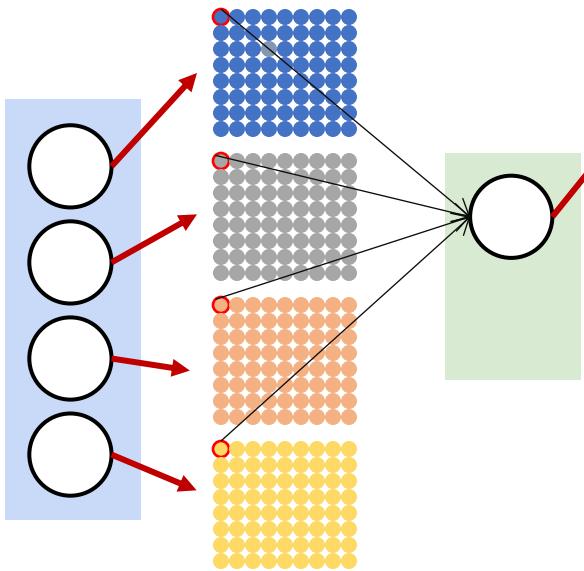


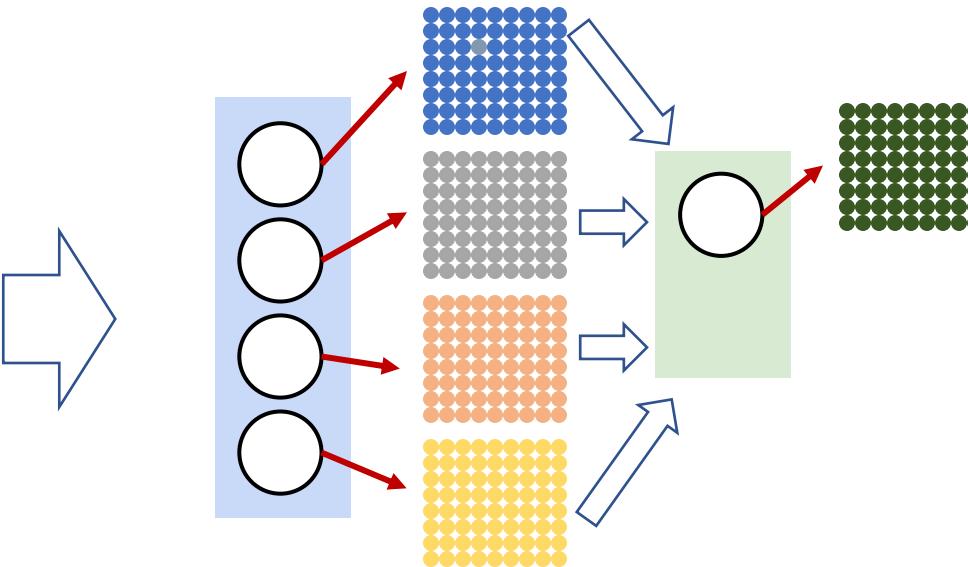


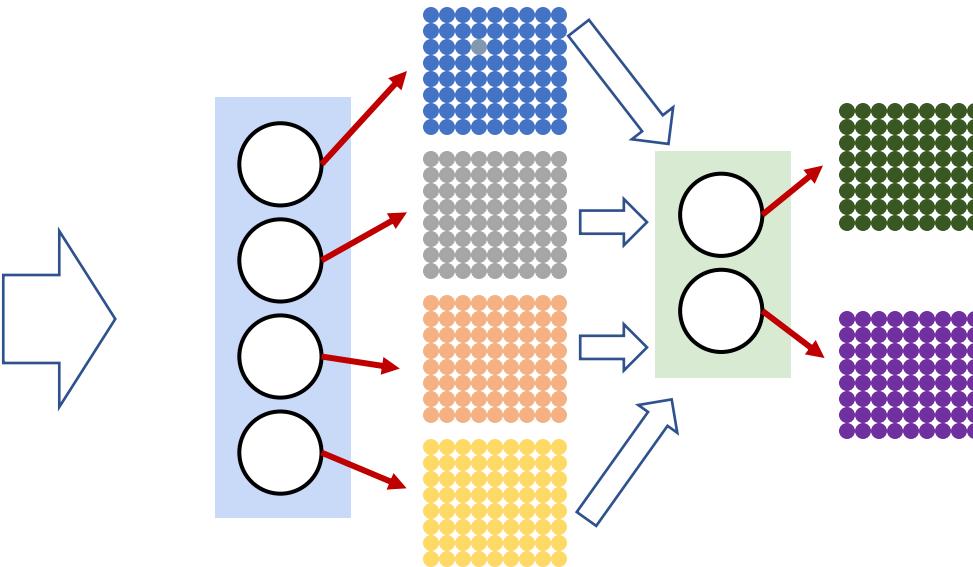


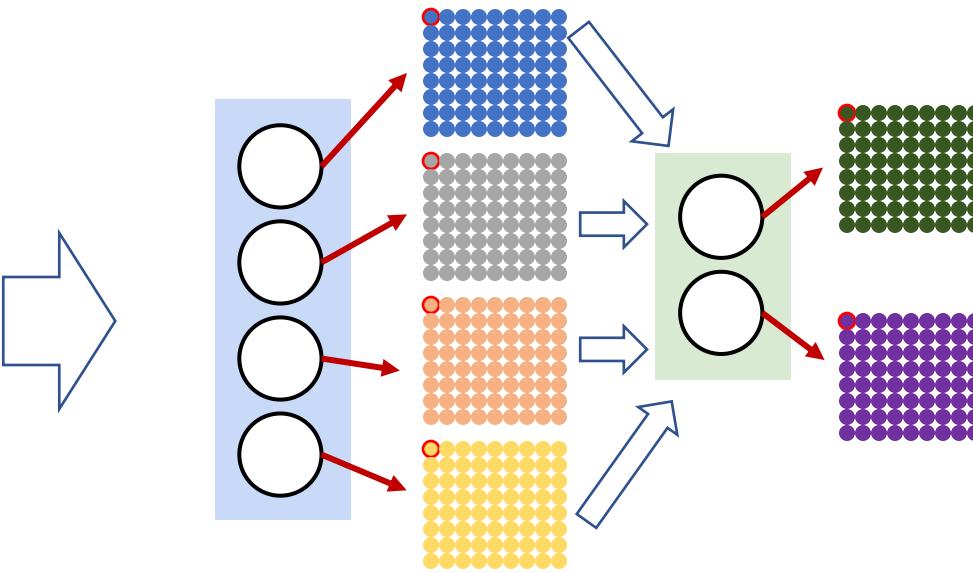


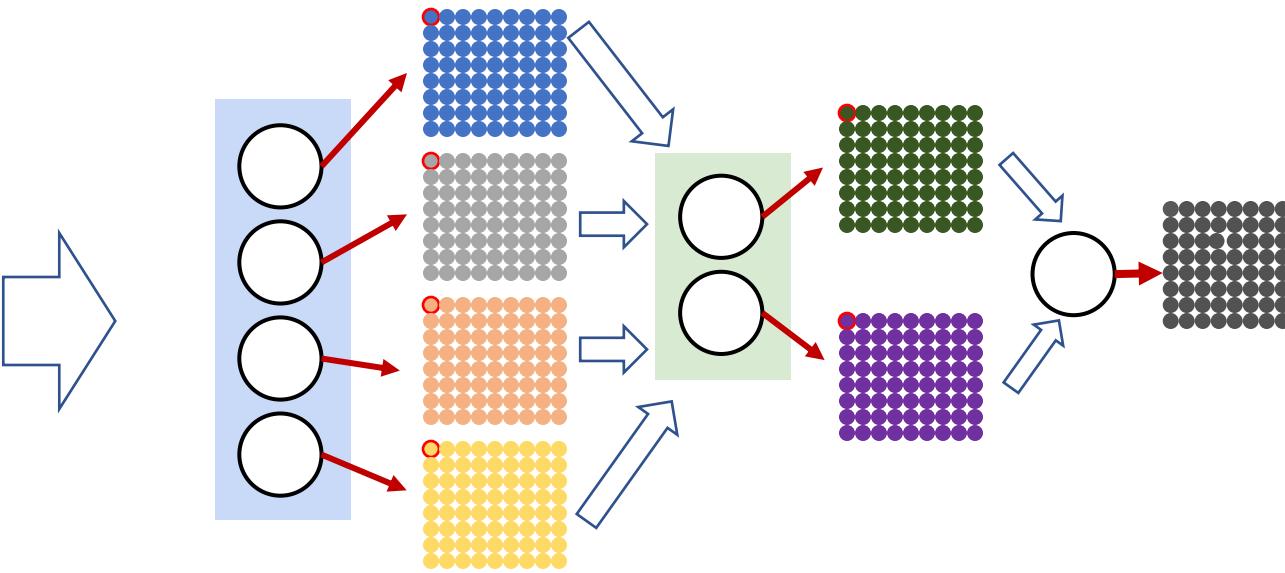


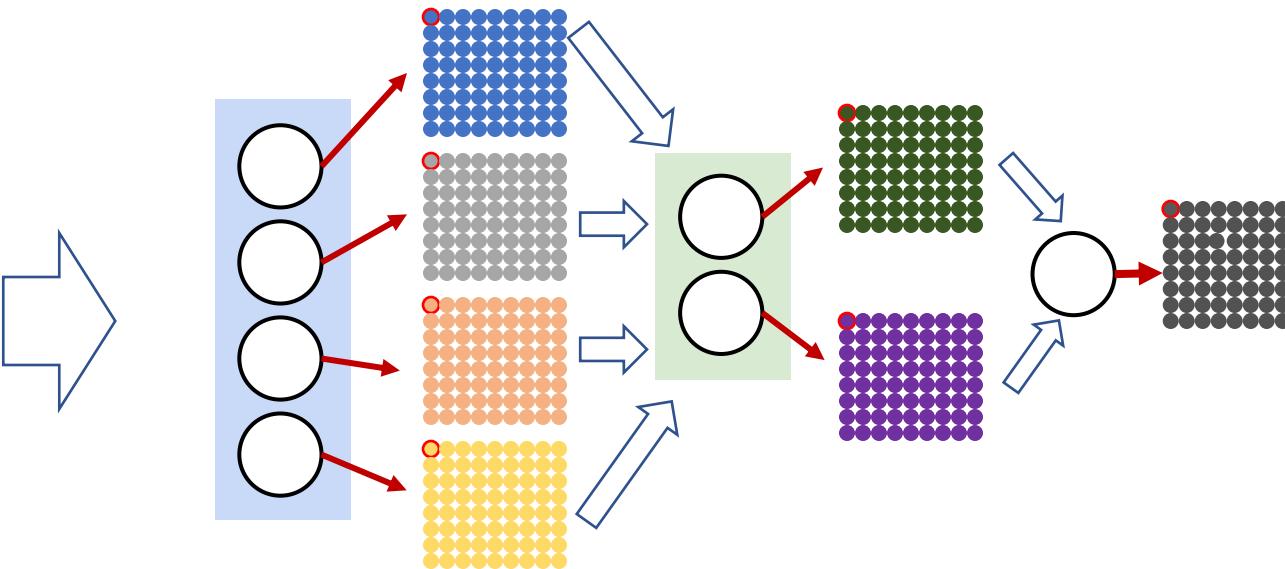


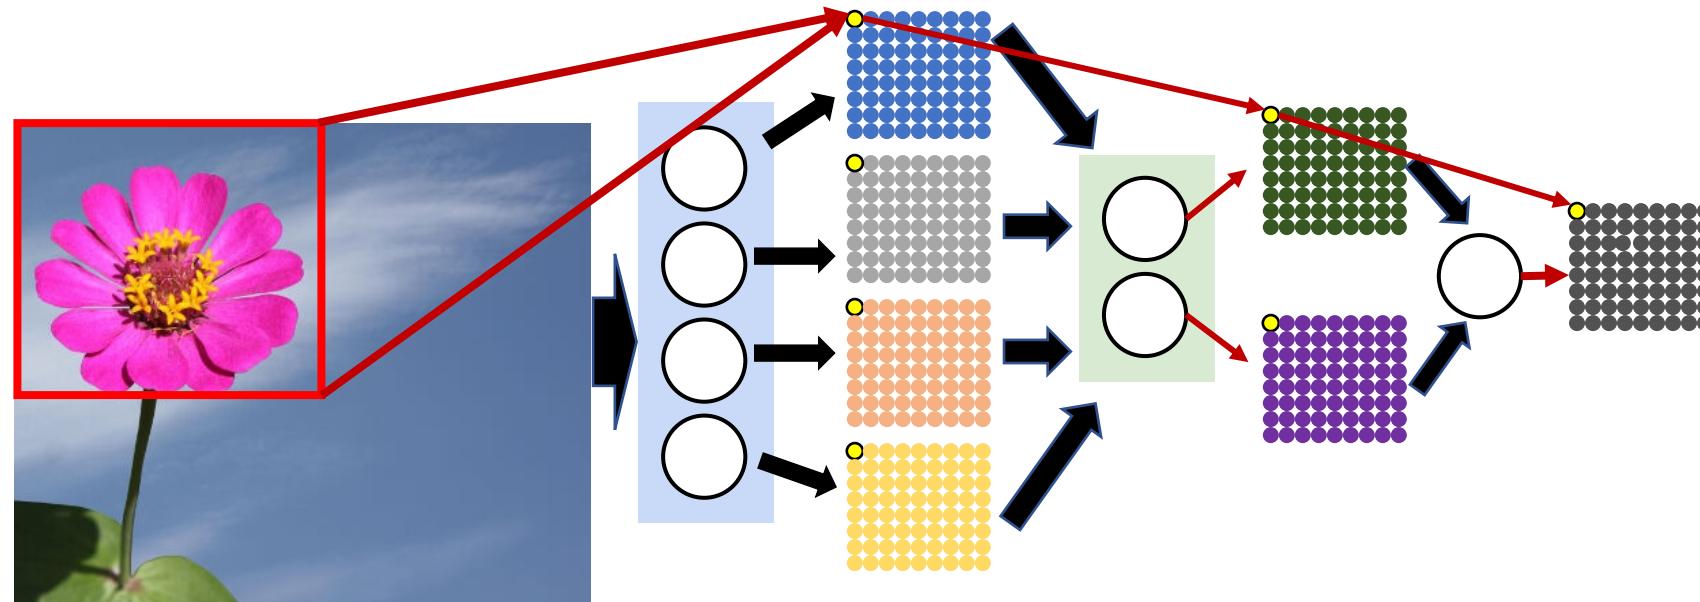




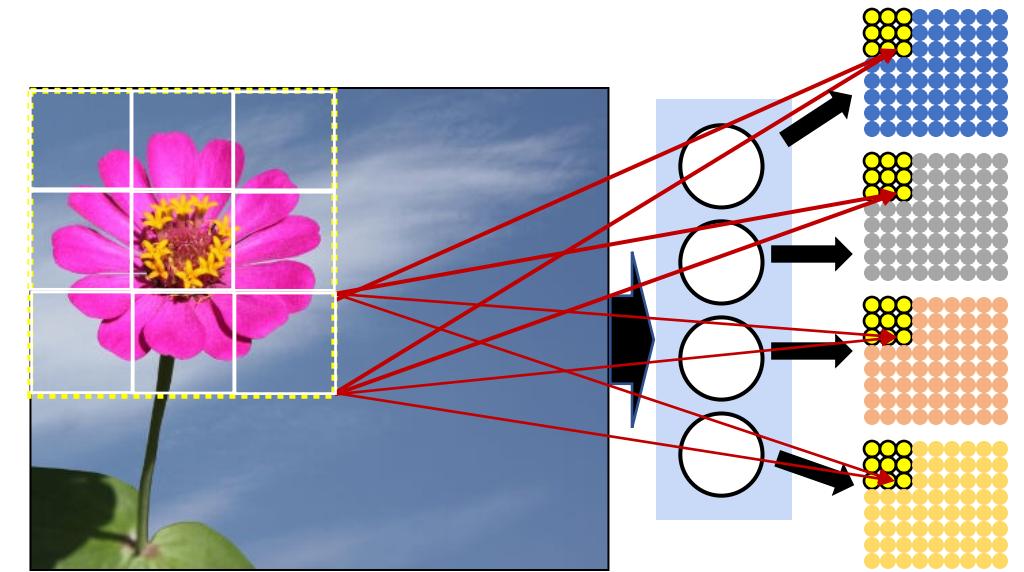






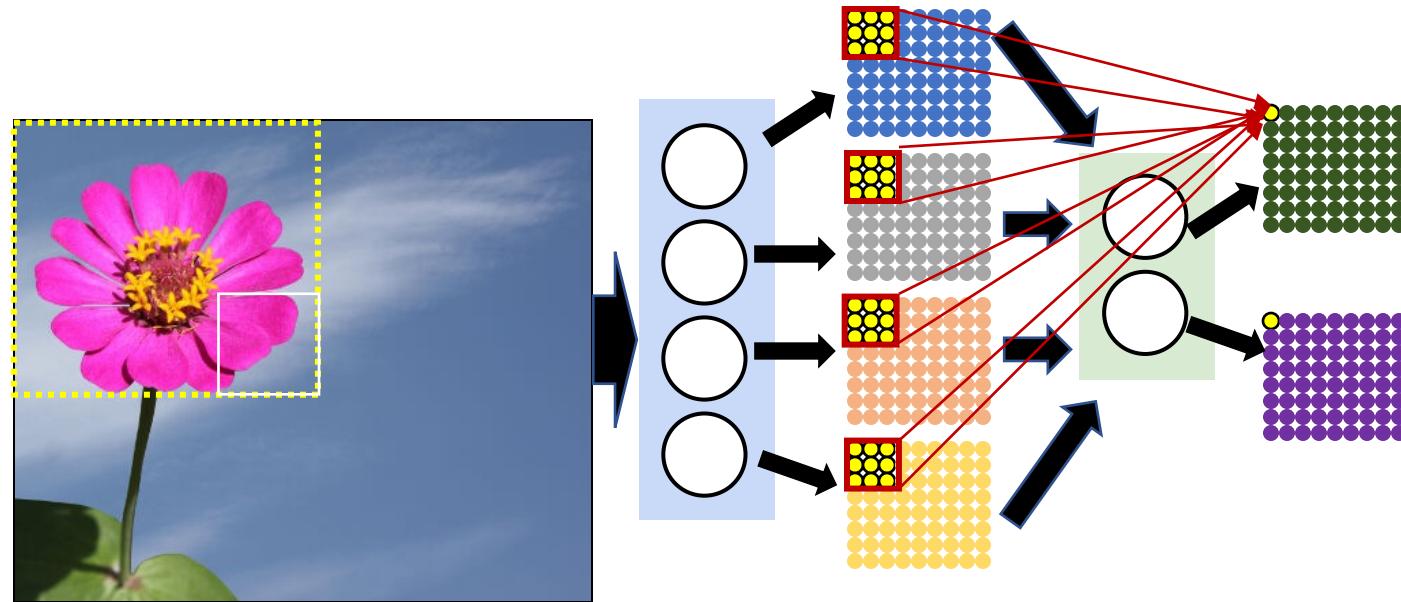


Flower Detection: Distributing the scan



The first layer evaluates smaller blocks of pixels

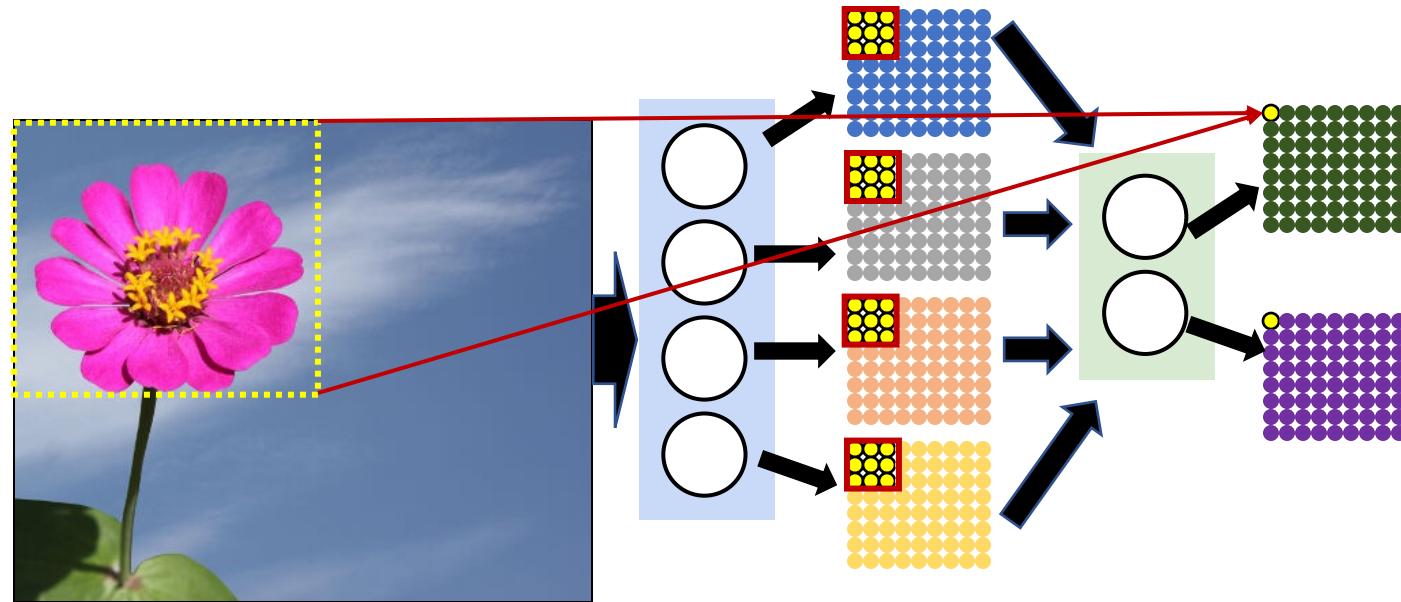
Flower Detection: Distributing the scan



The first layer evaluates smaller blocks of pixels

The next layer evaluates blocks of outputs from the first layer

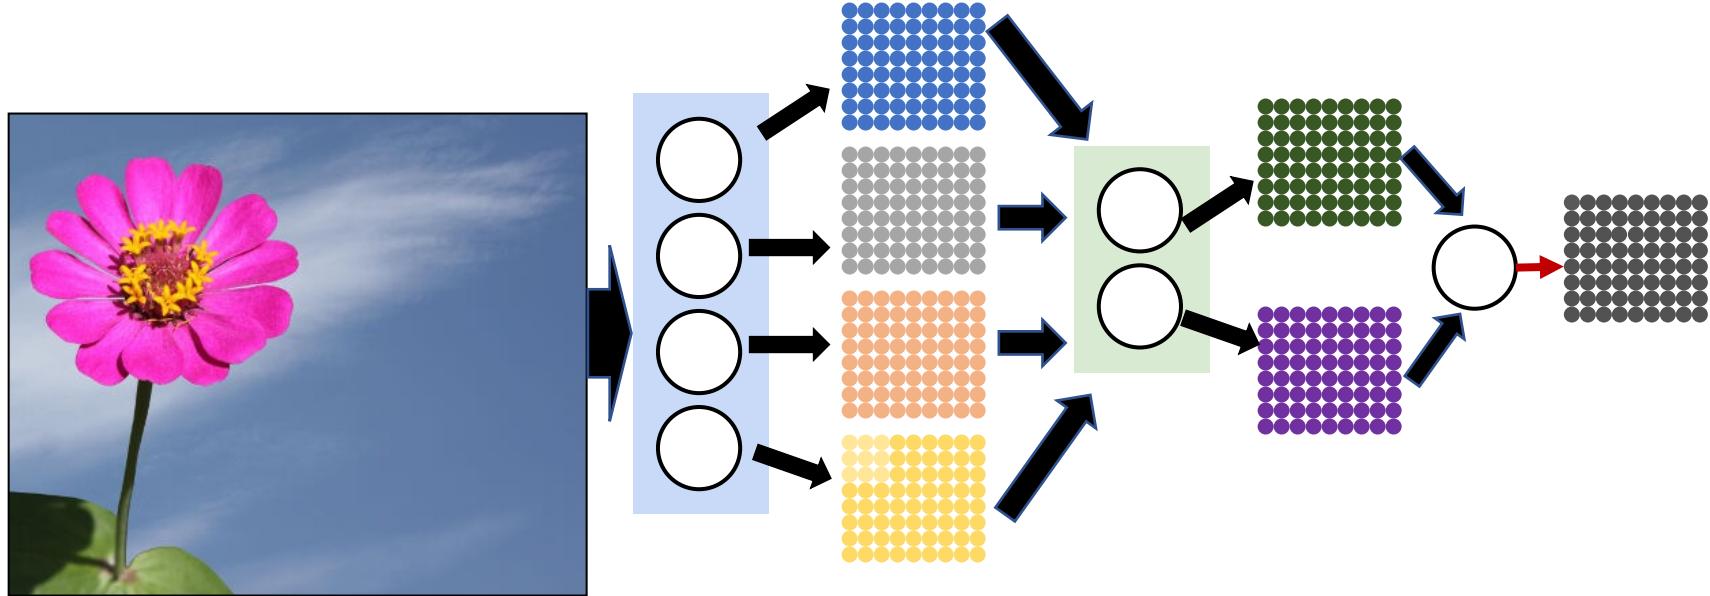
Flower Detection: Distributing the scan



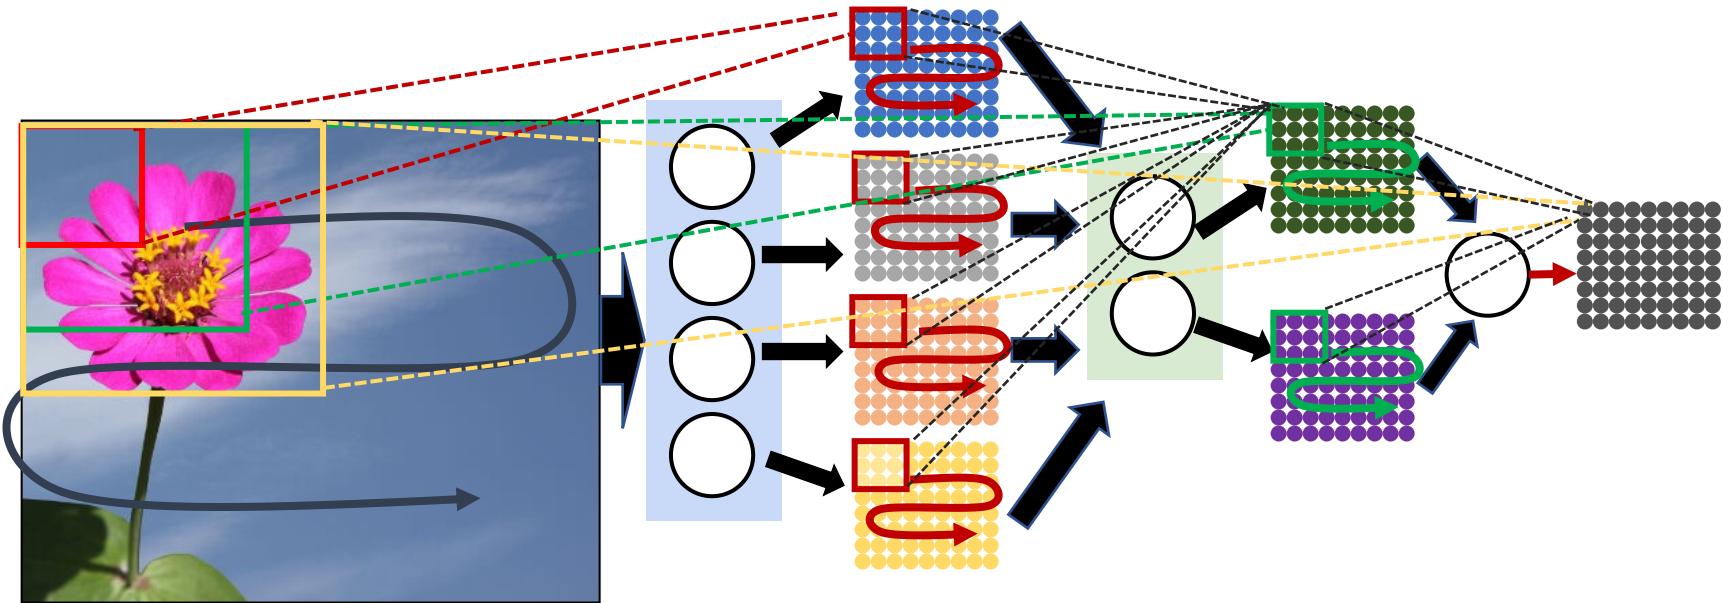
The first layer evaluates smaller blocks of pixels

The next layer evaluates blocks of outputs from the first layer

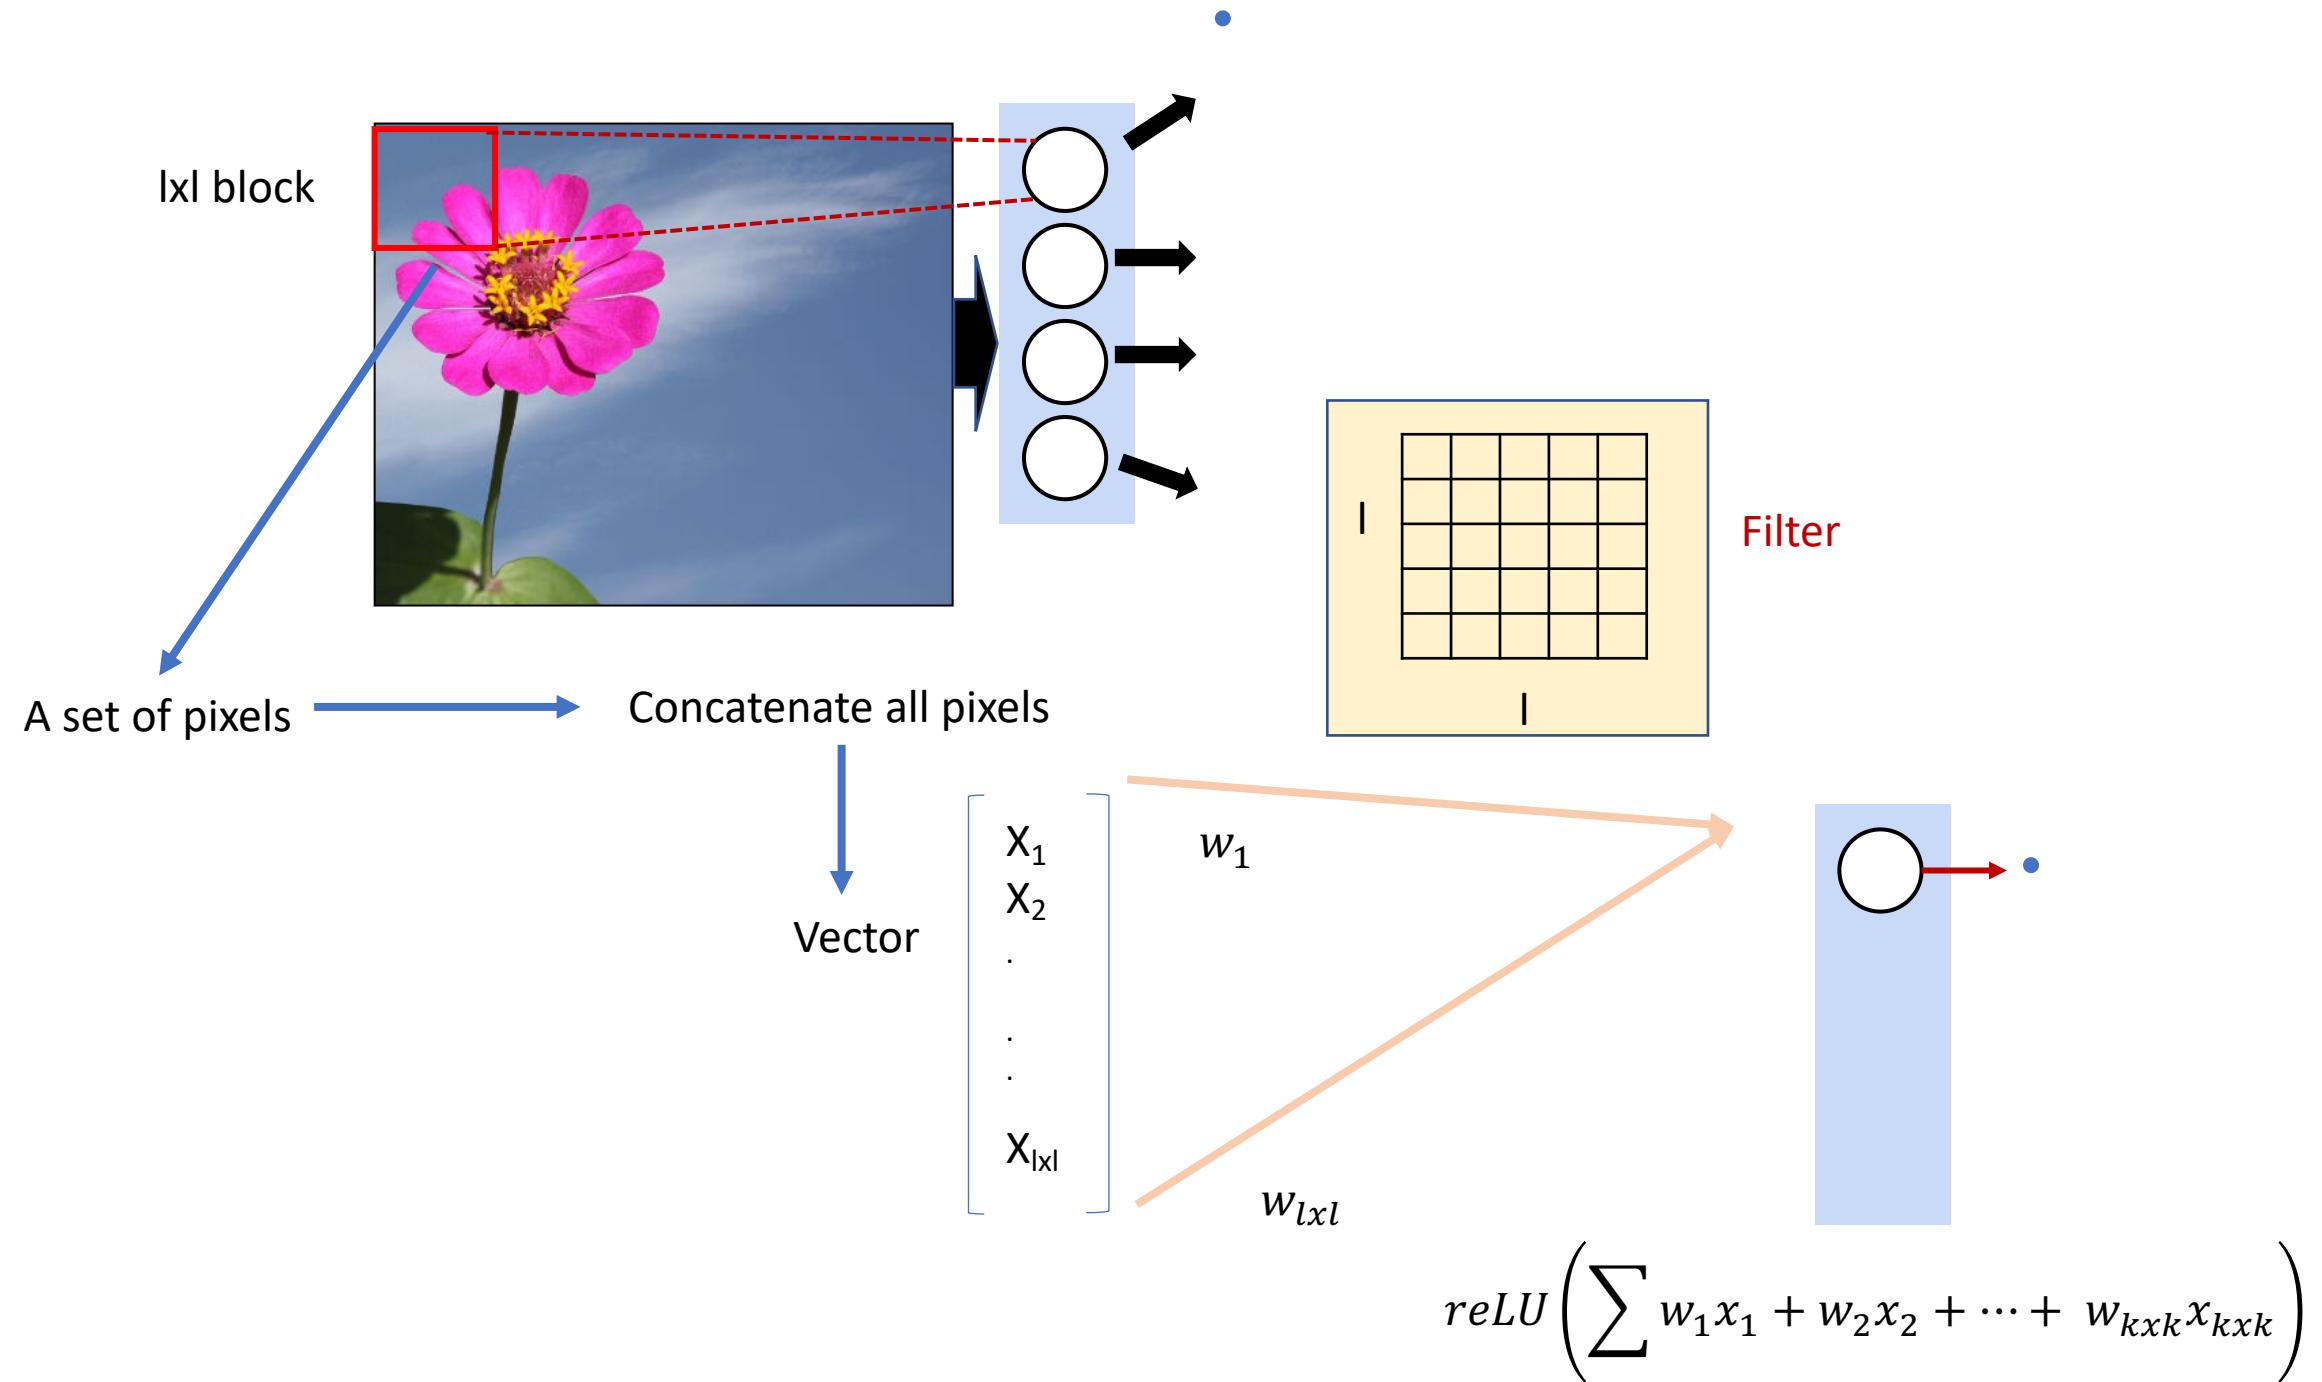
Distribute over 2 layers



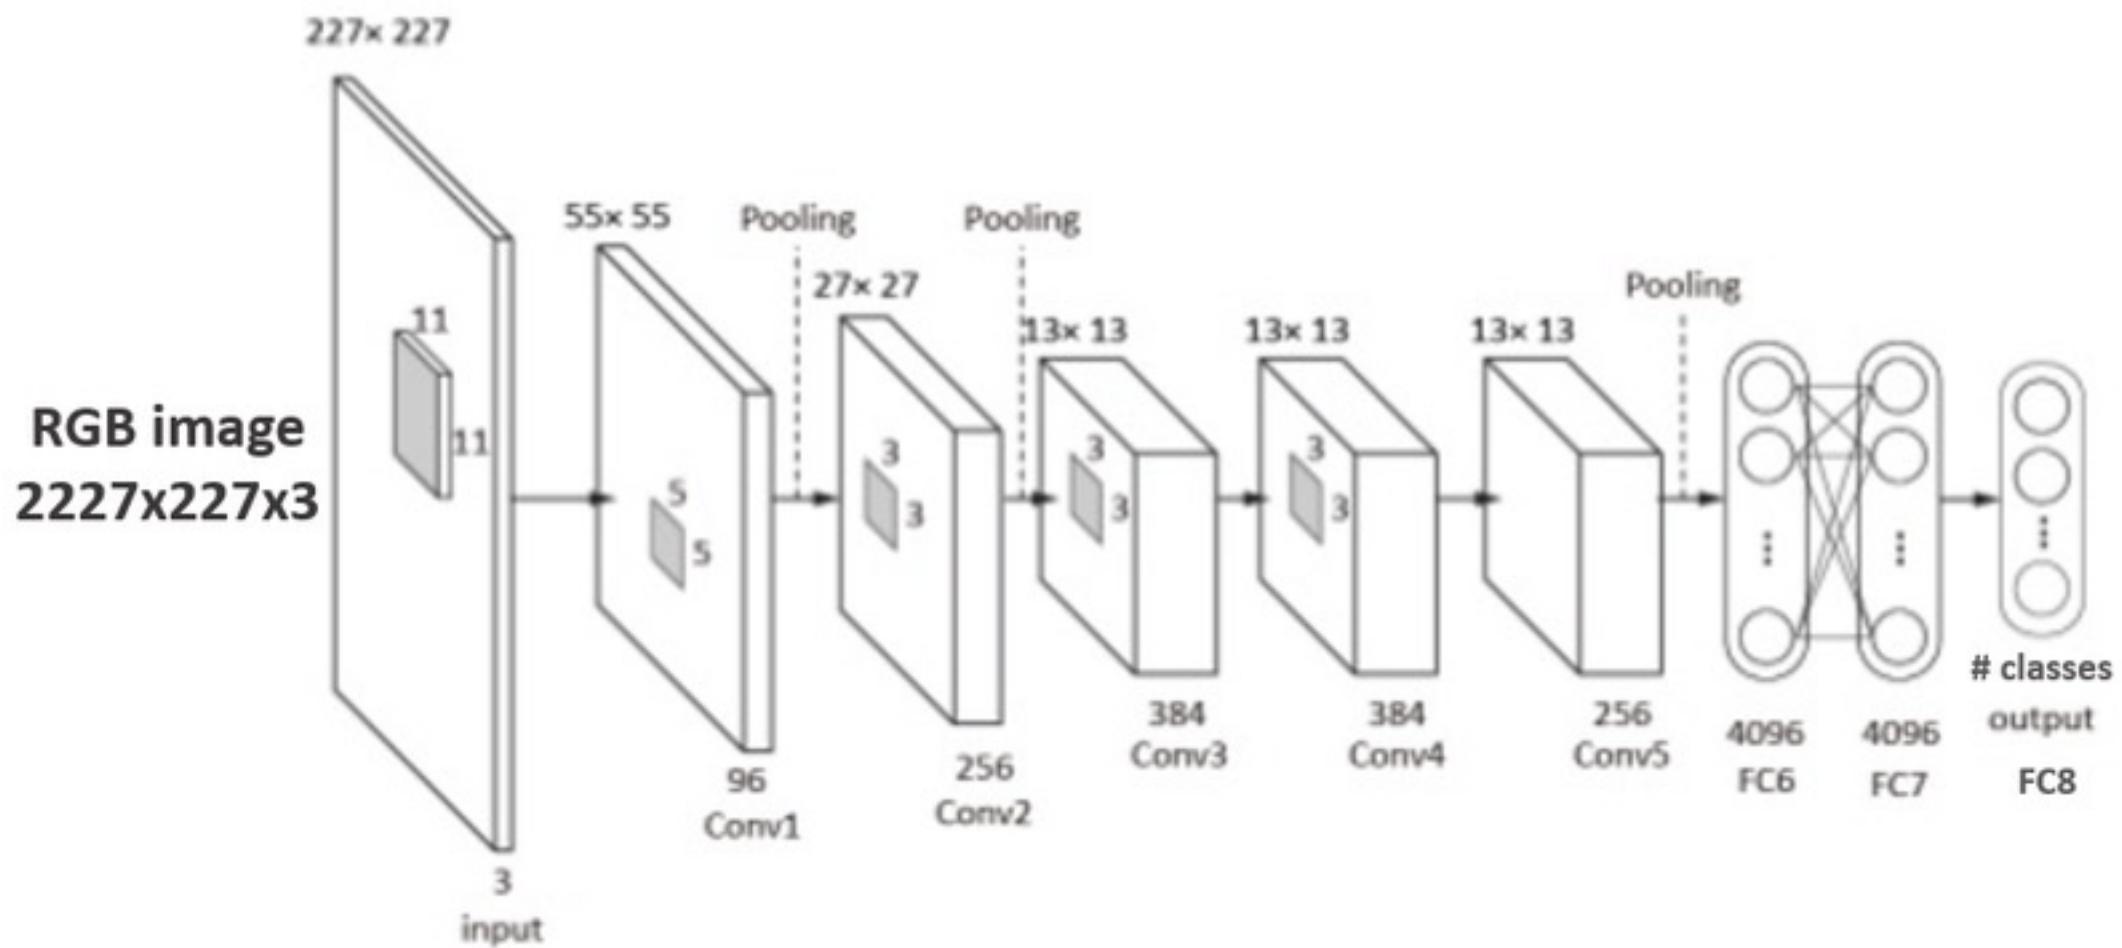
Distribute over 3 layers

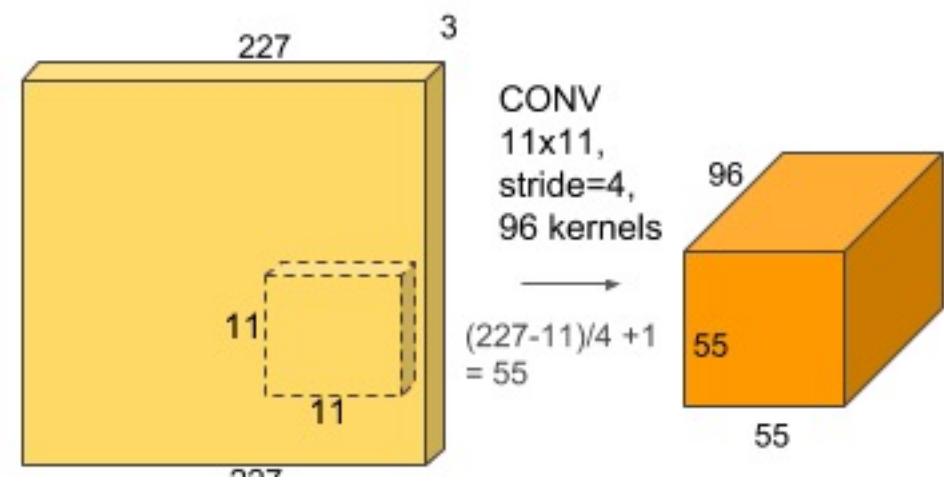


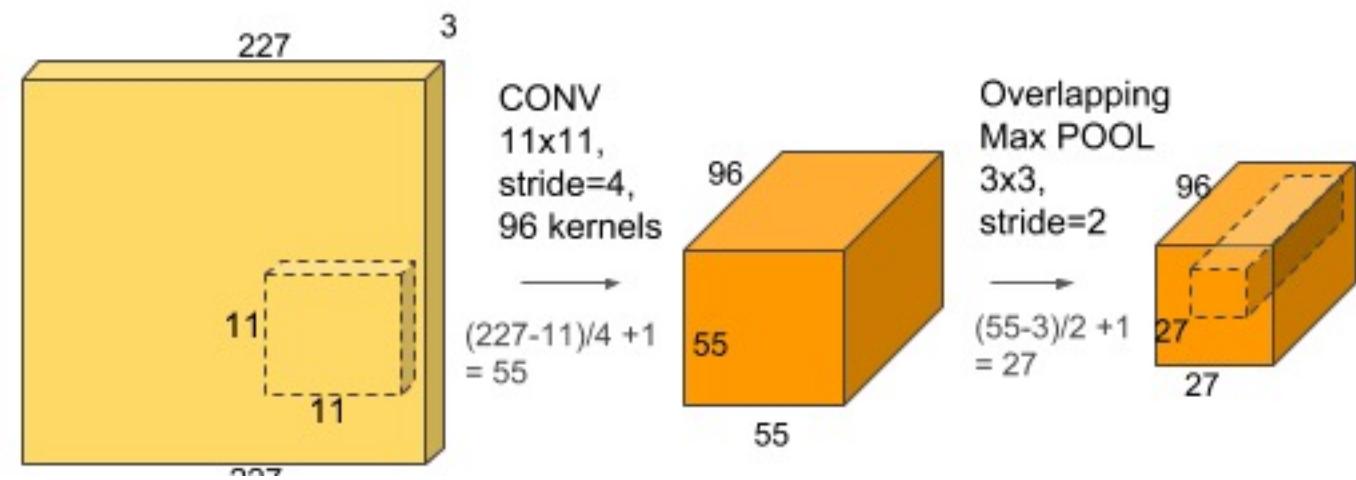
Convolutional Neural Networks (CNNs)

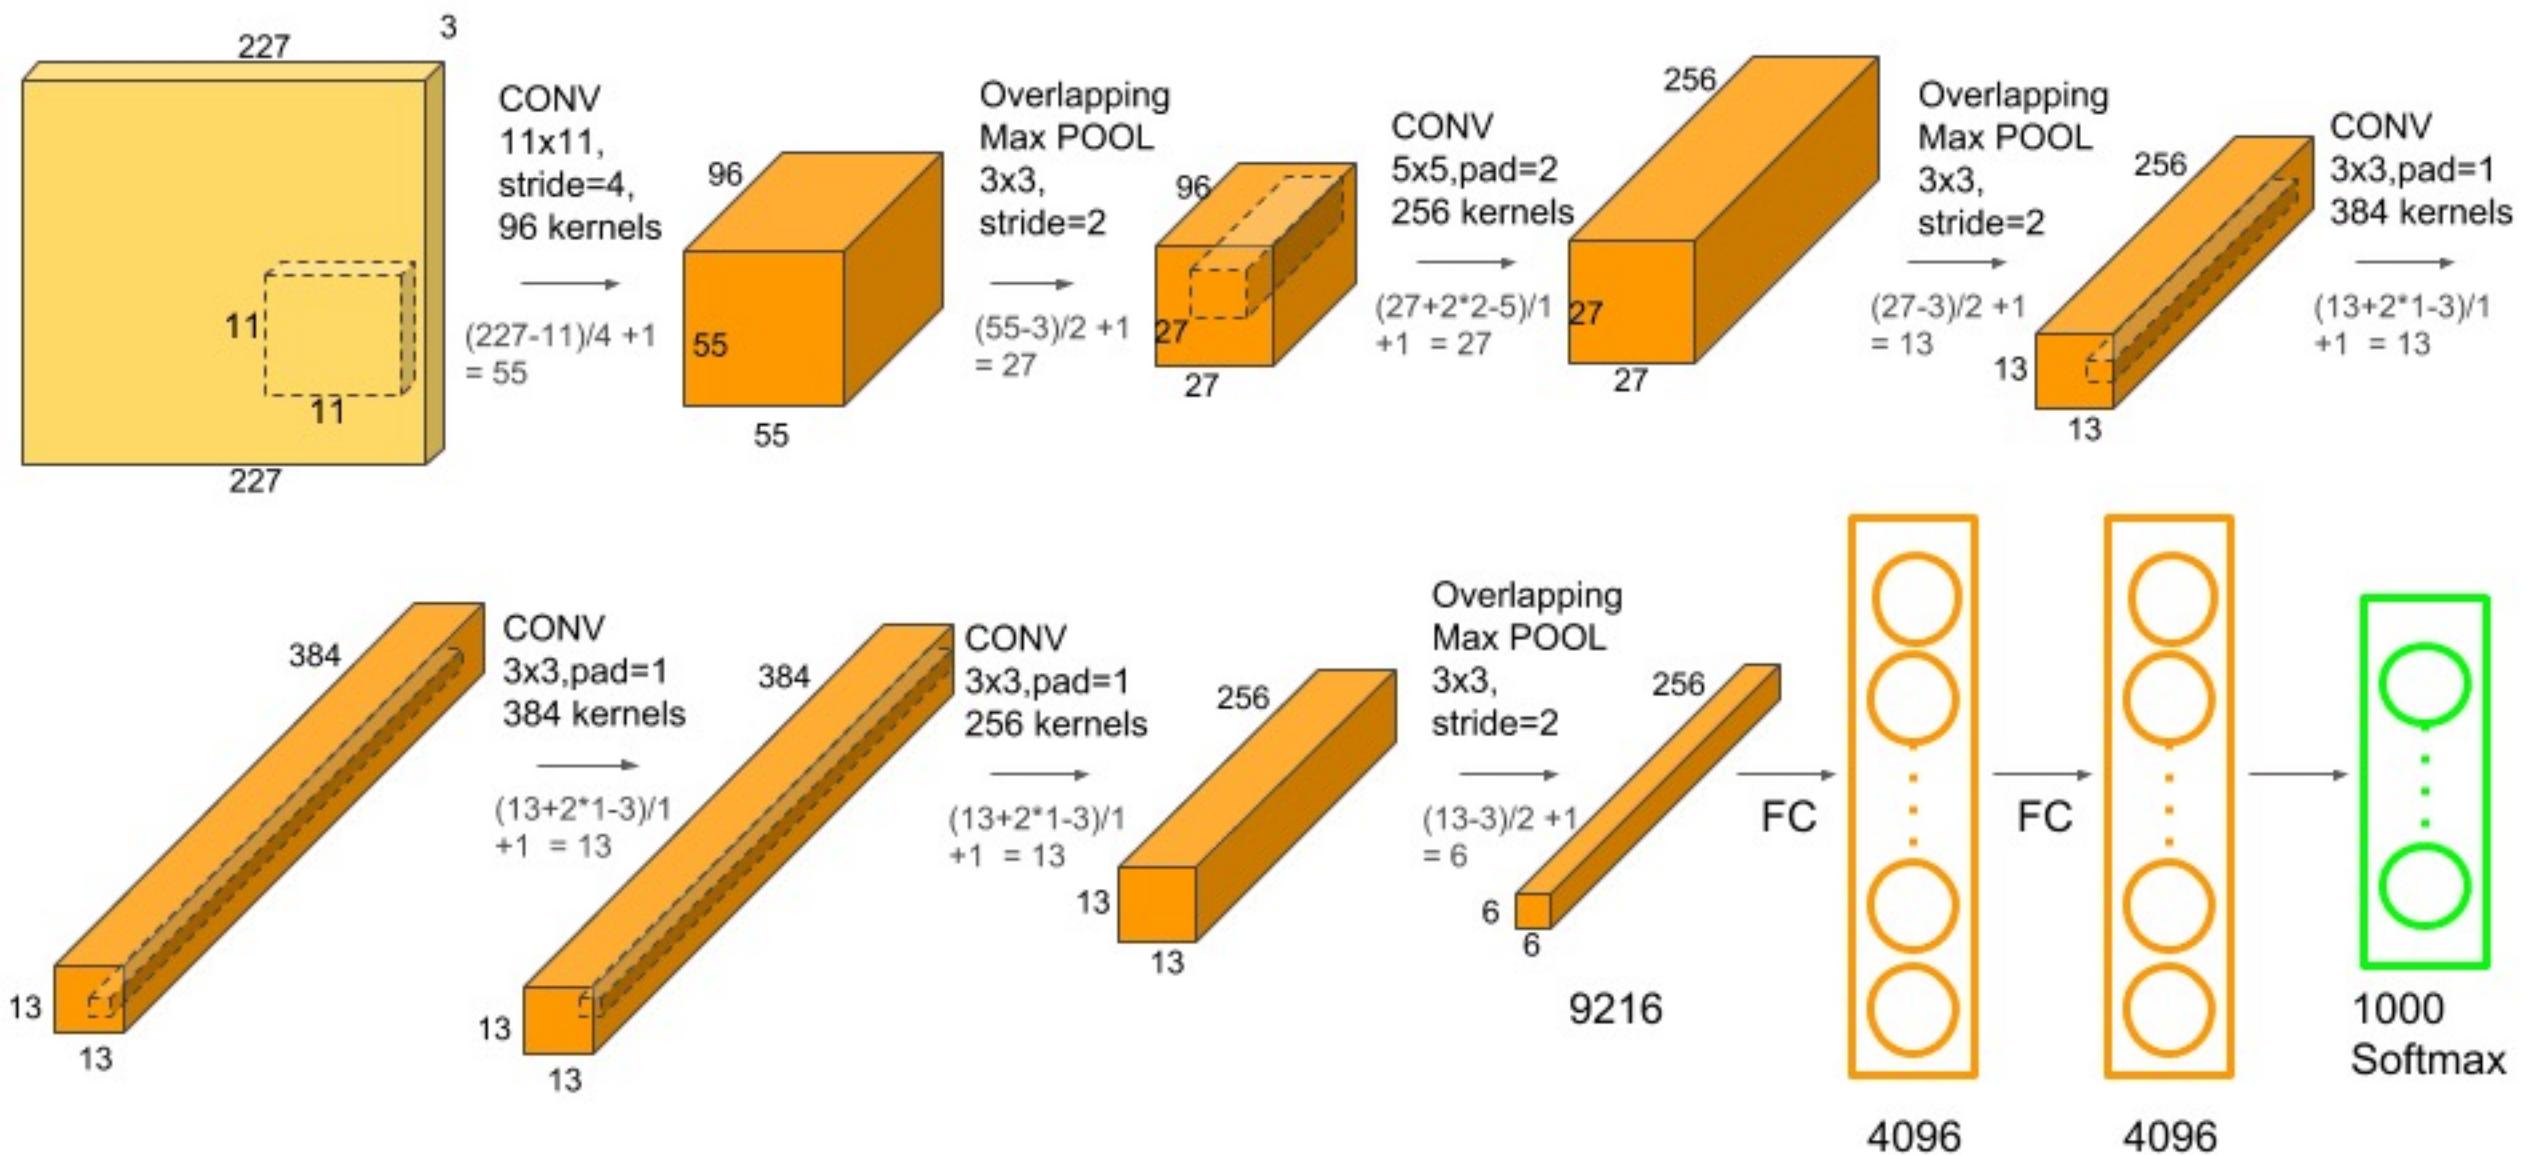


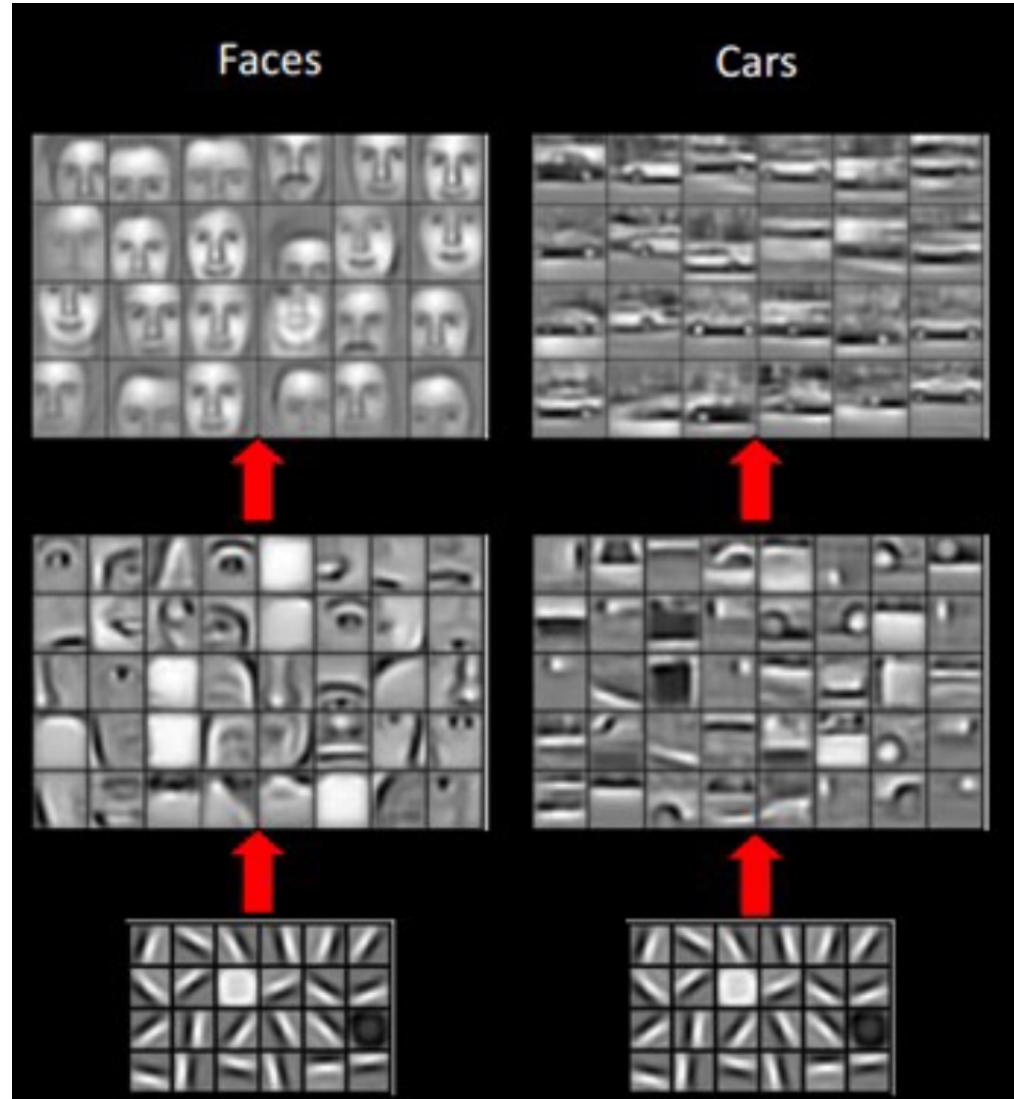
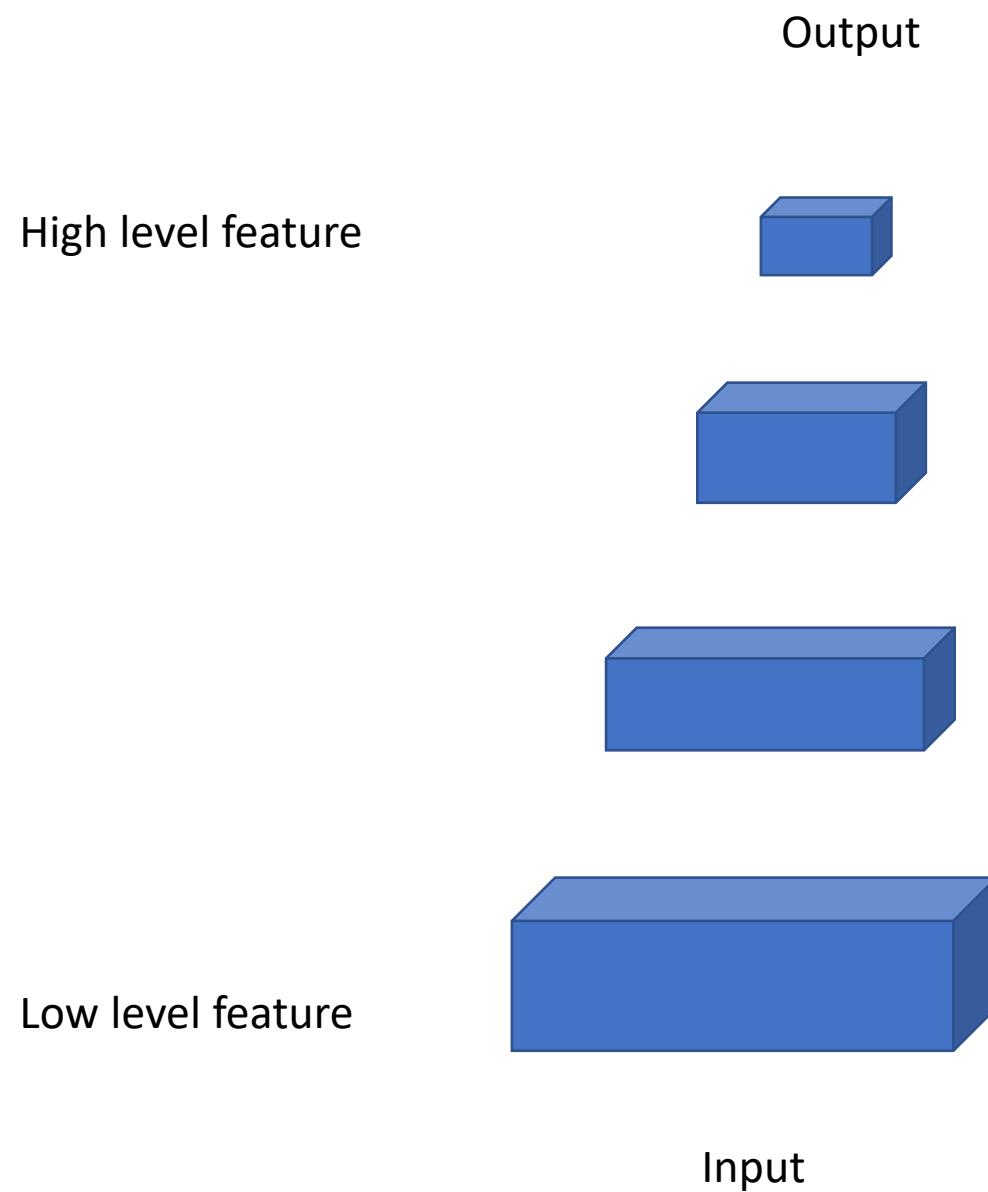
AlexNet







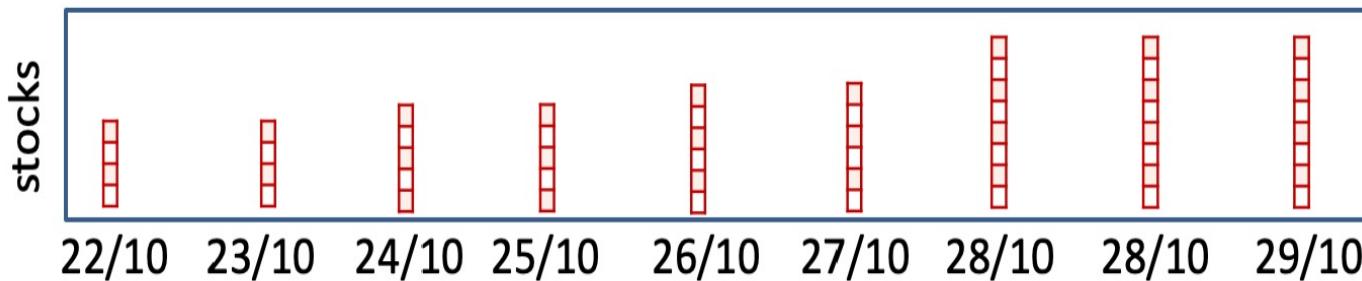




Recap

- Perceptron
- Multi-layer Perceptron
- Convolutional Neural Network
 - Detection ?
 - Classification ?
 - Prediction ?
 - Language Translation?

Recurrent Neural Networks

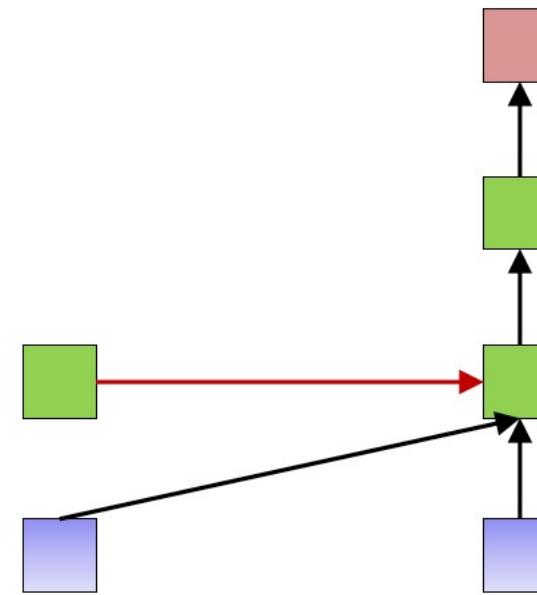
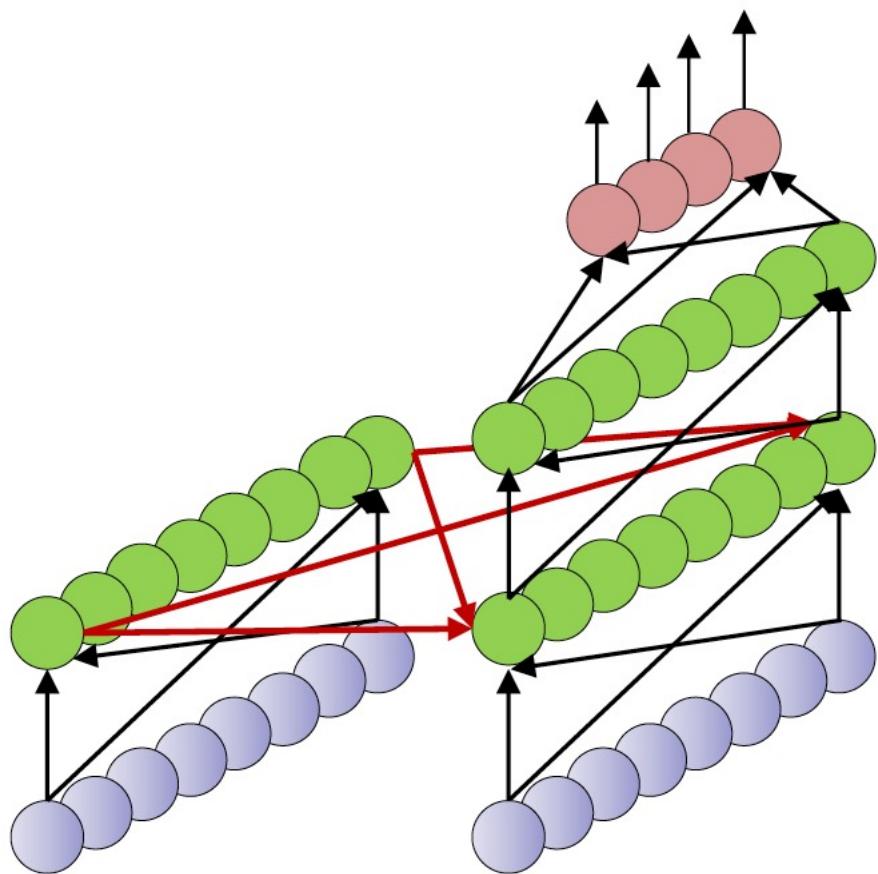


Stock Prediction:

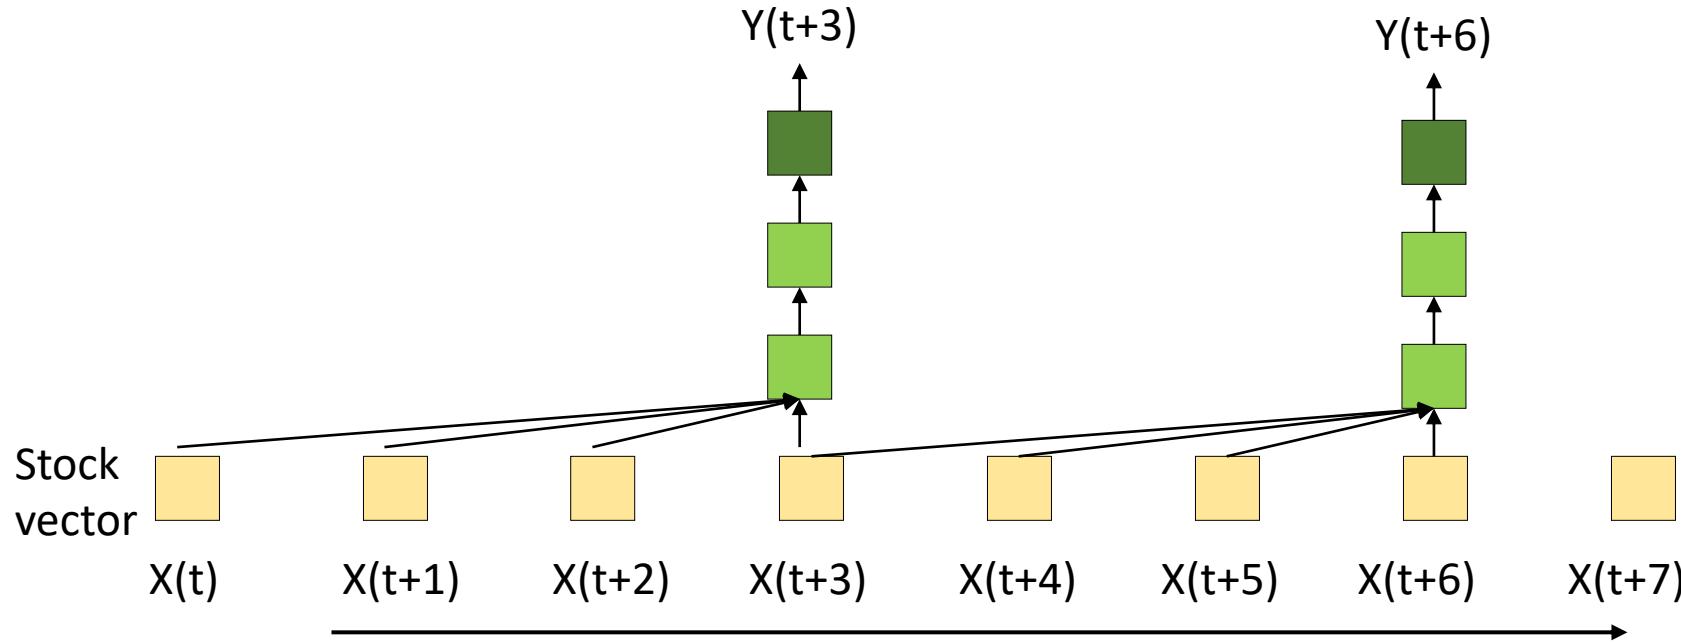
Inputs: are vectors.

Output: may be scalar or vector

Representation Shortcut

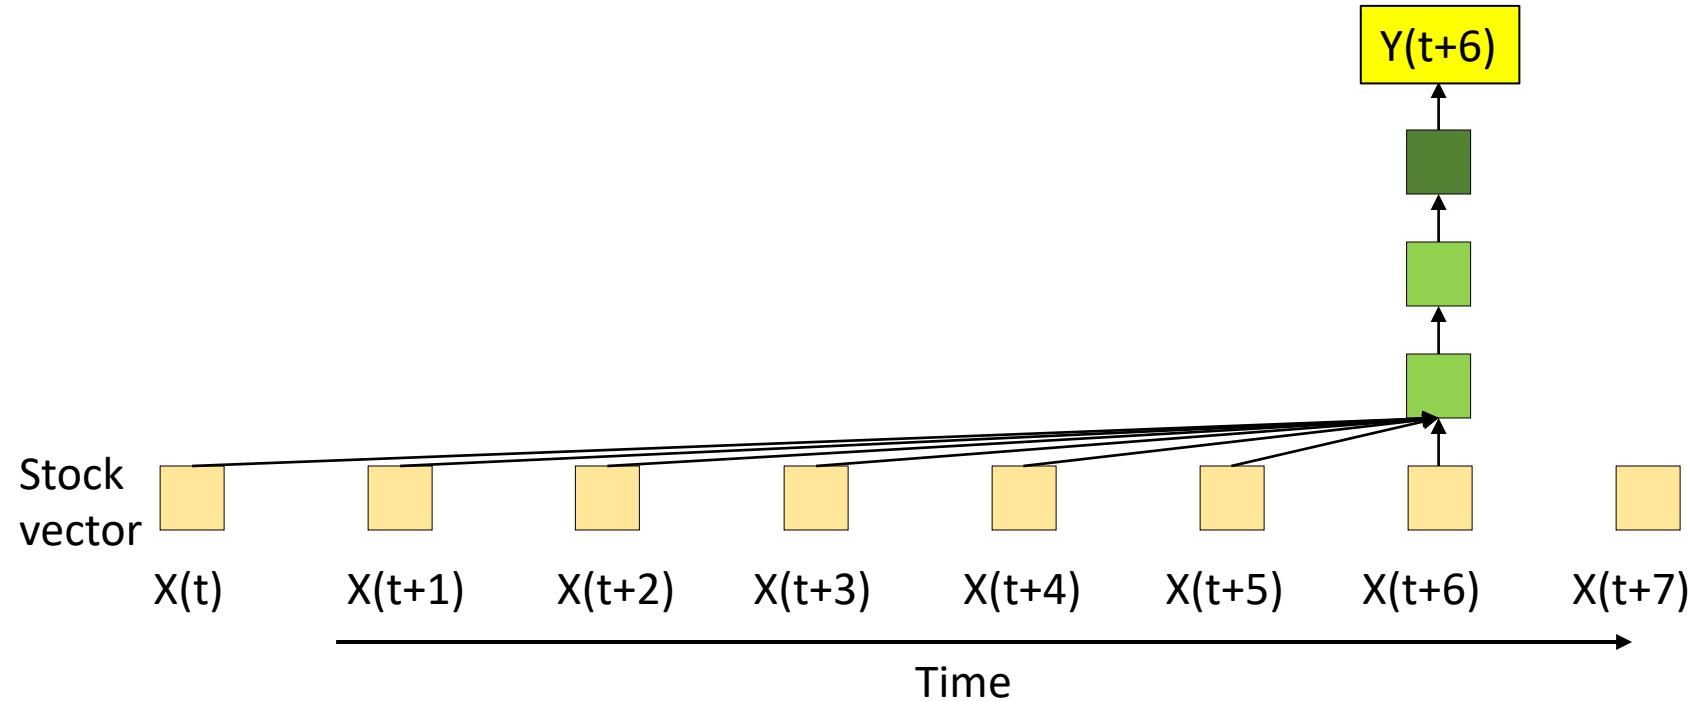


The stock predictor



$$Y_t = f(X_t, X_{t-1}, \dots, X_{t-N})$$

The stock predictor

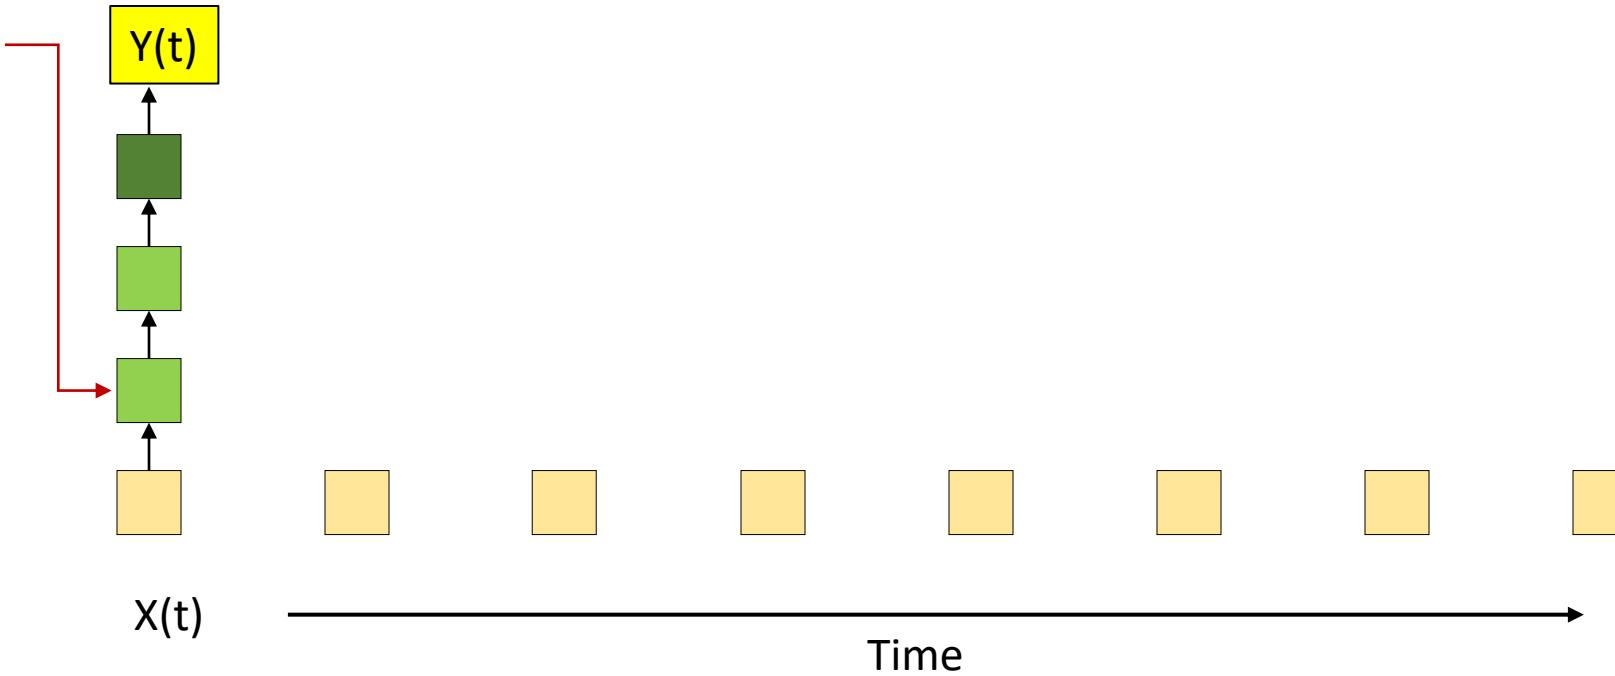


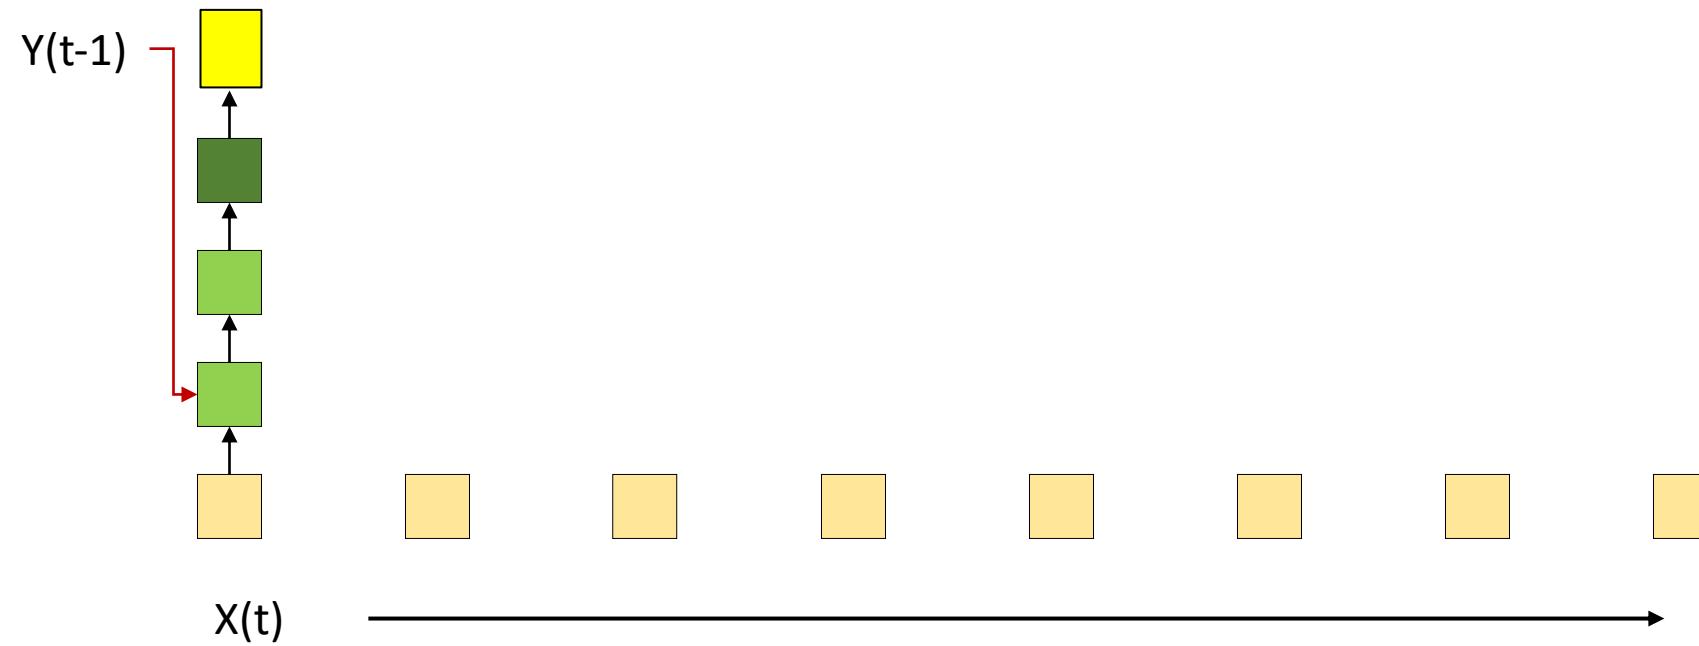
$$Y_t = f(X_t, X_{t-1}, \dots, X_{t-\infty})$$

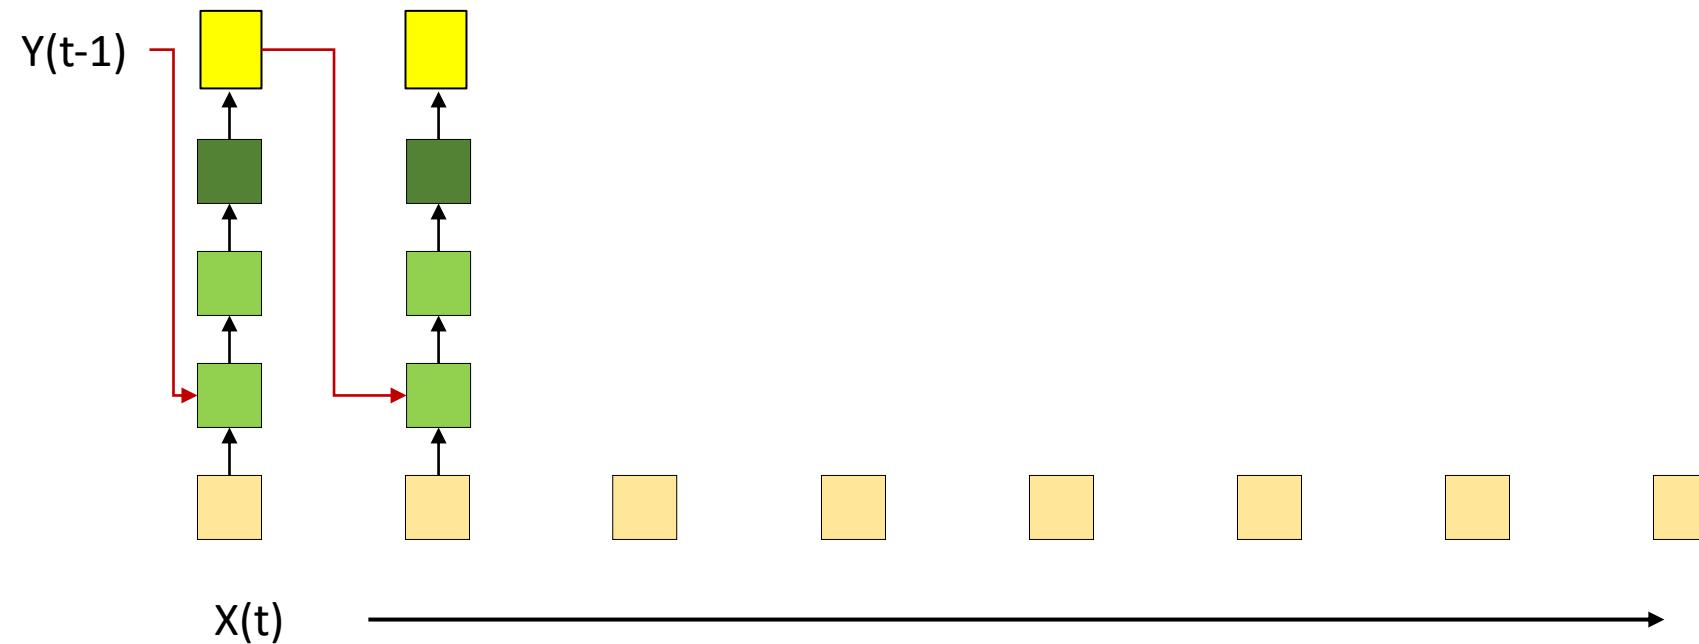
$$Y_t = f(X_t, X_{t-1}, \dots, X_{t-\infty})$$

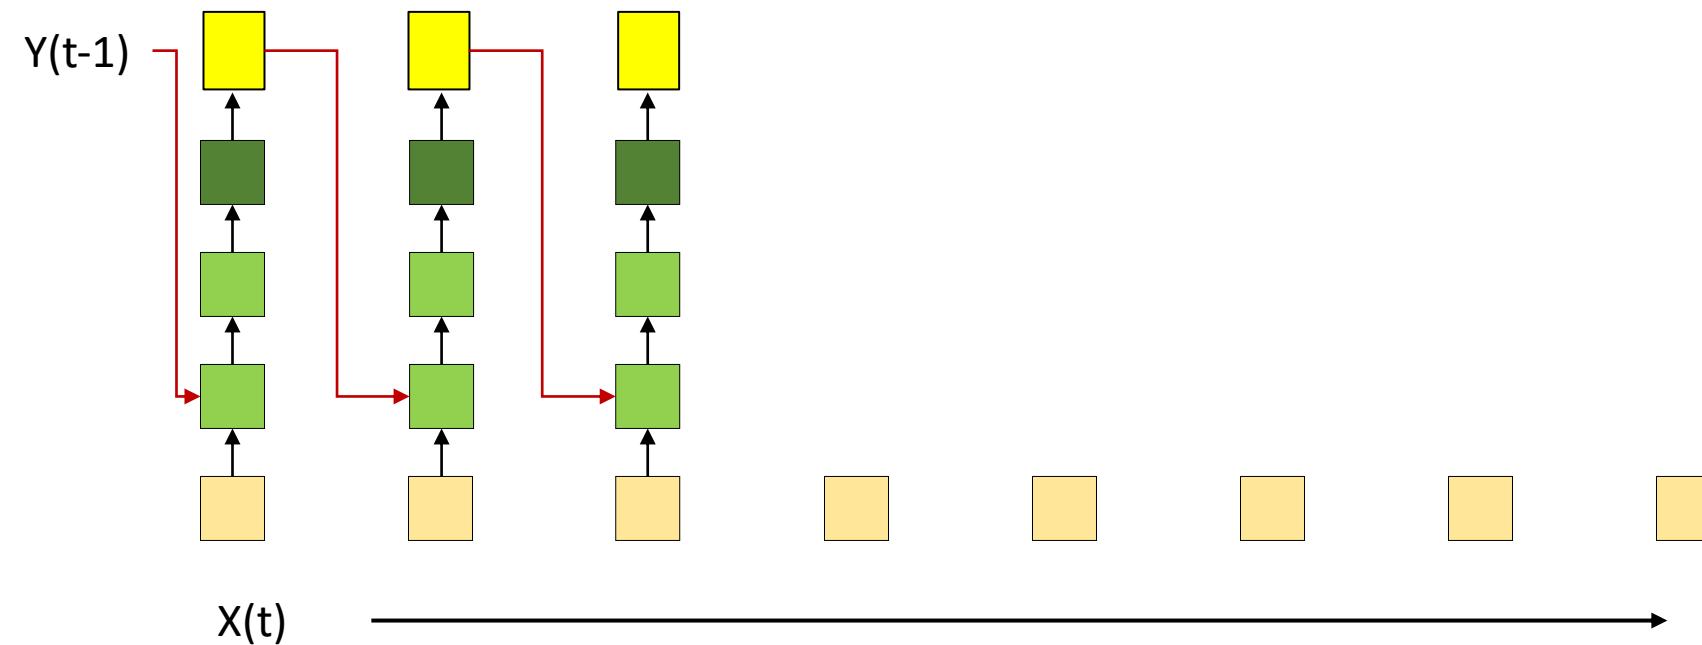
An input at X_0 at $t = 0$ produces Y_0

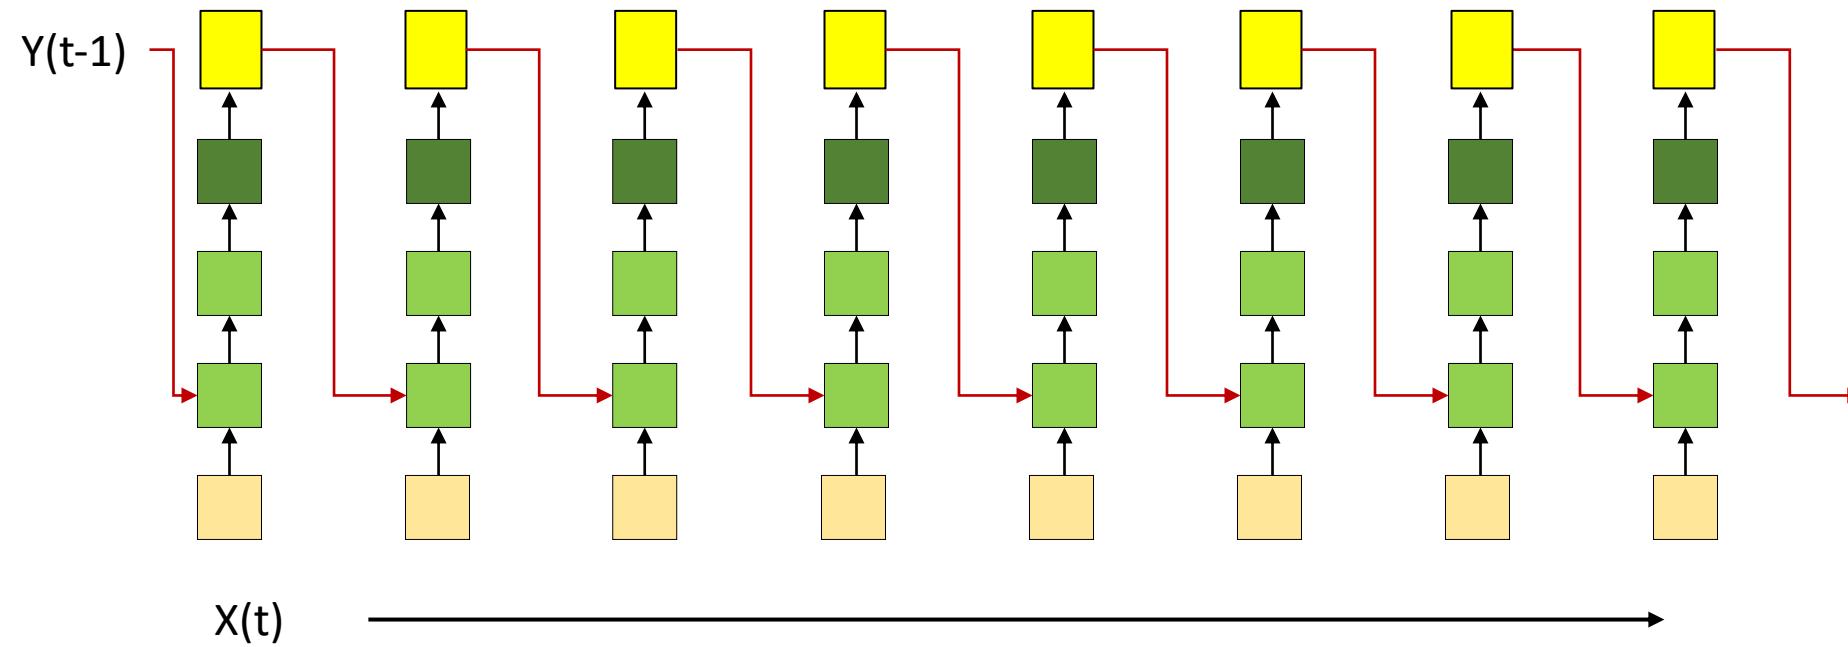
Y_0 produces Y_1 which produces Y_2 and so on until Y_∞

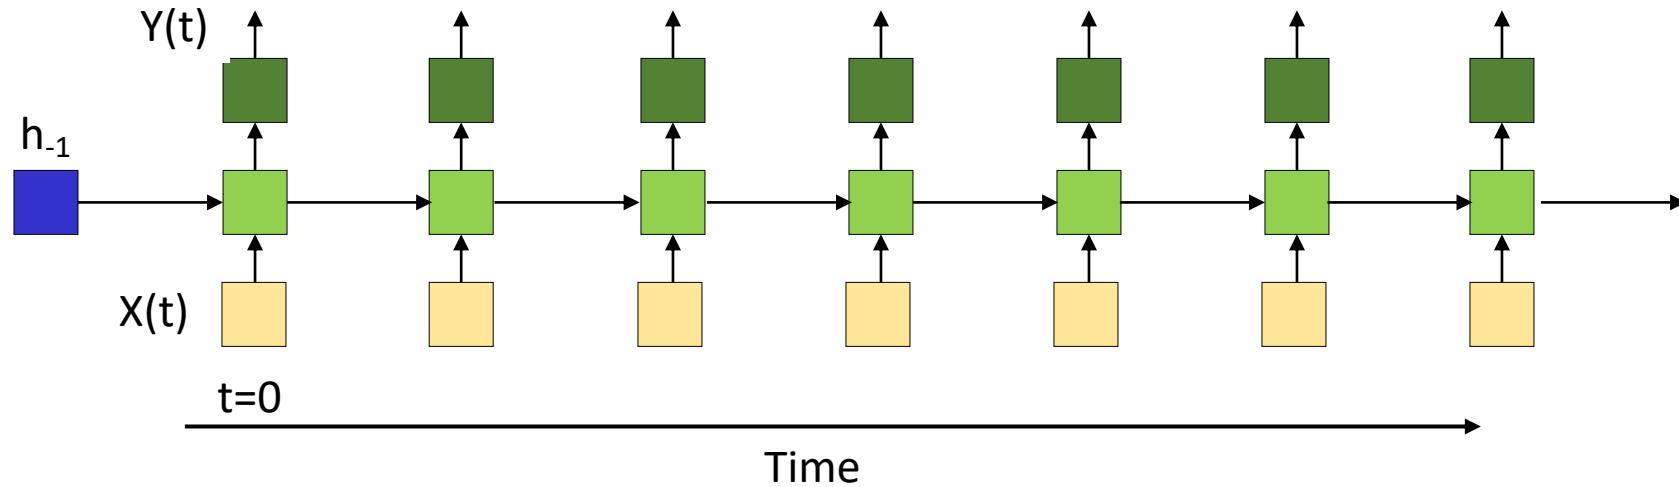






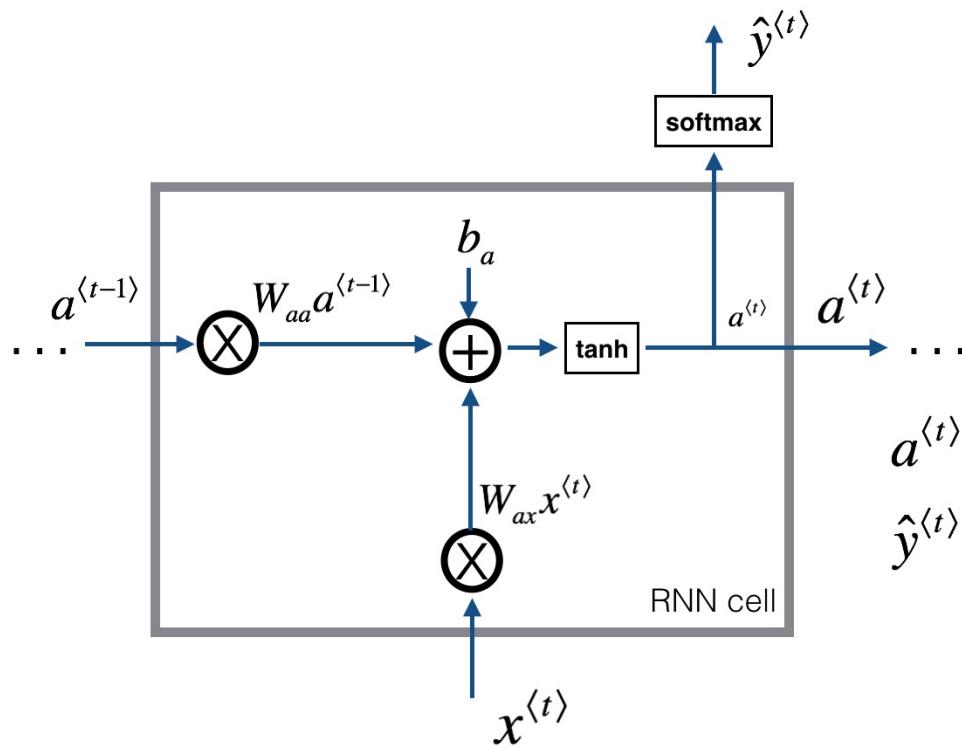
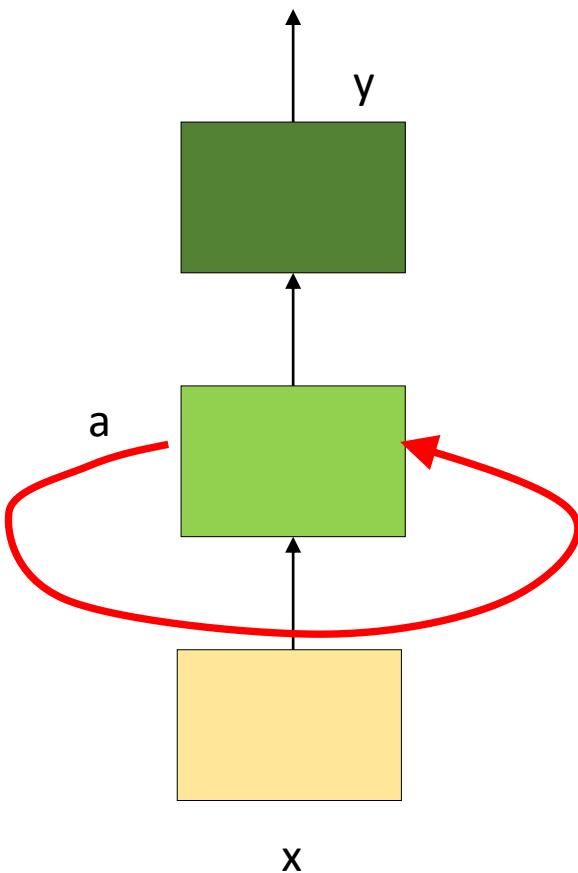






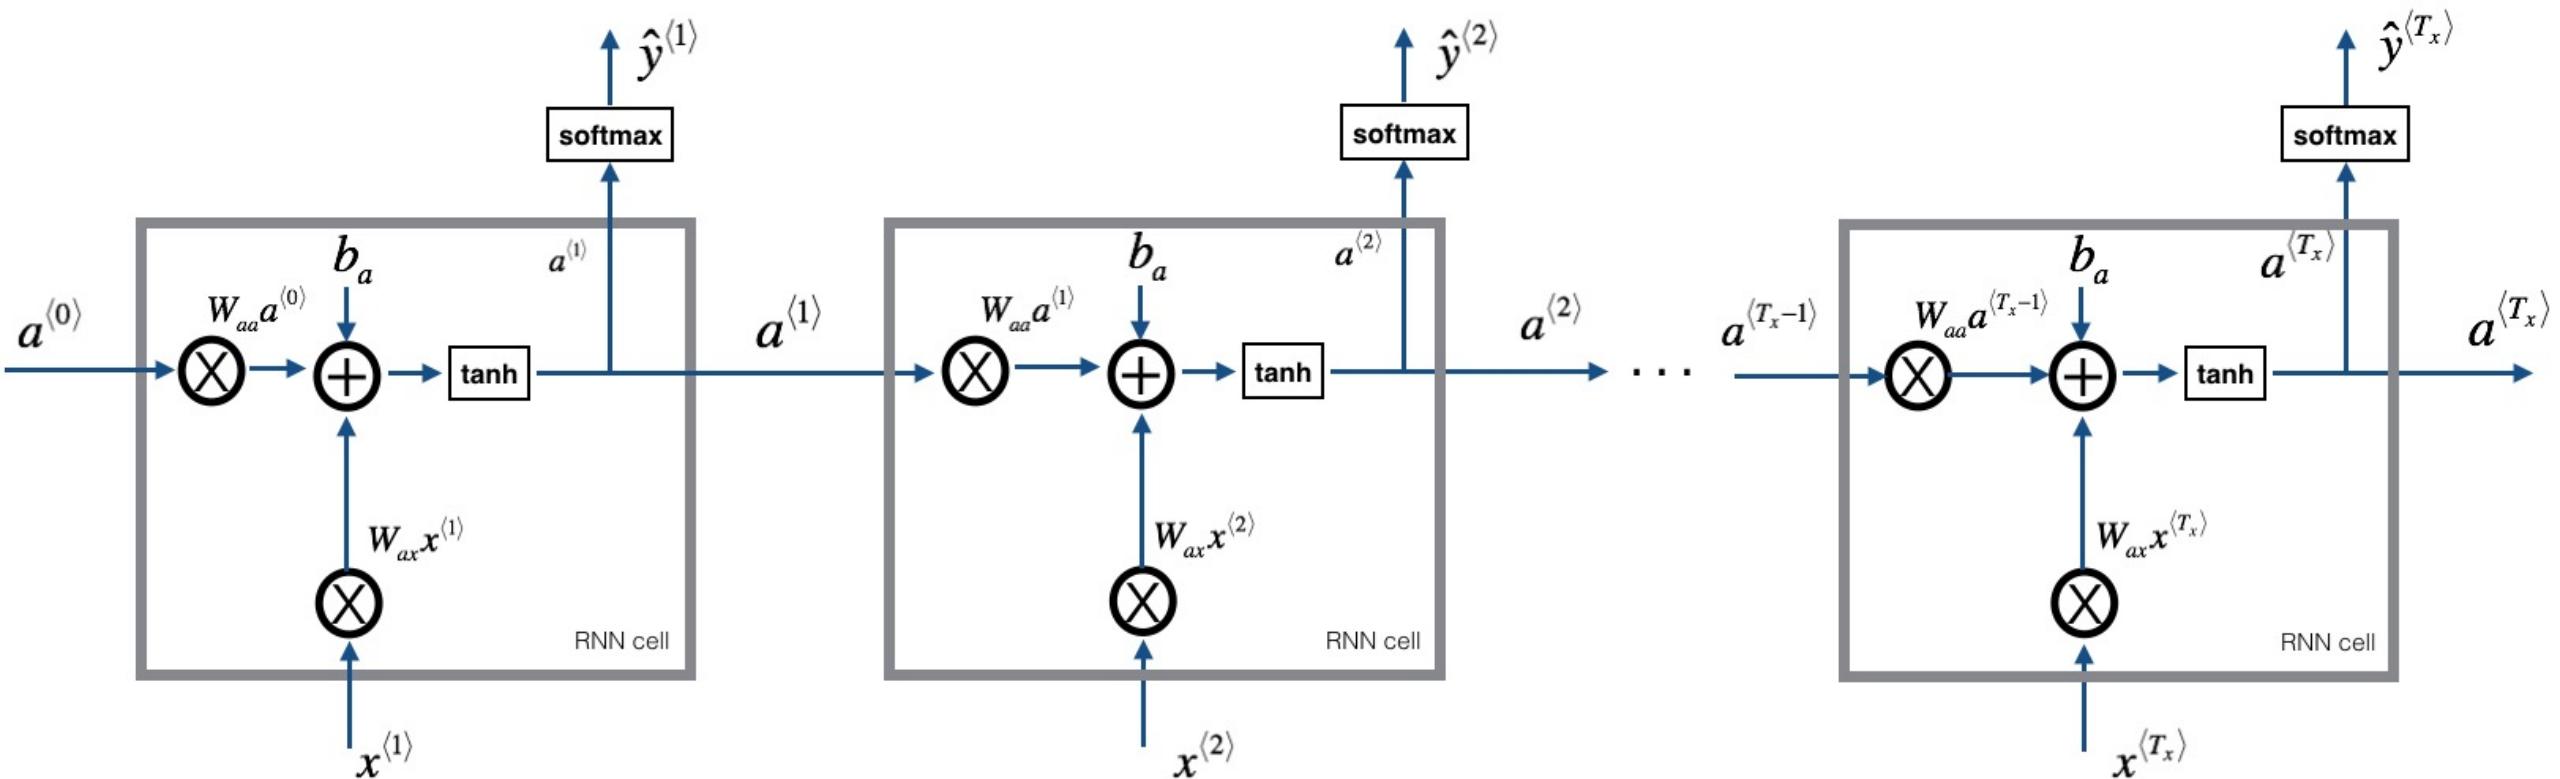
- Recurrent neural network
- All columns are identical
- *An input at $t=0$ affects outputs forever*

Recurrent Neural Networks



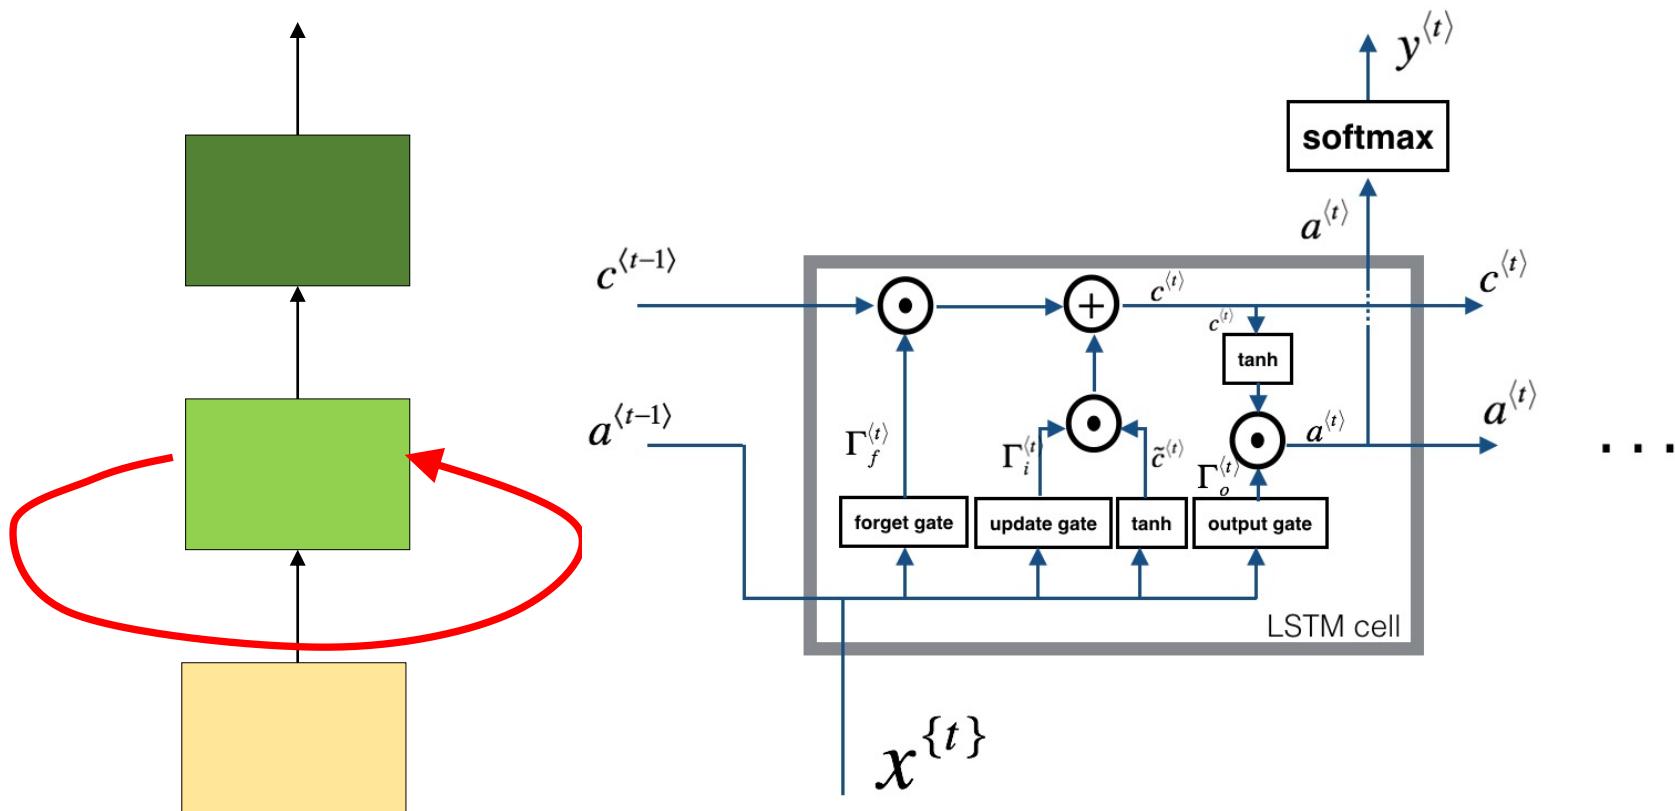
$$a^{(t)} = \tanh(W_{ax}x^{(t)} + W_{aa}a^{(t-1)} + b_a)$$
$$\hat{y}^{(t)} = \text{softmax}(W_{ya}a^{(t)} + b_y)$$

Recurrent Neural Networks



$$a^{(t)} = \tanh(W_{ax}x^{(t)} + W_{aa}a^{(t-1)} + b_a)$$
$$\hat{y}^{(t)} = \text{softmax}(W_{ya}a^{(t)} + b_y)$$

LSTM



$$\Gamma_f^{(t)} = \sigma(W_f[a^{(t-1)}, x^{(t)}] + b_f)$$

$$\Gamma_u^{(t)} = \sigma(W_u[a^{(t-1)}, x^{(t)}] + b_u)$$

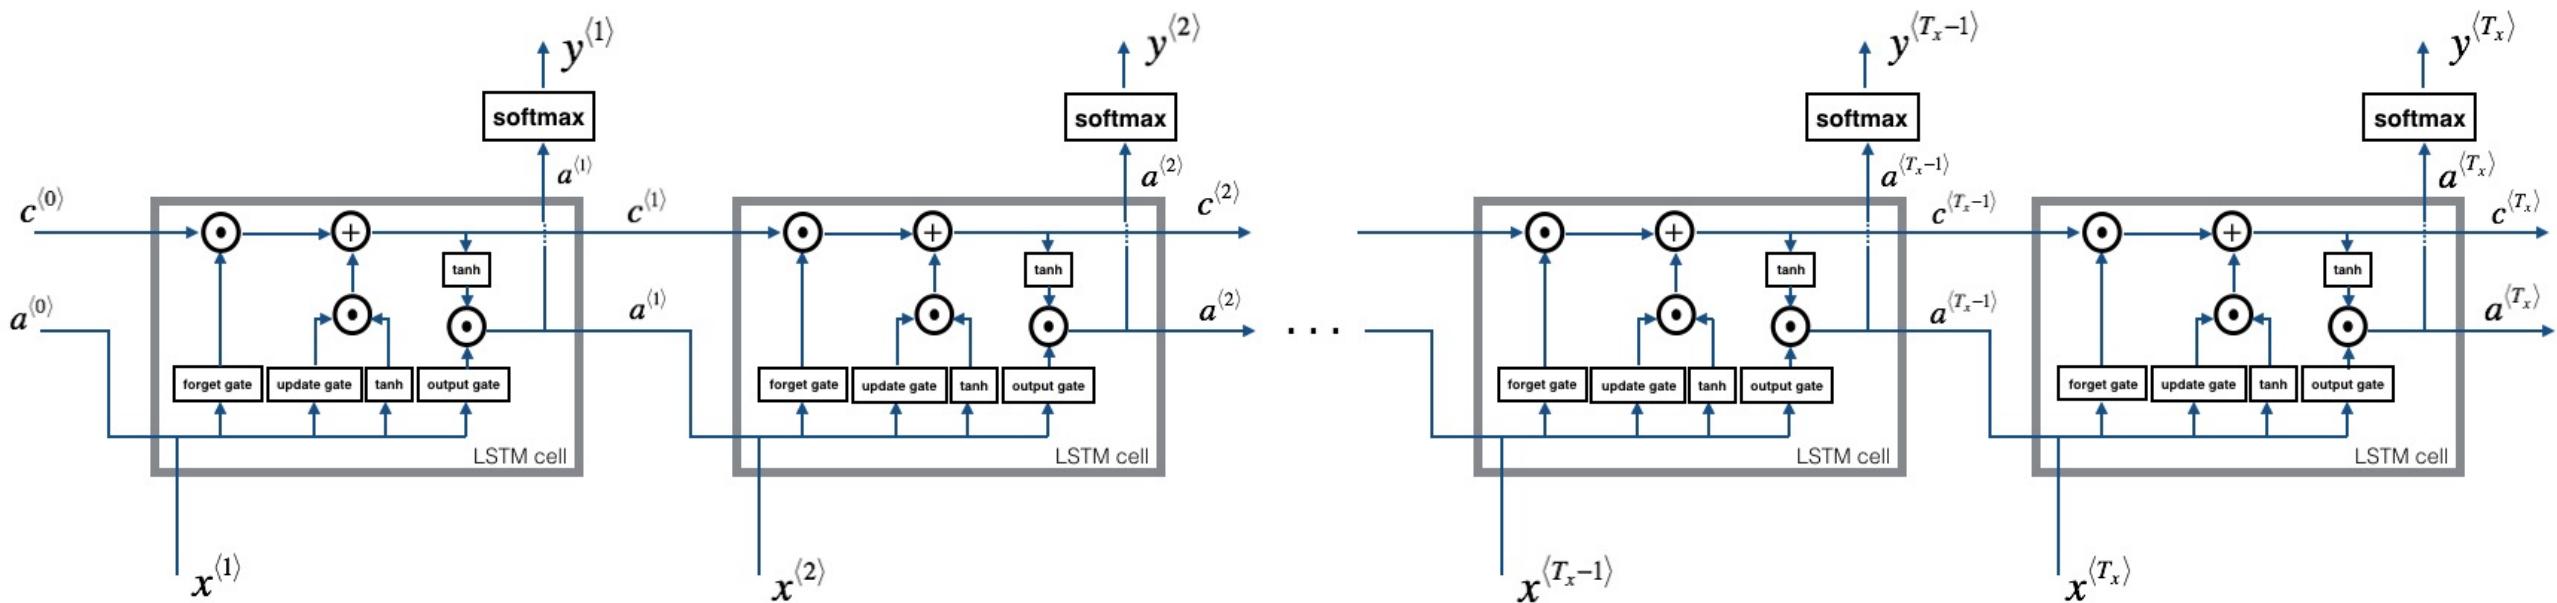
$$\tilde{c}^{(t)} = \tanh(W_C[a^{(t-1)}, x^{(t)}] + b_C)$$

$$c^{(t)} = \Gamma_f^{(t)} \circ c^{(t-1)} + \Gamma_u^{(t)} \circ \tilde{c}^{(t)}$$

$$\Gamma_o^{(t)} = \sigma(W_o[a^{(t-1)}, x^{(t)}] + b_o)$$

$$a^{(t)} = \Gamma_o^{(t)} \circ \tanh(c^{(t)})$$

LSTM



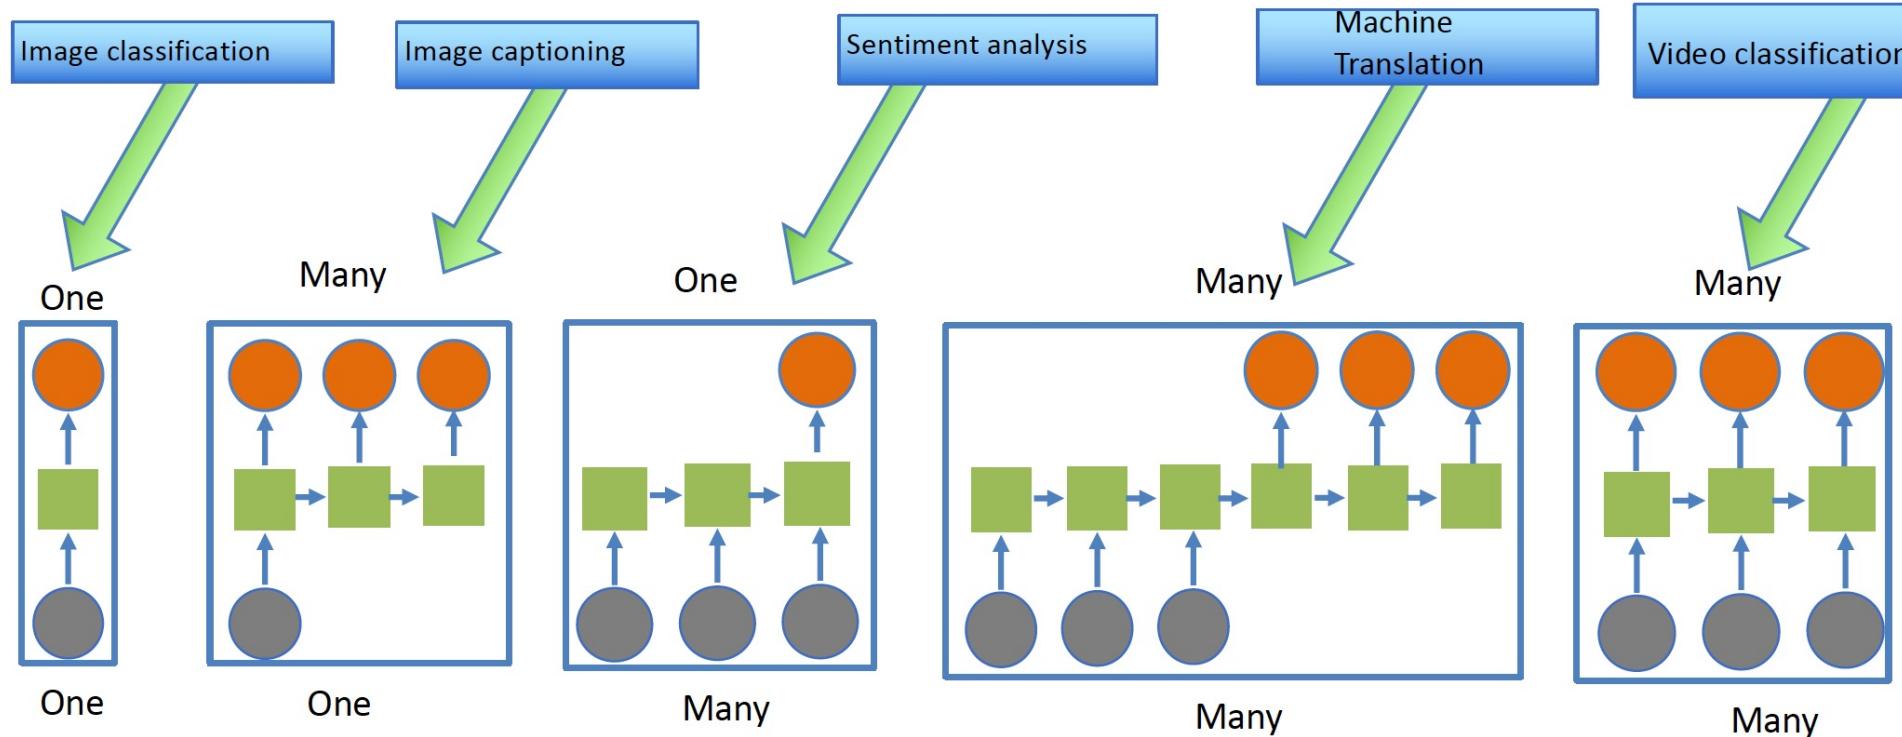
RNN VS. LSTM VS. GRU

<https://jhui.github.io/2017/03/15/RNN-LSTM-GRU/>

https://datascience-enthusiast.com/DL/Building_a_Recurrent_Neural_Network-Step_by_Step_v1.html

<https://blog.floydhub.com/a-beginners-guide-on-recurrent-neural-networks-with-pytorch/>

Variants on recurrent nets



Input

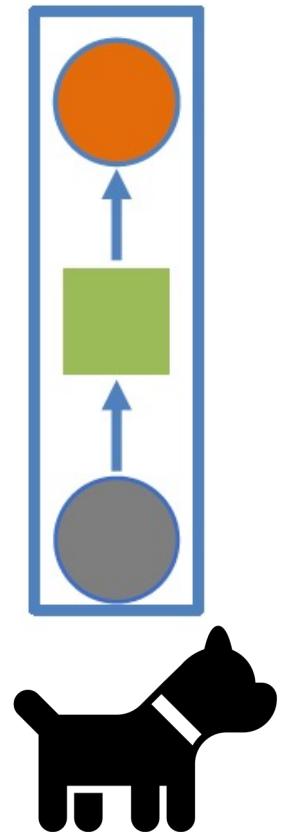
Recurrent core (Hidden unit)

Output

One-One

Label: dog

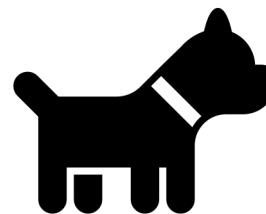
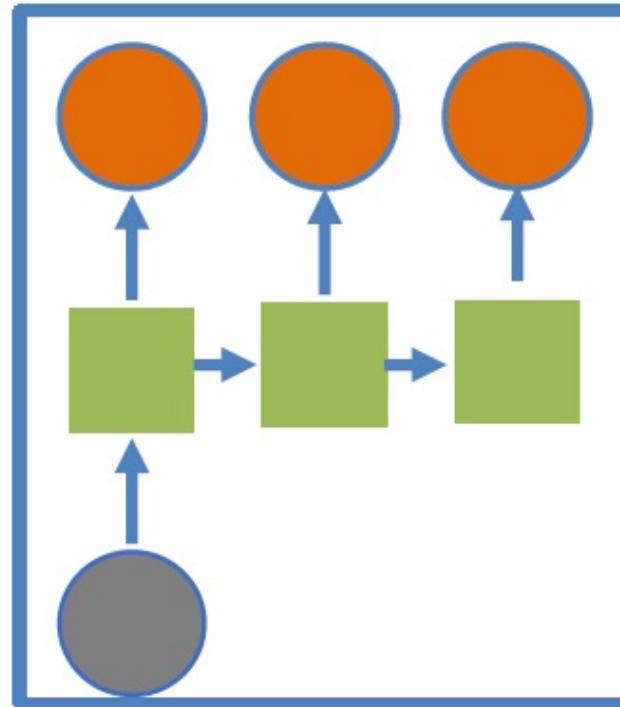
Image Classification

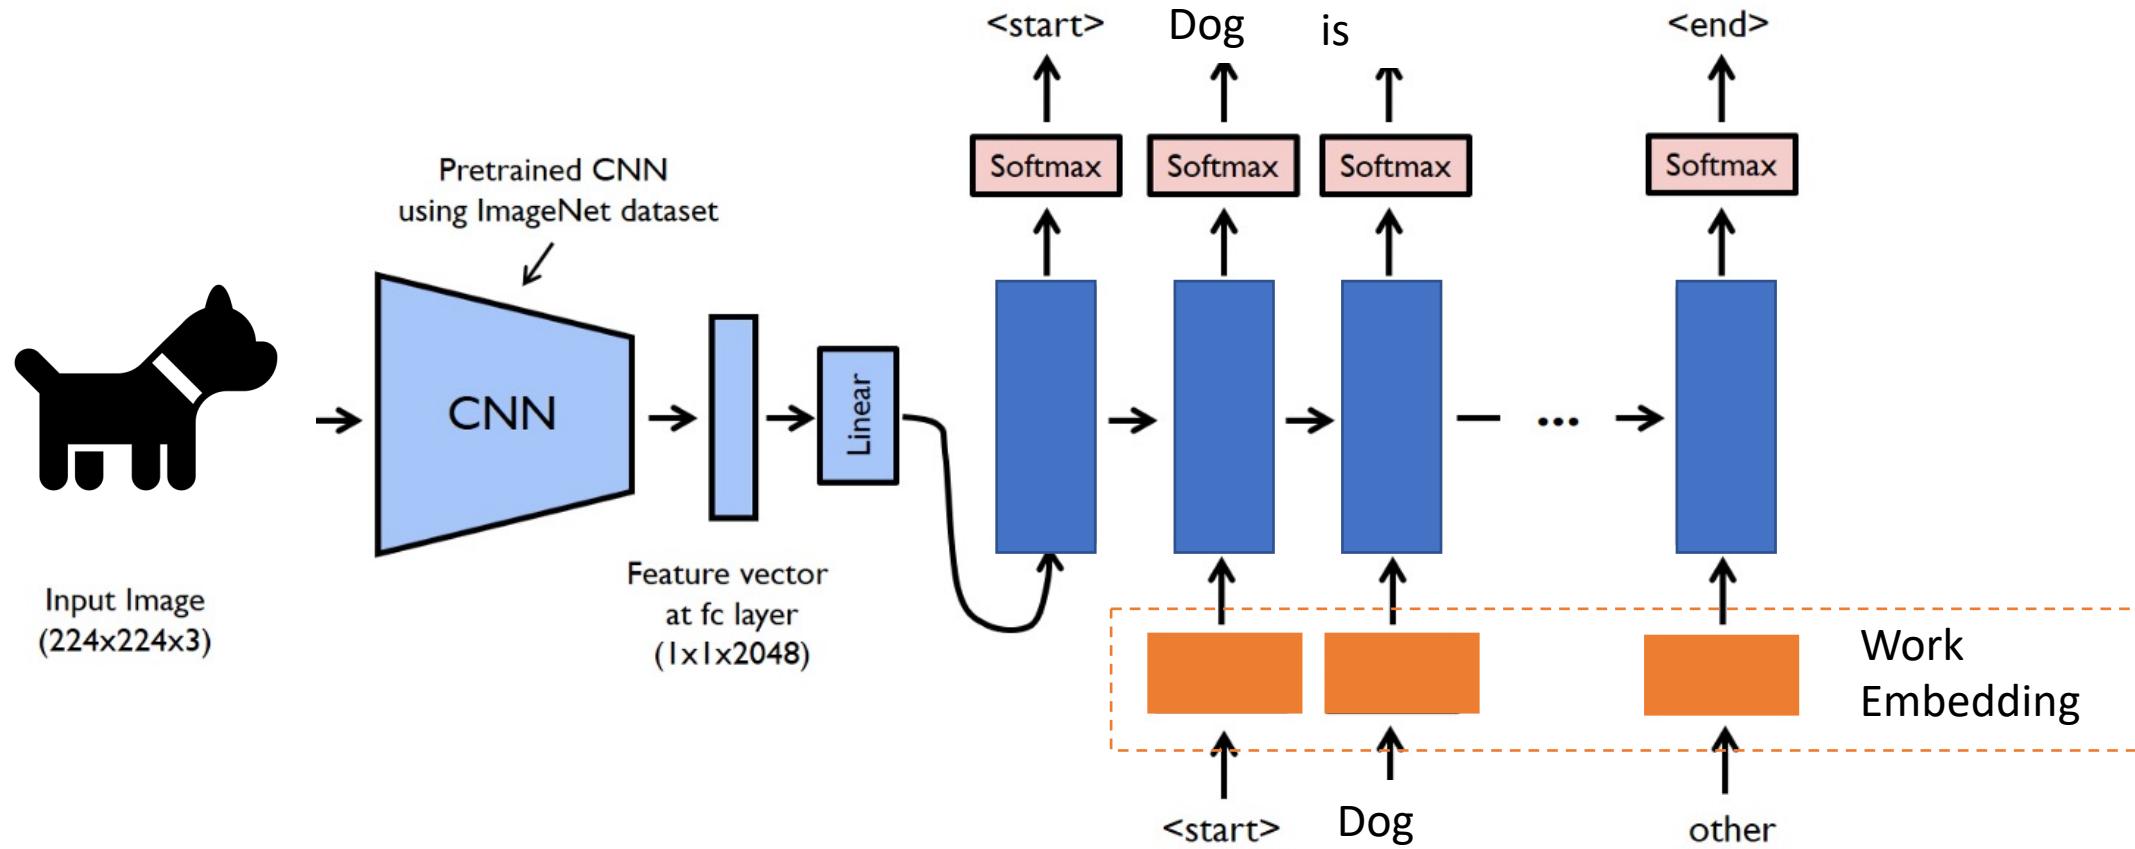


One-Many

Image Captioning

Dog is standing

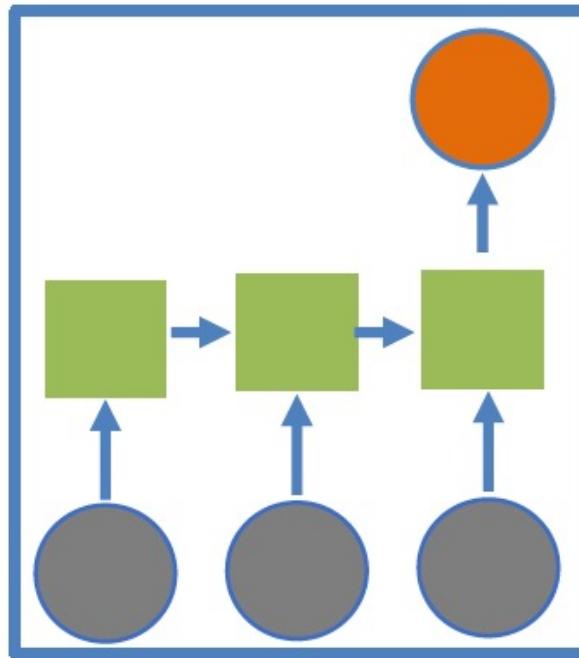




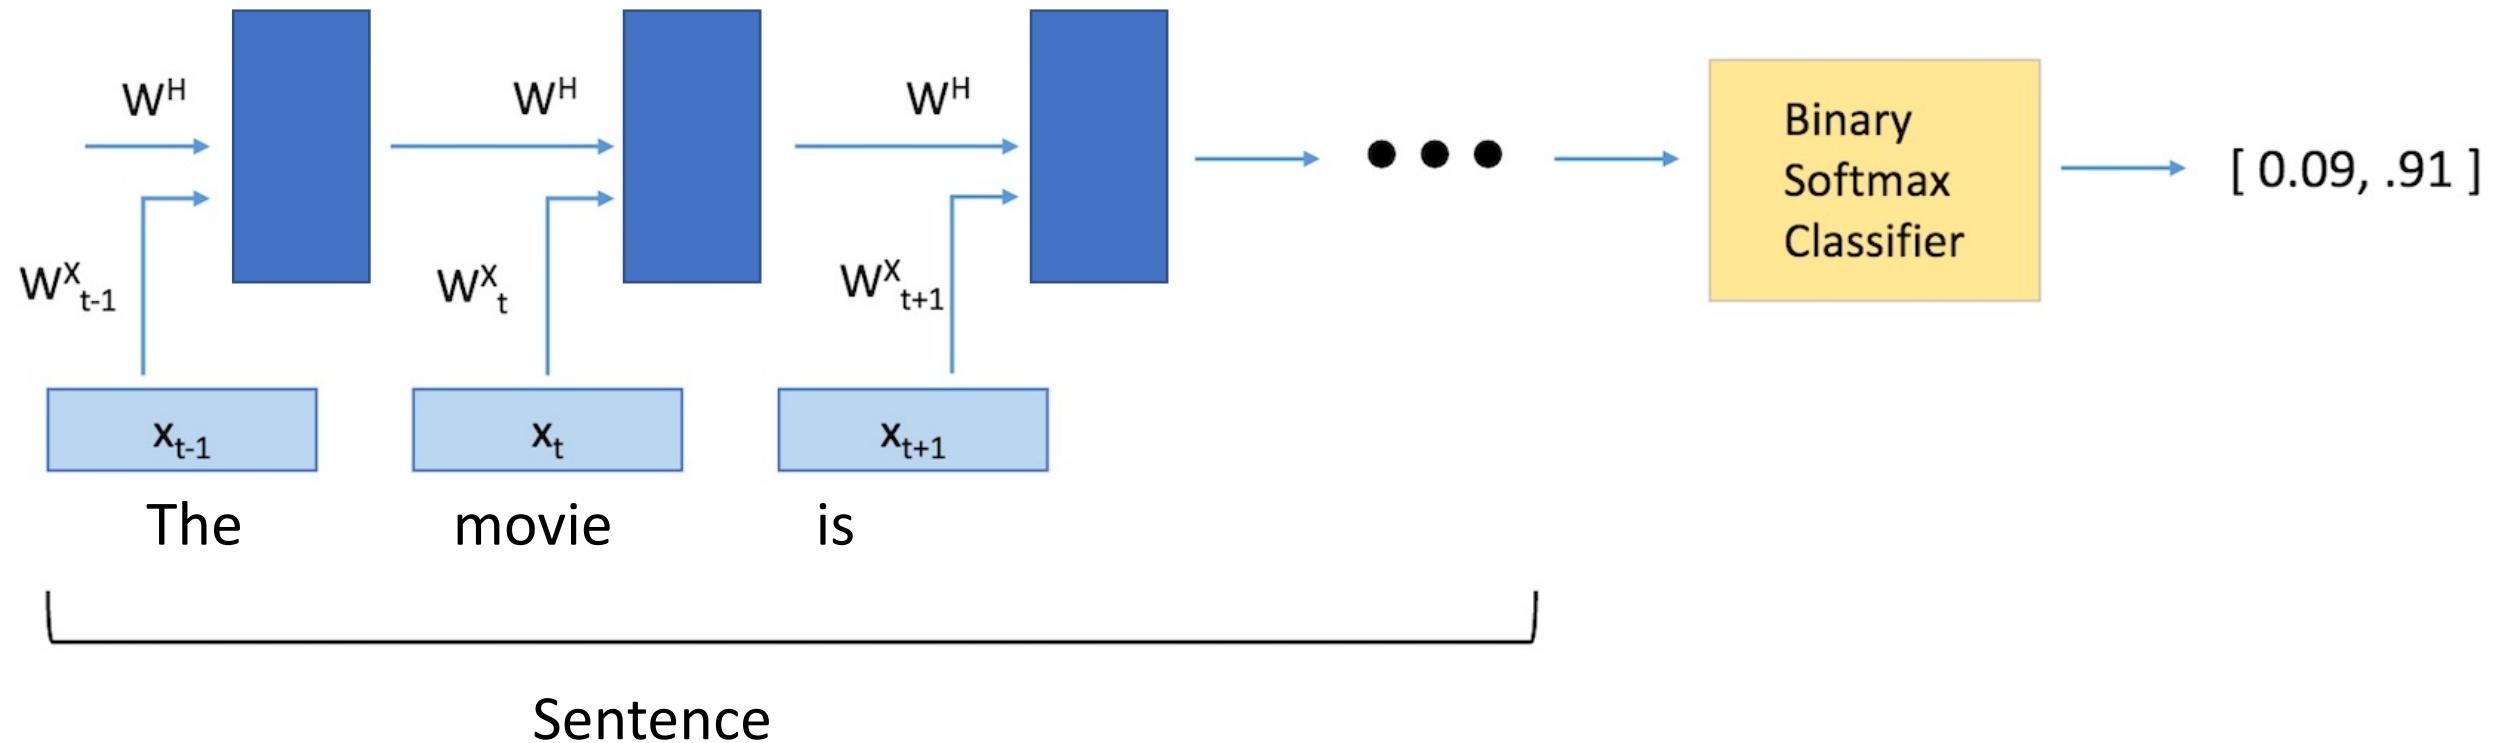
Many-One

Sentiment Analysis

Ranking: 0.8

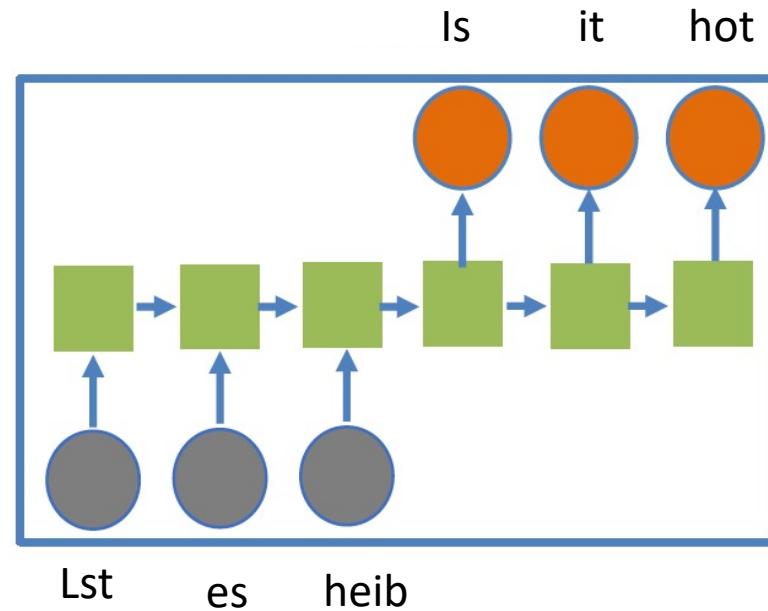


The movie is interesting

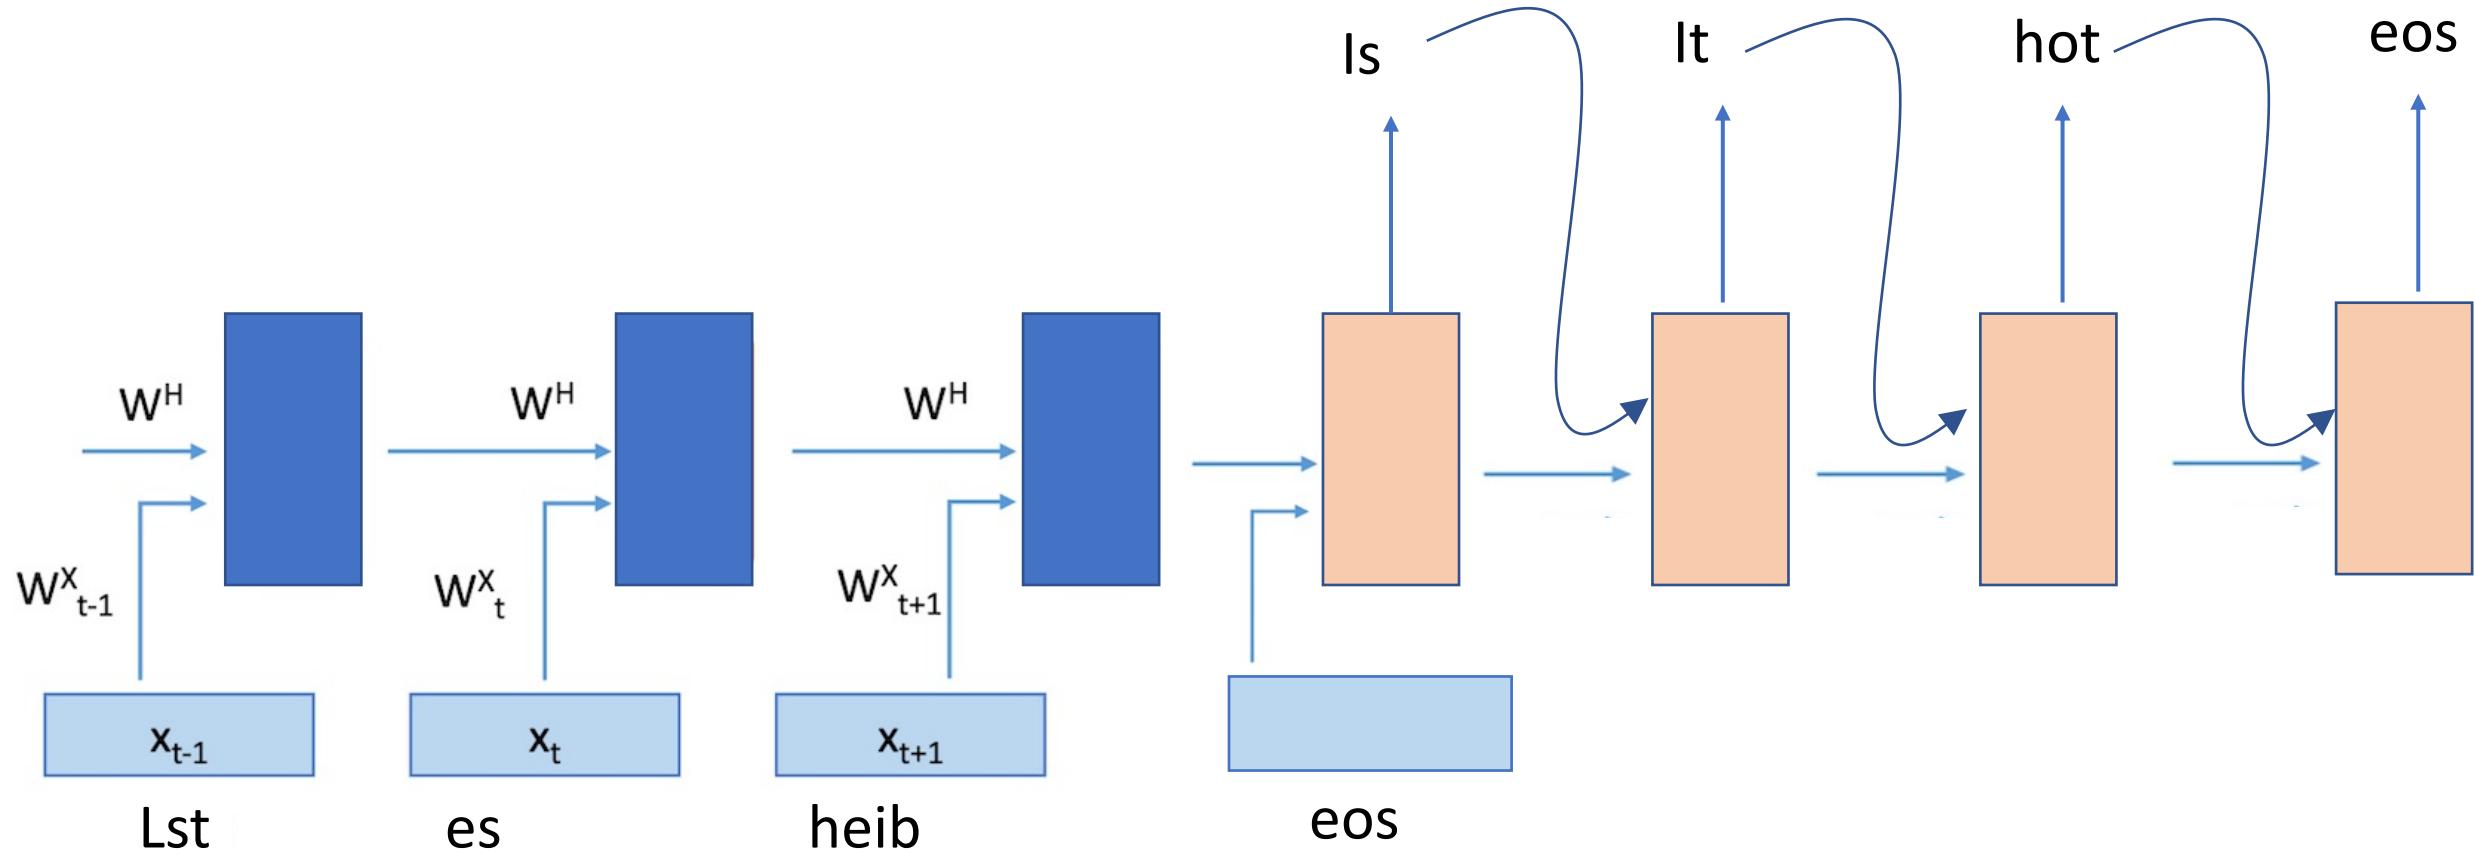


Many-Many

Language Translation



Many-Many



Vanilla RNN

```
model = keras.models.Sequential([
    keras.layers.SimpleRNN(2,
        return_sequences=True,
        input_shape=[None, 1]),
    keras.layers.Dropout(0.3),
    keras.layers.SimpleRNN(1),
    keras.layers.Dense(1)
])

model.compile(
    loss='mse',
    optimizer='Adam',
    metrics=['mae', 'mse'],
)
```

LSTM

```
model = keras.models.Sequential([
    keras.layers.LSTM(20,
                      return_sequences=True,
                      input_shape=[None, 1]),
    keras.layers.Dropout(0.2),
    keras.layers.LSTM(10),
    keras.layers.Dropout(0.2),
    keras.layers.Dense(1)
])

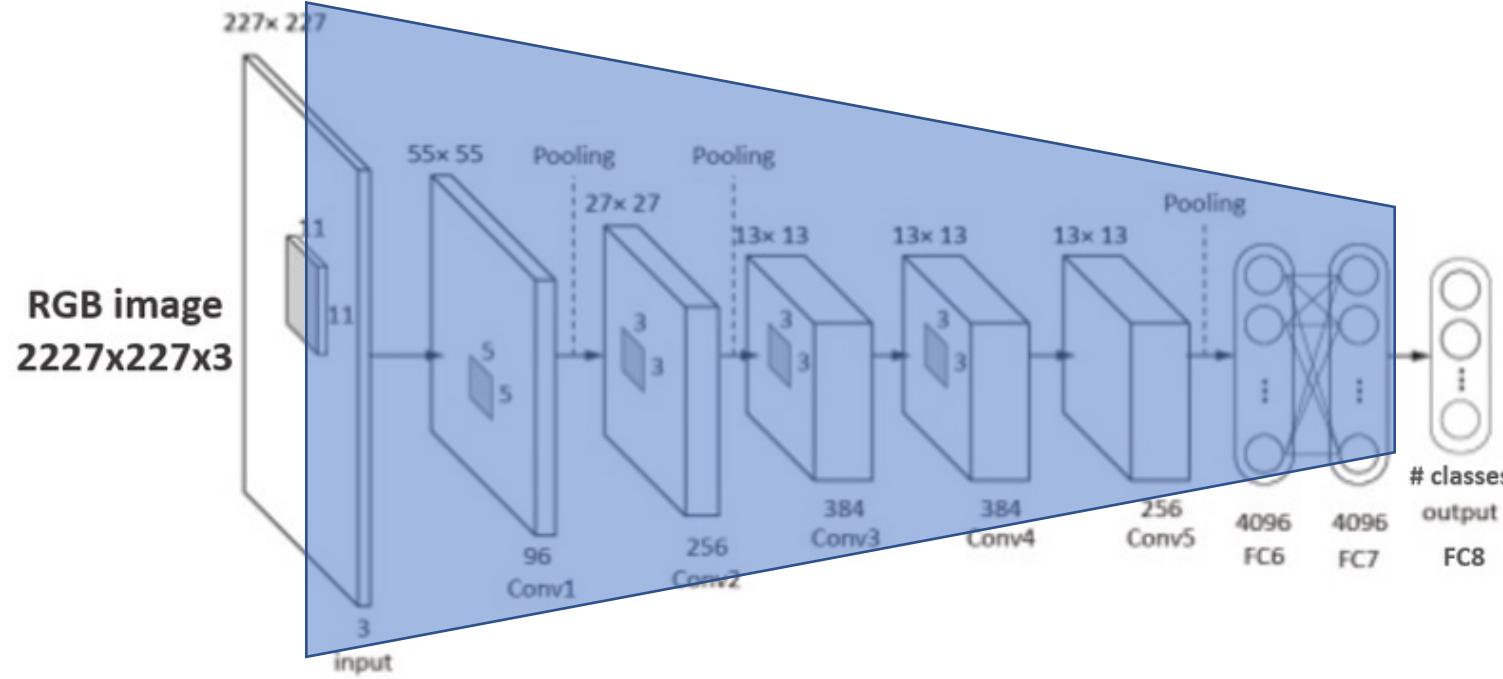
model.compile(
    loss='mse',
    optimizer='Adam',
    metrics=['mae', 'mse'],
)
```

GRU

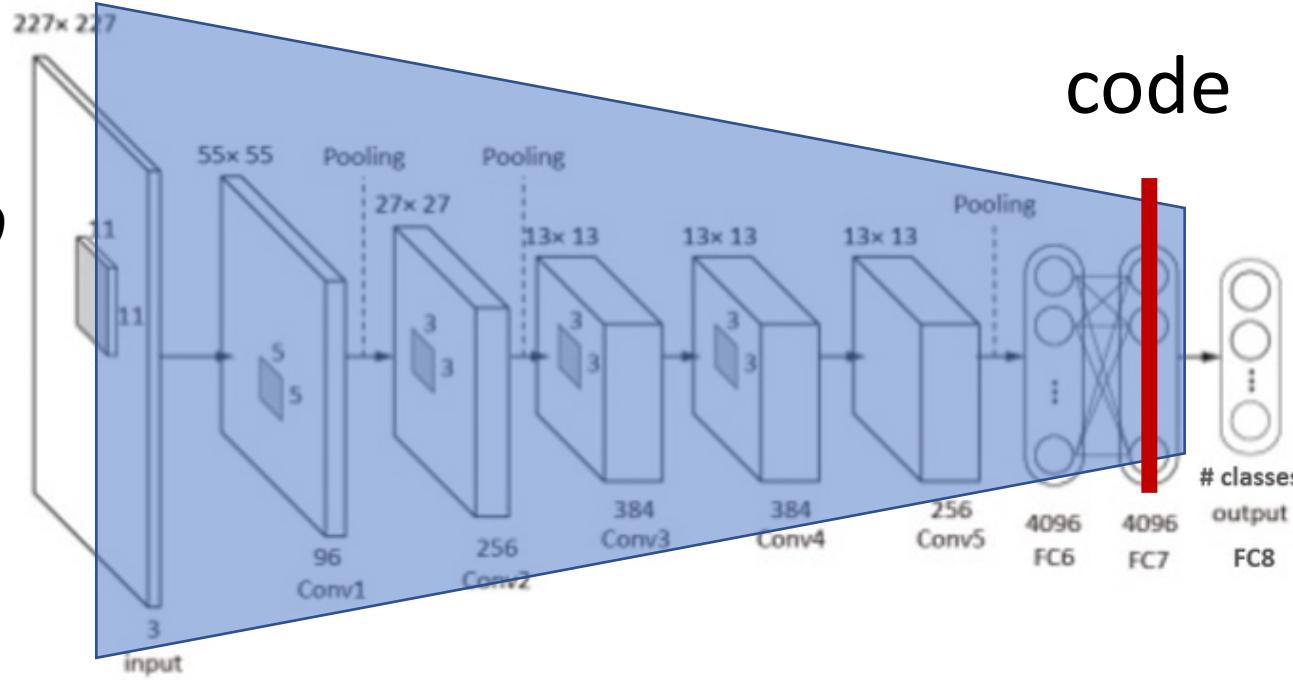
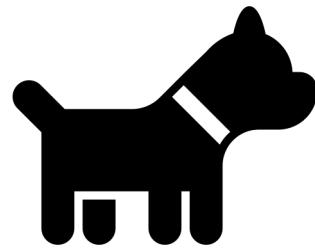
```
model = keras.models.Sequential([
    keras.layers.Conv1D(filters=20, kernel_size=4,
                        strides=2, padding="valid",
                        input_shape=[None, 1]),
    keras.layers.GRU(2, input_shape=[None, 1],
                     activation='relu'),
    keras.layers.Dropout(0.2),
    keras.layers.Dense(1),
])

model.compile(
    loss='mse',
    optimizer='Adam',
    metrics=['mae', 'mse'],
)
```

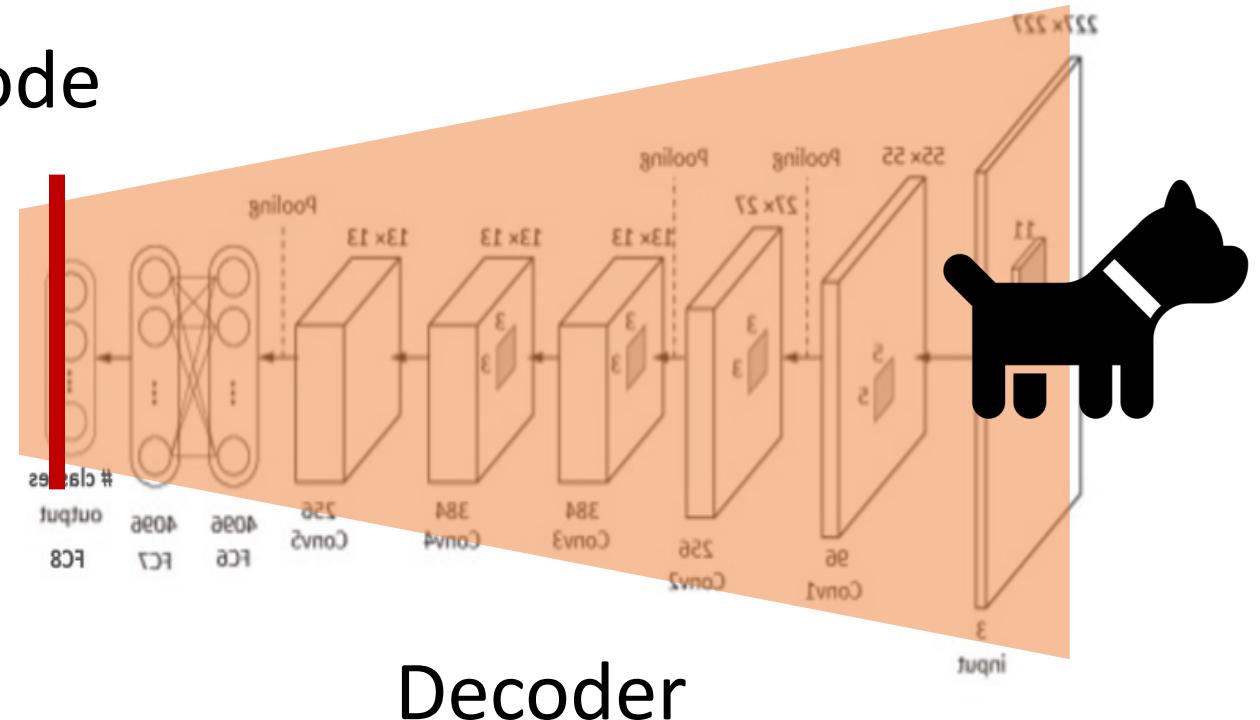
Autoencoder



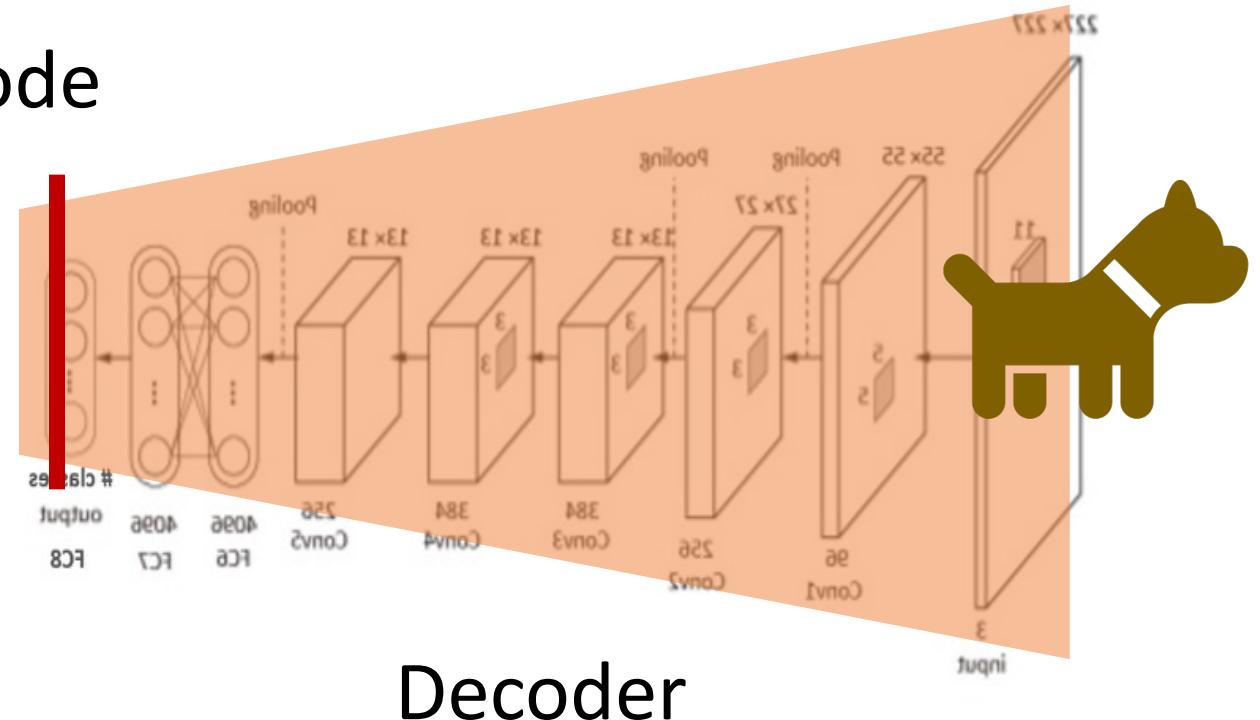
Encoder

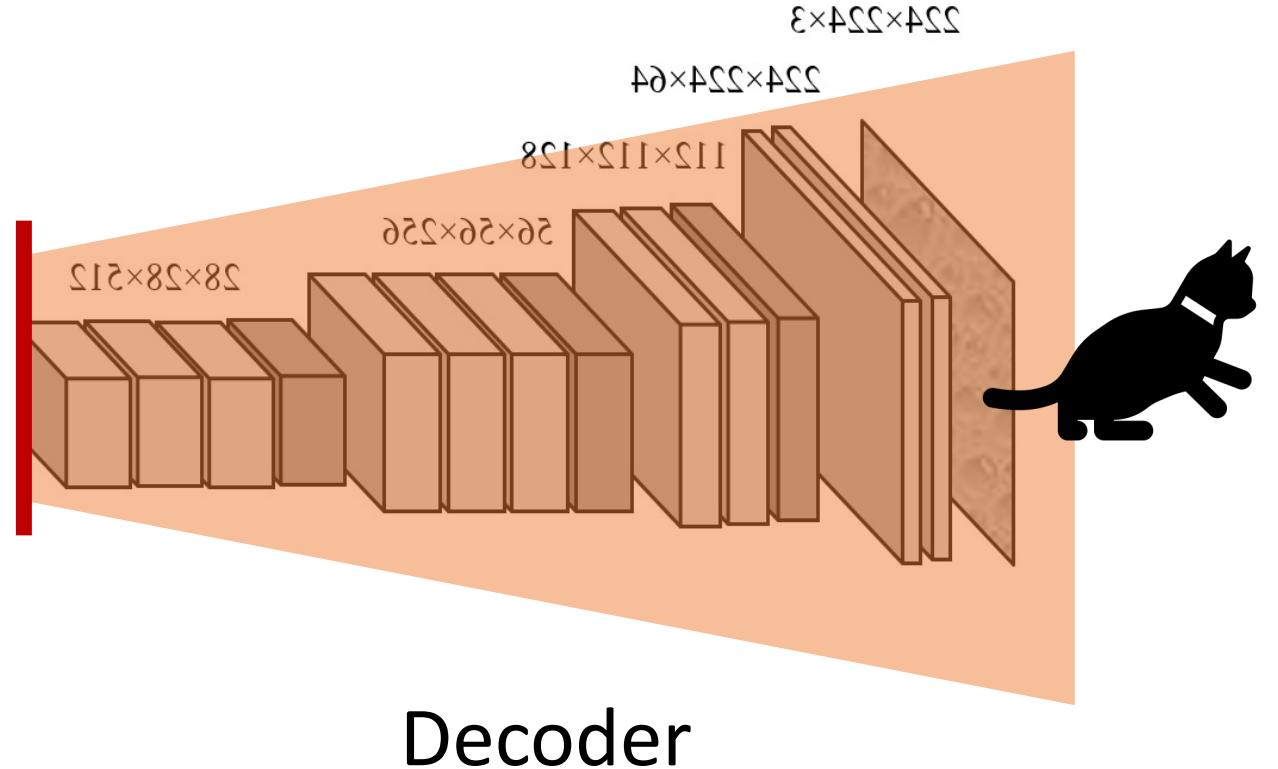


code

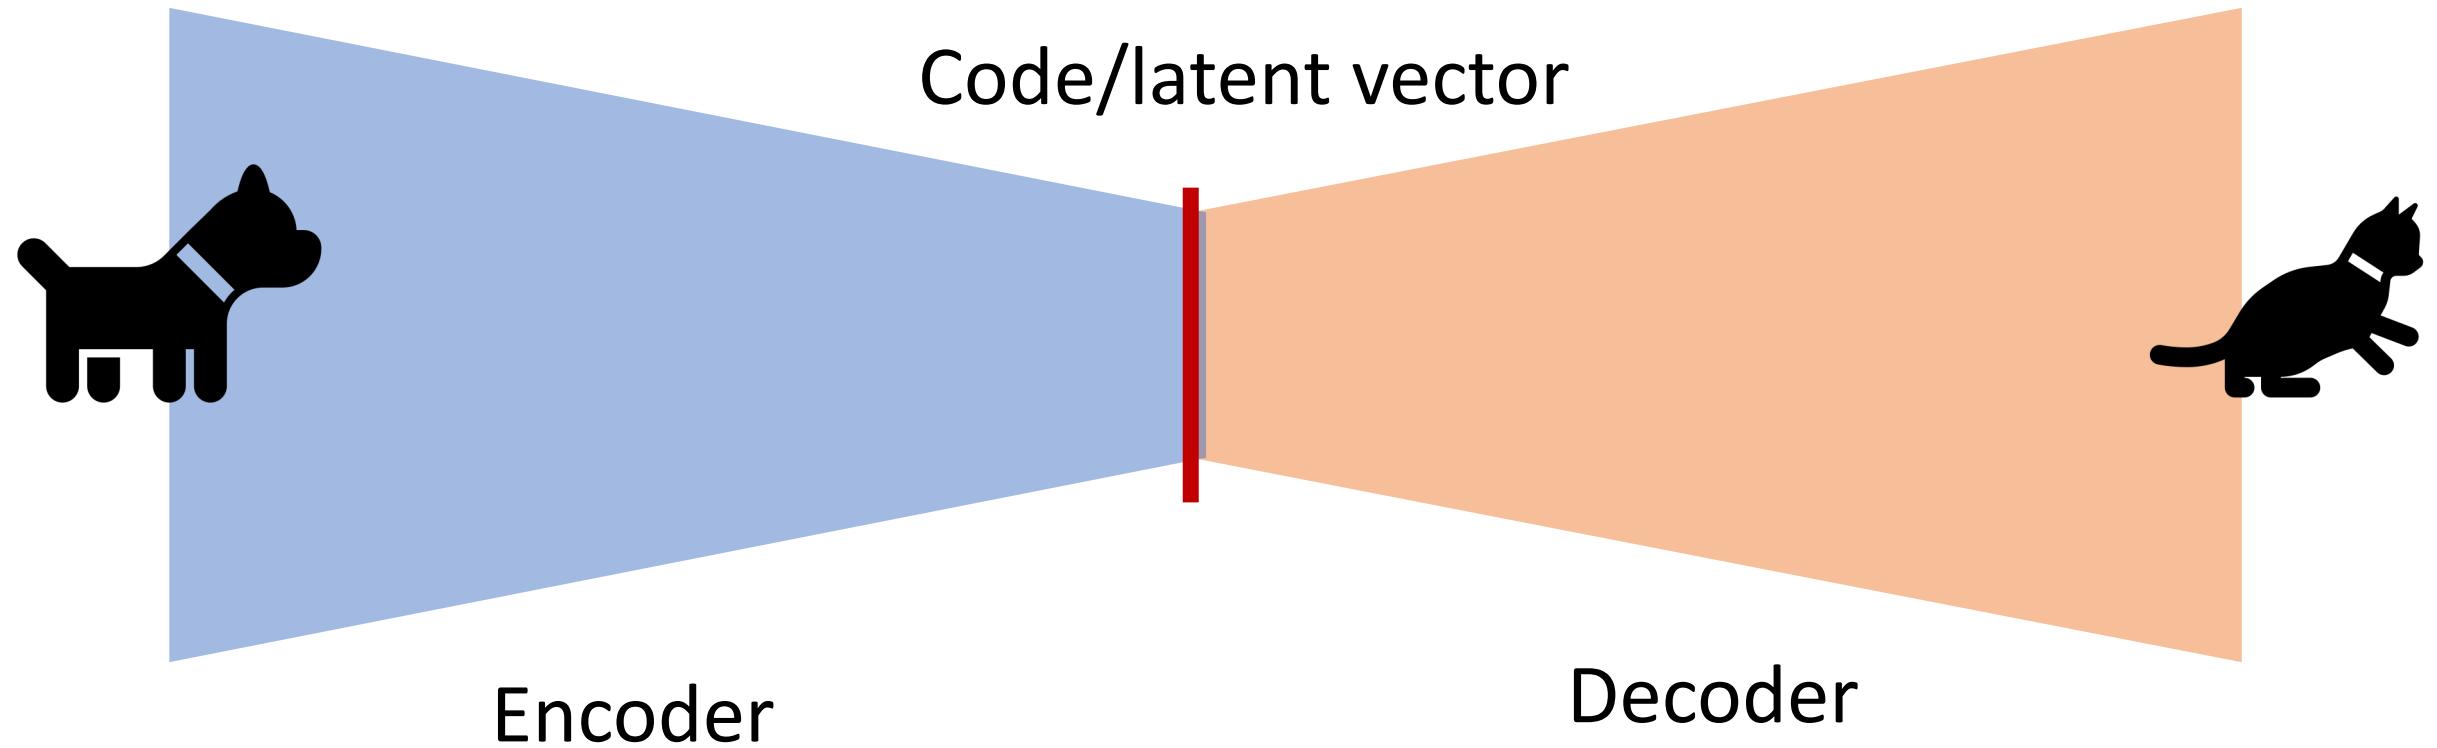


code





Autoencoder



```
filter_1 = 64
filter_2 = 32
filter_3 = 16
kernel_size = (4, 4)
pool_size = (2, 2)
latent_dim = 16

input_layer = Input(shape=(60, 80, 3))

encoder_conv1 = Conv2D(
    filter_1,
    kernel_size,
    activation='sigmoid',
    padding='same')(input_layer)

encoder_pool1 = MaxPool2D(
    pool_size,
    padding='same')(encoder_conv1)

encoder_conv2 = Conv2D(
    filter_2,
    kernel_size,
    activation='sigmoid',
    padding='same')(encoder_pool1)

code_output = MaxPool2D(
    pool_size,
    padding='same')(encoder_conv2)

encoder = Model(input_layer, code_output, name='encoder')
```

```
code_inputs = Input(shape=(15, 20, 32), name="decoder_input")

decoder_conv1 = Conv2D(
    filter_2,
    kernel_size,
    activation='sigmoid',
    padding='same')(code_inputs)

decoder_up1 = UpSampling2D(
    pool_size)(decoder_conv1)

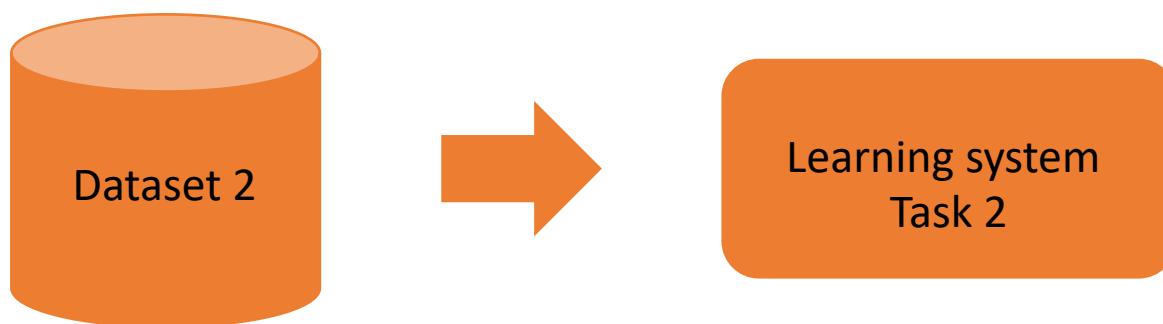
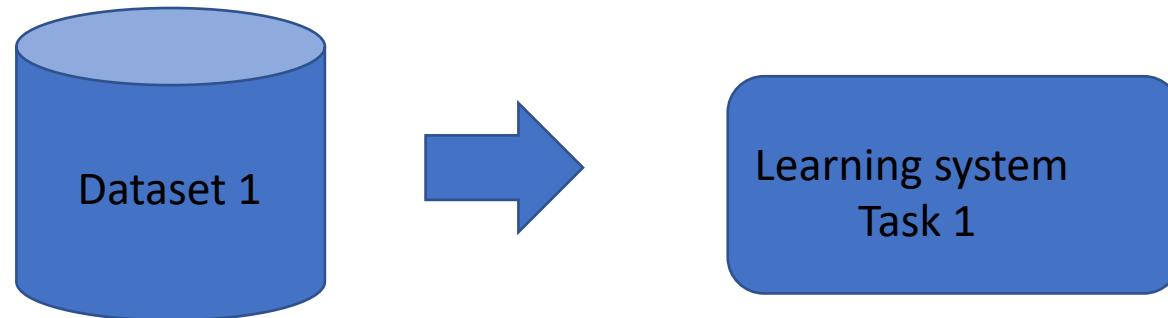
decoder_conv2 = Conv2D(
    filter_1,
    (1,1),
    activation='sigmoid',
    name="conv3")(decoder_up1)

decoder_up2 = UpSampling2D(
    pool_size,
    name="up3")(decoder_conv2)

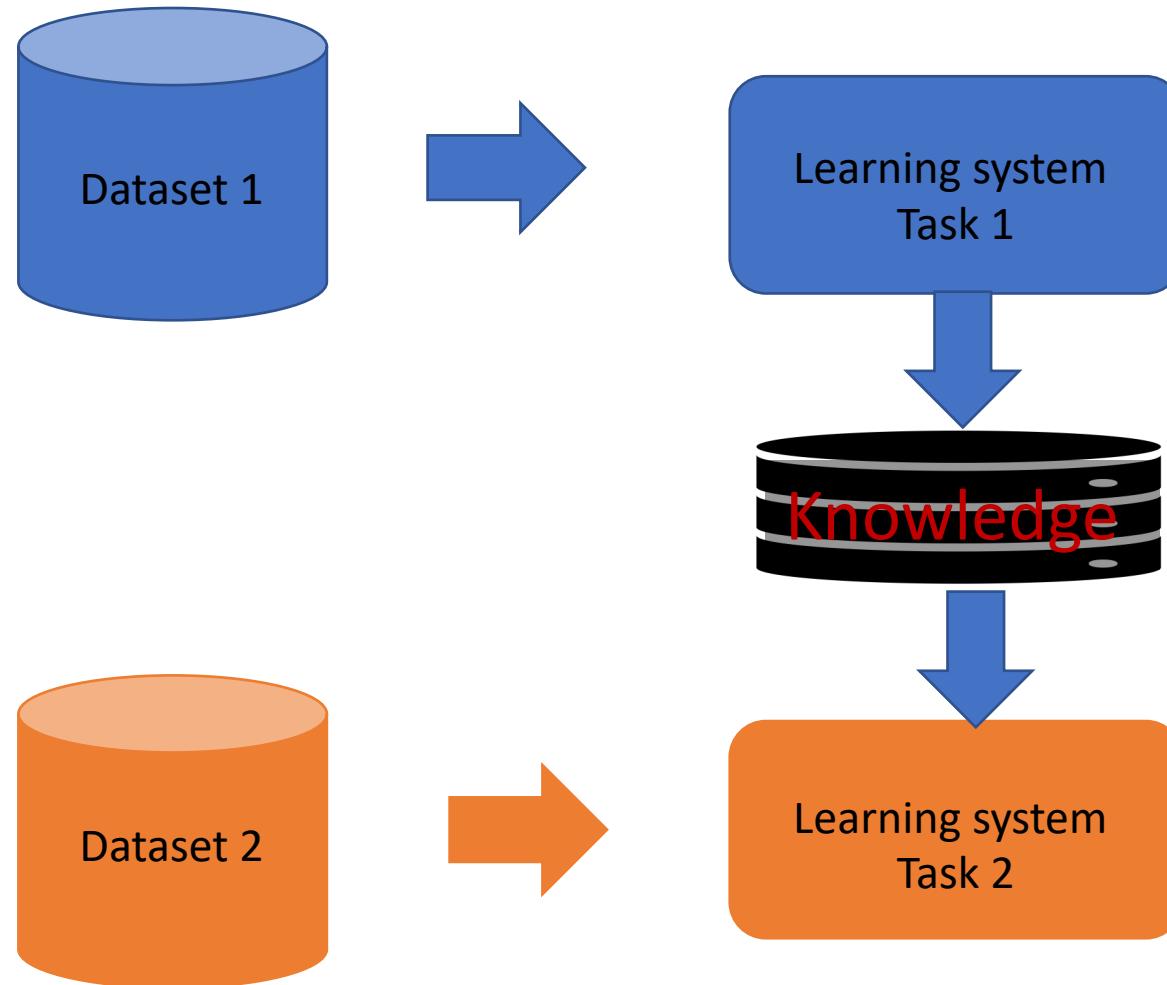
output_layer = Conv2D(
    3,
    (kernel_size),
    padding='same',
    activation='sigmoid')(decoder_up2)

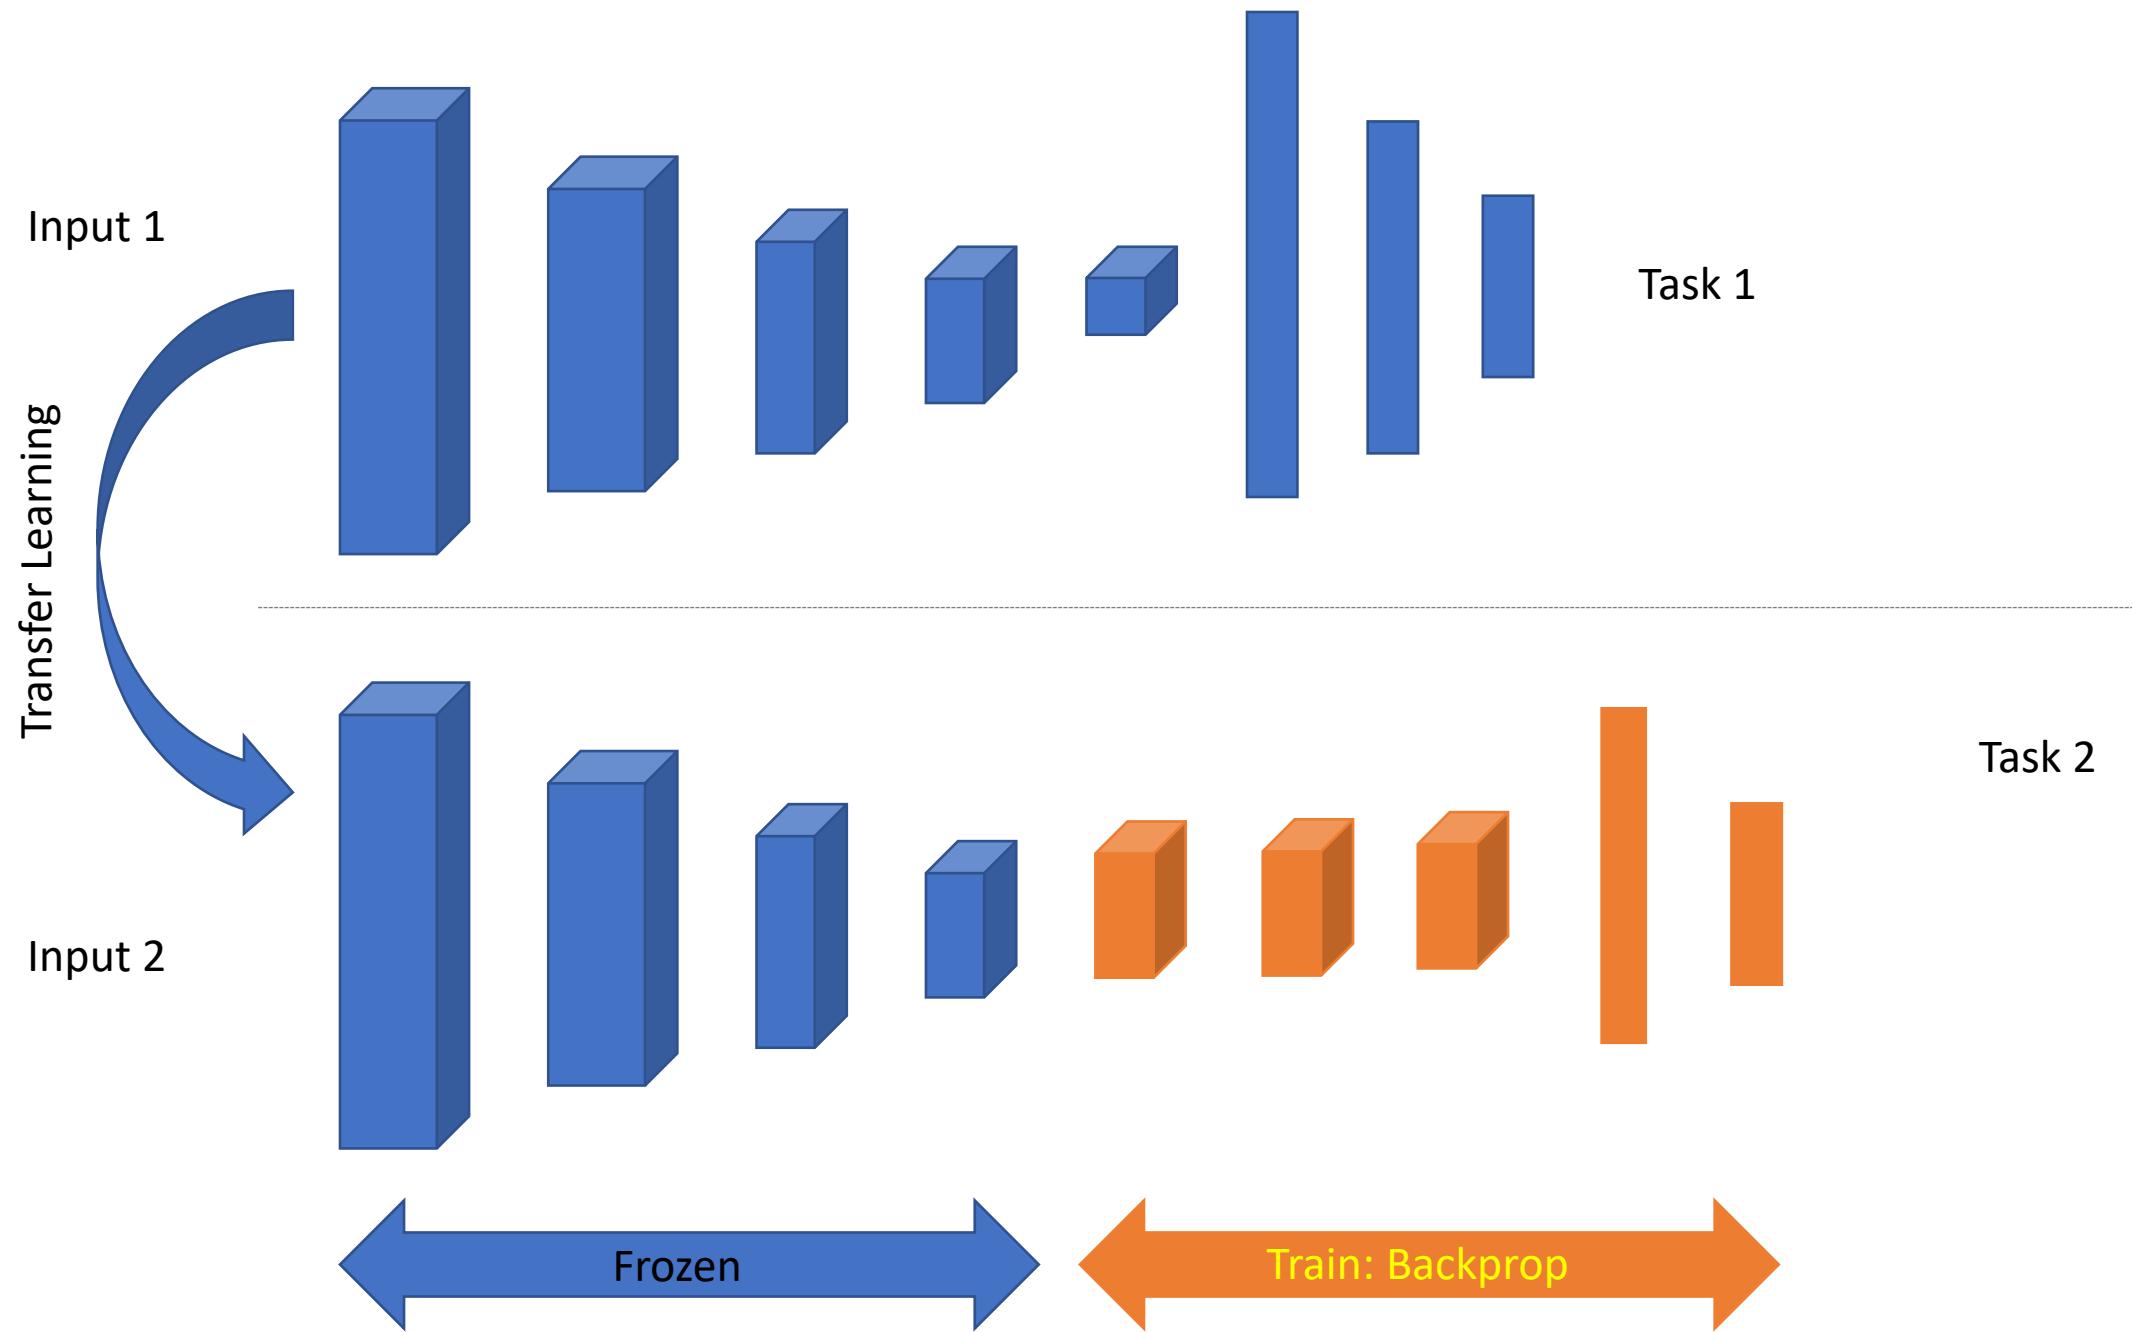
decoder = Model(code_inputs, output_layer, name='decoder')
```

Transfer Learning



Transfer Learning





```
# Load `DenseNet201`.
```

```
1 densenet = tf.keras.applications.DenseNet201(input_shape=IMG_SHAPE,
2                                              weights='imagenet',
3                                              include_top=False)
```

```
# Extend the model.
```

```
model = tf.keras.Sequential([
    densenet,
    tf.keras.layers.GlobalAveragePooling2D(),
    tf.keras.layers.Dense(101, activation='softmax'),
])
```

